

732A91: Lab 1

Bayesian Learning

Sarah Alsaadi, Carles Sans Fuentes

April 13, 2017

Question 1

Let $y_1, \dots, y_n | \theta \sim \text{Bern}(\theta)$, and assume that you have obtained a sample with $s = 14$ successes in $n = 20$ trials. Assume a $\text{Beta}(\alpha_0, \beta_0)$ prior for θ and let $\alpha_0 = \beta_0 = 2$.

- Draw random numbers from posterior $\theta | y_1, \dots, y_n \sim \text{Beta}(\alpha_0 + s, \beta_0 + f)$ and verify graphically that the posterior mean and standard deviation converges to the true mean $E[\theta] = \frac{\alpha_0 + s}{\alpha_0 + s + \beta_0 + f} \approx 0.66$ and true standard deviation $\text{Var}(\theta) = \frac{(\alpha_0 + s)(\beta_0 + f)}{(\alpha_0 + s + \beta_0 + f)^2 * (\alpha_0 + s + \beta_0 + f + 1)} \approx 0.09$. Figure 1 and 2 shows how the mean and standard deviation of samples randomly drawn converges to the true mean and true standard deviation with larger samples.




Figure 1: A plot representing the mean of 10000 samples where the sample size goes from 1 to 10000. The theoretical mean is plotted as a line.

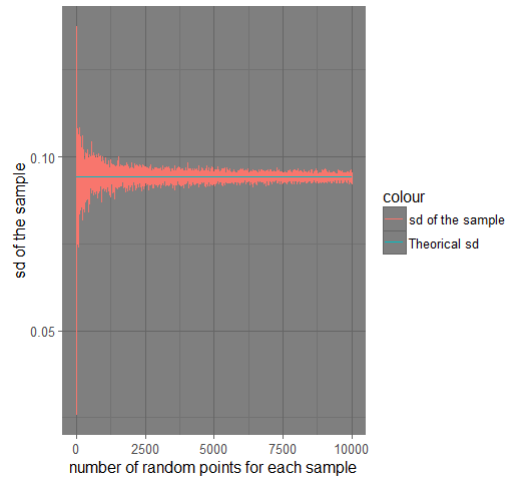


Figure 2: A plot representing the standard deviation of 10000 where the sample size goes from 1 to 10000. The theoretical standard deviation (sd) is plotted as a line.

- Use simulation (nDraws= 10000) to compute the posterior probability $Pr(\theta < 0.4|y)$ and compare with the exact value $Pr(\theta < 0.4|y) = 0.00397$. The simulated value is given below, we get 0.0036 which is close enough to the exact value.

```
1 pbeta(0.4,16,8)
2 [1] 0.003972563
3 simulated pbeta
4 [1] 0.0036
```

- Compute the posterior distribution of the log-odds $\phi = \log(\frac{\theta}{1-\theta})$ by simulation with nDraws= 10000. Figure 3 below shows the histogram together with the kernel density of the data simulated from the posterior distribution of the log-odds $\phi = \log(\frac{\theta}{1-\theta})$ with nDraws= 10000.




Figure 3: Histogram and kernel density of the data simulated from the posterior distribution of the log-odds $\phi = \log(\frac{\theta}{1-\theta})$ with nDraws= 10000.

Question 2

Assume that you have asked 10 randomly selected persons about their monthly income (in thousands Swedish Krona) and obtained the following 10 observations

$(y_1, \dots, y_{10}) = (14, 25, 45, 25, 30, 33, 19, 50, 34, 67)$. Assume $y_1, \dots, y_n | \mu, \sigma^2 \sim \text{logNormal}(\mu, \sigma^2)$,

$\mu = 3.5$ and a non-informative prior $p(\sigma^2) \propto \frac{1}{\sigma^2}$. It can be shown that the posterior for σ^2 is $Inv - \chi^2(n, \tau^2)$ (scaled), where

$$\tau^2 = \frac{\sum_{i=1}^n (\log(y_i) - \mu)}{n}$$

- Simulate 10000 draws from the posterior of σ^2 and compare with the theoretical $Inv - \chi^2(n, \tau^2) = Inv - \chi^2(10, 0.198)$. Figure 4 shows the histogram of the data simulated from $Inv - \chi^2(10, 0.198)$ and the density of the theoretical posterior distribution $Inv - \chi^2(10, 0.198)$. Since the simulated sample is so big, the histogram and the density is quite alike.




Figure 4: Histogram of simulated data and theoretical density of the σ^2 with draws= 10000.

- The most common measure of income inequality is the Gini coefficient, G , where $0 \leq G \leq 1$. $G=0$ means a completely equal income distribution, whereas $G=1$ means complete income inequality. It can be shown that $G = 2\Phi(\frac{\sigma}{\sqrt{2}}) - 1$ when incomes follow a $logNormal(\mu, \sigma^2)$ distribution. Use the draws in a) to compute the posterior of the Gini coefficient for the current data set. Figure 5 shows the histogram of G and the kernel density. It shows that G is far from 1 so the income distribution seems to be closer to equal than inequal.

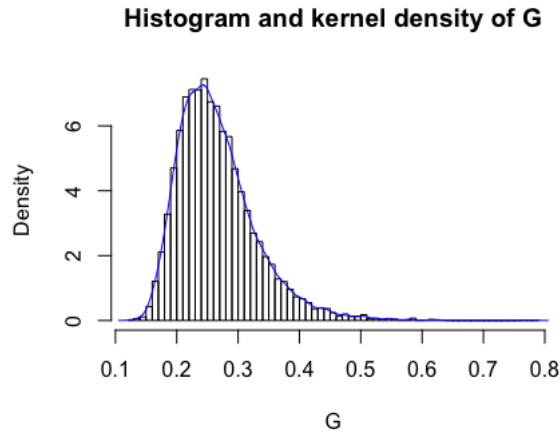


Figure 5: Histogram of simulated data and kernel density of G with draws= 10000.

- Use the posterior draws from b) to compute a 95% equal tail credible set for G . Also, do a kernel density estimate of the posterior of G and use it to compute a 95% Highest Posterior Density set for G . The Highest Posterior Density set (HPD) is a credible set that consists of

θ -values having the Highest Posterior Density. Below we have a 95% equal tail credible set for G and HPD. For a symmetric distribution, the intervals would be the same but since the distribution of G is slightly skewed, the intervals differ.

```
1 quantile(G, probs = c(0.025, 0.975))
2 #      2.5%      97.5%
3 # 0.1739323 0.4152031
4 hdi(density(G), credMass=0.95)
5 #lower      upper
6 #0.1601422 0.3918670
```

Question 3

The following data is the observed wind directions at a given location on 10 different days. The data recorded in radians:

$$(y_1, \dots, y_{10}) = (-2.44, 2.14, 2.54, 1.83, 2.02, 2.33, -2.79, 2.23, 2.07, 2.02)$$

where $-\pi \leq y \leq \pi$.

Assume that the observations are independent and follow the von Mises distribution

$$p(y|\mu, \kappa) = \frac{\exp(\kappa * \cos(y - \mu))}{2\pi I_0(\kappa)}$$

$-\pi \leq y \leq \pi$.

Furthermore, assume that $\mu = 2.39$ and let $\kappa \sim \text{Exp}(\lambda = 1)$ apriori.

- Plot the posterior distribution of κ for the data given above. The posterior is given by the following:

$$p(\kappa|\mu, y) = \exp(-\kappa) \prod_{i=1}^{10} \frac{\exp(\kappa * \cos(y_i - 2.39))}{2\pi I_0(\kappa)}$$

Figure 6 shows the posterior of κ .




Figure 6: Posterior of κ .

- The approximate point where the posterior has its mode is $(2.125, 4.7 * 10^{-6})$.

Contributions

All results and comments presented have been developed and discussed together by the members of the group.

Appendix

Question 1

```
1
2 install.packages("geoR")
3 install.packages("HDInterval")
4 library("geoR")
5 library("HDInterval")
6 #lab 1 Bayesian learning
7
8 #1a
9 x<-rbeta(10000,2+14,2+6)
10 h<-hist(x,breaks = seq(0,1,0.02))
11 xfit<-seq(min(x),max(x),length=100)
12 yfit<-dbeta(xfit,16,8)
13 yfit <- yfit*diff(h$mids[1:2])*length(x)
14 lines(xfit, yfit, col="blue", lwd=2)
15
16 #1b
17 pbeta(0.4,16,8)
18 y<-x[x<0.4]
19 prob<-length(y)/length(x)
20
21
22 #1c
23 xodds<-log(x/(1-x))
24 hist(xodds,breaks = 50,prob=TRUE, main = 'Histogram and density', xlab = expression(paste(log(
  theta/(1-theta))))
25 lines(density(xodds), col="blue", lwd=2)
26
27 #2a
28
29 data<-c(14,25,45,25,30,33,19,50,34,67)
30
31 tao2<-function(data)
32 {
33   sum((log(data)-3.5)^2)/length(data)
34 }
35
36 tao<-tao2(data=data)
37
38
39 sigma2<-rinvchisq(10000,df=10,scale=tao)
40 hist(sigma2,breaks = 100, prob=TRUE)
41 x<-seq(from=0, to=10000, by=0.001)
42 hx<-dinvchisq(x,df=10,scale=tao)
43
44 sigma2.histogram = hist(sigma2, breaks = 100, freq = F)
45 sigma2.ylim.normal = range(0, sigma2.histogram$density, dinvchisq(sigma2,df=10,scale=tao), na.
  rm = T)
46 hist(sigma2, breaks = 100, freq = F, ylim = c(0, 5.5), main=expression("Histogram and
  theoretical density of " ~ sigma^2), xlab = expression(paste(sigma^2)), ylab = 'Density')
47 curve(dinvchisq(x,df=10,scale=tao), add = T,col="blue")
48
49 #2.b
50 sigma<-sqrt(sigma2)/sqrt(2)
51 G<-2*pnorm(sigma,0,1)-1
52 hist(G,freq=F,breaks=50, main='Histogram and kernel density of G')
53 lines(density(G),col="blue")
54
55 #2.c
56 quantile(G, probs = c(0.025, 0.975))
57 # 2.5% 97.5%
58 # 0.1739323 0.4152031
59
60 hdi(density(G),credMass=0.95)
61
62 #lower upper
63 #0.1601422 0.3918670
64
65 #3.a
66
67 windradians<-c(-2.44,2.14,2.54,1.83,2.02,2.33,-2.79,2.23,2.07,2.02)
68 posterior<-function(k)
69 {
70   return(exp(-k)*prod(exp(k*cos(windradians-2.39))/(2*pi*besselI(x=k,nu=0))))
71 }
72
73
74
75 values<-sapply(k,posterior)
76 plot(k,values,type="l",main="Posterior of " ~ kappa, xlab=expression(paste(kappa)), ylab="
  Posterior")
77
```

```
78 #3.b
79 max(values)
80 #4.727694e-06
81 k[which.max(value)]
82 #2.125
```

732A91: Lab 2

Bayesian Learning

Sarah Alsaadi, Carles Sans Fuentes

May 22, 2017

Linear and polynomial regression

In this exercise we use the dataset TempLinkoping.txt in our analysis. The dataset contains daily temperatures (in Celsius degrees) at Malmsslatt and Linköping over the course of the year 2016, having as a response variable temp and covariate time.

First we determine the prior distribution of the model parameter. Given that our likelihood is a quadratic regression (see the summary of the quadratic regression model below) and that its conjugate prior is a multivariate normal distribution (because there is more than one Beta parameter), the prior hyperparameters are chosen as following:

1. μ_0 is the linear coefficients from our data (our best guess is just the computation of the parameters as if it was a quadratic regression of our data),
2. ω_0 a diagonal matrix of 1s given that all data has the same importance,
3. v_0 equal 6 in order to not to give too much importance to our prior and
4. σ_0^2 equal 16 (similar variance of the quadratic regression model)

```
1 > summary(lm)
2
3 Call:
4 lm(formula = temp ~ time + I(time^2), data = data)
5
6 Residuals:
7     Min       1Q   Median       3Q      Max
8 -10.0408  -2.6971  -0.1414   2.5157  12.2085
9
10 Coefficients:
11             Estimate Std. Error t value Pr(>|t|)
12 (Intercept) -10.6754     0.6475  -16.49  <2e-16 ***
13 time         93.5980     2.9822   31.39  <2e-16 ***
14 I(time^2)   -85.8311     2.8801  -29.80  <2e-16 ***
15 ---
16
17 Residual standard error: 4.107 on 363 degrees of freedom
18 Multiple R-squared:  0.7318, Adjusted R-squared:  0.7304
19 F-statistic: 495.3 on 2 and 363 DF, p-value: < 2.2e-16
```

In order to check whether our prior is sensible, we have simulated 1000 draws from the joint prior of all parameters and for every draw compute the regression curve. In order to check whether our prior look properly, we have evaluated the distributions of the Beta parameters and comparing it with the one from the quadratic regression. Below we have the mean of the beta coefficients.

```
1 apply(betasprior, 2, mean)
2     beta0     beta1     beta2
3  -9.991929  92.984489 -85.085894
```


It can be seen that given that our prior is really similar to our data, the Beta mean coefficients are almost the same.

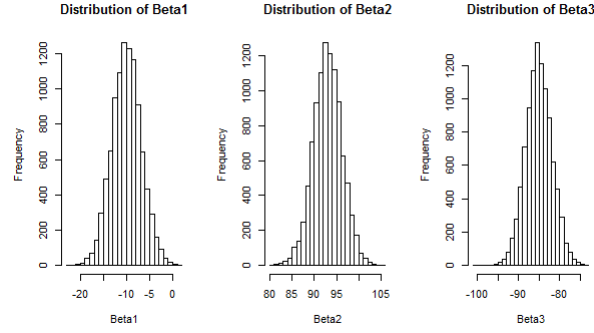


Figure 1: Histogram plot distribution of the our 3 Beta prior after 10000 simulations

It can be seen that the distribution looks normally distributed through the mean parameter of our data. This is good given the prior used. For this reason, we can say the curve looks reasonable.

Now, we have written a program that simulates from the joint posterior distribution of $\beta_0, \beta_1, \beta_2$ and σ_2 . We have produced a scatter plot of the temperature data and overlay a curve for the posterior mean of the regression function $f(time) = \beta_0 + \beta_1 * time + \beta_2 * time^2$ as well as the 95% equal tail posterior probability intervals for every value of time and then connect the lower and upper limits of the interval by curves (see figure 2)




Figure 2: Histogram of the data with the posterior mean distribution of the data (Blue line) and its 95% credible interval (red lines)

It can be seen that the variance of the data increases for the last half of the prediction.

In 1d we are asked to locate the time with the highest expected temperature (that is, the time where $E(temp—time)$ is maximal). We have used the previous betas to simulate the new data from the highest value of temp given time and we have plot the distribution of this points (see figure 3 below). It is seen that the mean of the maximum point is around 14-16. This makes sense since it is similar to the data we got.




Figure 3: Scatterplot of the data together with the computed temperature (blue curve) from the posterior mean of $\beta_0, \beta_1, \beta_2$ and the lower 2.5% and the upper 97.5% posterior credible interval

Finally, we are asked to estimate a polynomial model of order 7 and we are told that higher order terms may not be needed. Since we have a strong belief that higher order terms are not needed, we specify the prior parameters for those coefficients to reflect that. This would be a an expected value close to 0 and a very small variance.

Question 2

In activity 2 we are asked to write the model with the logistic function a data set to predict the probability that a woman will work given some variables that defines her.

For that, first we fit our data to the logistic model directly, and we get the following:

```

1 > summary(glmModel)
2
3 Call:
4 glm(formula = Work ~ 0 + ., family = binomial, data = Womenwork)
5
6 Deviance Residuals:
7     Min       1Q   Median       3Q      Max
8 -2.1662  -0.9299   0.4391   0.9494   2.0582
9
10 Coefficients:
11             Estimate Std. Error z value Pr(>|z|)
12 Constant      0.64430    1.52307   0.423 0.672274
13 HusbandInc   -0.01977    0.01590  -1.243 0.213752
14 EducYears     0.17988    0.07914   2.273 0.023024 *
15 ExpYears      0.16751    0.06600   2.538 0.011144 *
16 ExpYears2    -0.14436    0.23585  -0.612 0.540489
17 Age          -0.08234    0.02699  -3.050 0.002285 **
18 NSmallChild  -1.36250    0.38996  -3.494 0.000476 ***
19 NBigChild    -0.02543    0.14172  -0.179 0.857592
20 ---
21
22 (Dispersion parameter for binomial family taken to be 1)
23
24     Null deviance: 277.26  on 200  degrees of freedom
25 Residual deviance: 222.73  on 192  degrees of freedom
26 AIC: 238.73
27
28 Number of Fisher Scoring iterations: 4

```

Now we are asked to approximate the posterior distribution of the 8-dim parameter vector β with a multivariate normal distribution. For that, we have implemented a logistic function and used the `optim()` from R to find the Hessian and the best parameters for the best case given optimizing the posterior. The results is the following one:

```

1 > OptimResults
2 $par
3 [1] -0.020734822  0.198537870  0.170970267 -0.157370056 -0.073756571 -1.308417158
4 [7]  0.002373846
5
6 $value
7 [1] -134.015

```

```

8
9 $counts
10 function gradient
11      59      14
12
13 $convergence
14 [1] 0
15
16 $message
17 NULL
18
19 $hessian
20      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
21 [1,] -21691.9214 -10335.9656 -7724.06469 -1095.904557 -33606.6413 -190.535283
22 [2,] -10335.9656 -6069.9347 -4546.30708 -647.404832 -19675.4015 -127.202981
23 [3,] -7724.0647 -4546.3071 -5353.80184 -982.743357 -16236.0124 -75.200938
24 [4,] -1095.9046 -647.4048 -982.74336 -214.881014 -2470.3813 -7.716988
25 [5,] -33606.6413 -19675.4015 -16236.01237 -2470.381294 -68796.8982 -321.680446
26 [6,] -190.5353 -127.2030 -75.20094 -7.716988 -321.6804 -11.942247
27 [7,] -996.0942 -584.7677 -373.77395 -42.860366 -1894.1299 -12.925833
28      [,7]
29 [1,] -996.09418
30 [2,] -584.76769
31 [3,] -373.77395
32 [4,] -42.86037
33 [5,] -1894.12992
34 [6,] -12.92583
35 [7,] -123.33156

```

Also, the 95% CI for the variable NSmallChild must be found (see below) in order to define whether the variable is important or not

```

1 lowbound highbound
2 -2.028515 -0.588319

```

It can be seen that the credible interval is far from 0, which means that the variable is an important feature.

Finally, we are asked to Write a function that simulates from the predictive distribution of the response variable in a logistic regression. We have used our previously normal approximation to simulate and plot the predictive distribution for the Work variable for a 40 year old women, with two children (3 and 9 years old), 8 years of education, 10 years of experience. and a husband with an income of 10. The results is the following histogram:




Figure 4: Histogram of the predictive distribution of the response variable in a logistic regression

The result shows that it is quite much probable that the Woman will not be able to work.

Contributions

All results and comments presented have been developed and discussed together by the members of the group.

Appendix

Question 1

```
1
2 #lab 2 Bayesian learning
3
4 # install.packages("mvtnorm")
5 # install.packages("MASS")
6 library("mvtnorm")
7 library("geoR")
8 library("MASS")
9
10 data<-read.table("~/Google Drive/Kurser/Bayesian learning/Lab2/TempLinkoping2016.txt",head=TRUE
11 )
12 data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part2/TempLinkoping2016.txt", sep ="
13 \t")
14
15 #1.a
16 lm<-lm(temp~time+I(time^2),data)
17 summary(lm)
18
19 #1.b
20 mu_0<-c(-10,93,-85)
21 sigma_0<-16
22 v_0<-6
23 omega_0<-diag(3)
24 sigma_prior<-rinvchisq(1,df=v_0,scale=sigma_0)
25 beta_prior<-rmvnorm(n=1, mean =mu_0, sigma = sigma_prior*ginv(omega_0))
26 regress_prior<-beta_prior[1]+beta_prior[2]*data$time+beta_prior[3]*(data$time)^2
27 plot(y=data$temp,x=data$time)
28 lines(y=regress_prior,x=data$time,col="blue")
29
30 betasprior<-function(n)
31 {
32   beta0<-numeric(n)
33   beta1<-numeric(n)
34   beta2<-numeric(n)
35
36   for (i in 1:n)
37   {
38     beta0[i]<-as.vector(rmvnorm(n=1, mean =mu_0, sigma = sigma_prior*ginv(omega_0)))[1]
39     beta1[i]<-as.vector(rmvnorm(n=1, mean =mu_0, sigma = sigma_prior*ginv(omega_0)))[2]
40     beta2[i]<-as.vector(rmvnorm(n=1, mean =mu_0, sigma = sigma_prior*ginv(omega_0)))[3]
41   }
42   return(data.frame(beta0, beta1, beta2))
43 }
44 betasprior<-betasprior(10000)
45 apply(betasprior, 2, mean)
46 dim(betasprior)
47 par(mfrow=c(1,3))
48 hist(betasprior[,1], breaks = 30, xlab = "Beta1", main = "Distribution of Beta1")
49 hist(betasprior[,2], breaks = 30, xlab = "Beta2", main = "Distribution of Beta2")
50 hist(betasprior[,3], breaks = 30, xlab = "Beta3", main = "Distribution of Beta3")
51
52
53
54 #1.c
55 X<-cbind(rep(1,nrow(data)),data$time,(data$time)^2)
56 #beta_hat<-ginv(t(X)%*%X)%*%t(X)%*%data$temp
57 mu_n<-ginv(t(X)%*%X+omega_0)%*%t(X)%*%data$temp+omega_0%*%mu_0
58 Sigma_n<-t(X)%*%X+omega_0
59 v_n<-v_0+nrow(data)
60 sigma_n<-(1/v_n)*(v_0*sigma_0+(t(data$temp)%*%data$temp+t(mu_0)%*%omega_0%*%mu_0-t(mu_n)%*%
61   Sigma_n)%*%mu_n))
62
63 sigma_posterior<-as.vector(rinvchisq(1,df=v_n,scale=sigma_n))
64 beta_posterior<-rmvnorm(n=1, mean =mu_n, sigma = sigma_posterior*ginv(Sigma_n))
65 regress_posterior<-beta_posterior[1]+beta_posterior[2]*data$time+beta_posterior[3]*(data$time)
66 ^2
67 plot(y=data$temp,x=data$time)
68 lines(y=regress_posterior,x=data$time,col="blue")
69
70 betas<-function(n)
71 {
72   beta0<-numeric(n)
73   beta1<-numeric(n)
74   beta2<-numeric(n)
75
76   for (i in 1:n)
77   {
78     beta0[i]<-as.vector(rmvnorm(n=1, mean =mu_n, sigma = sigma_posterior*ginv(Sigma_n)))[1]
79     beta1[i]<-as.vector(rmvnorm(n=1, mean =mu_n, sigma = sigma_posterior*ginv(Sigma_n)))[2]
```

```

78     beta2[i]<-as.vector(rmvnorm(n=1, mean =mu_n, sigma = sigma_posterior*ginv(Sigma_n)))[3]
79   }
80 }
81 return(data.frame(beta0,beta1, beta2))
82 }
83
84
85 betas<-betas(1000)
86 betas0<-betas[,1]
87 betas1<-betas[,2]
88 betas2<-betas[,3]
89
90
91 temp_posterior<-function()
92 {
93   n=1000
94   k=366
95   temp_posterior= matrix(data=NA, nrow=n, ncol=k)
96   for(j in 1:k){
97     for(i in 1:n){
98       temp_posterior[i,j] = betas0[i]+betas1[i]*data$time[j]+betas2[i]*(data$time[j])^2
99     }
100   }
101   return(temp_posterior)
102 }
103 temp_posterior<-temp_posterior()
104
105 quantile_f<-function()
106 {
107   k=366
108   q_lower<-numeric(k)
109   q_upper<-numeric(k)
110
111   for (i in 1:k)
112   {
113     q_lower[i]=quantile(temp_posterior[,i], probs = c(0.025, 0.975))[1]
114     q_upper[i]=quantile(temp_posterior[,i], probs = c(0.025, 0.975))[2]
115   }
116   return(data.frame(q_lower,q_upper))
117 }
118 }
119
120 quantile_f<-quantile_f()
121 plot(q_lower,type="l")
122
123
124 par(mfrow= c(1,1))
125 plot(y=data$temp,x=data$time)
126 lines( y =quantile_f[,2],x=data$time, col= "red")
127 lines(y =quantile_f[,1],x=data$time, col= "red")
128 lines(y=regress_posterior,x=data$time,col="blue")
129
130
131 myfunc<- function (data){
132   regress_posterior<-betas[1]+betas[2]*data+betas[3]*(data)^2
133   return(regress_posterior)
134 }
135 }
136
137 temp_mean<-mean(beta_posterior[,1])+mean(beta_posterior[,2])*data$time+mean(beta_posterior[,3])
138 *data$time^2
139 t<-c(timeMax=data$time[which.max(temp_mean)], tempMax=max(temp_mean))
140
141 temp_posterior_max<-function()
142 {
143   n=1000
144
145   temp_posterior_max= numeric(n)
146
147   for(i in 1:n){
148     temp_posterior_max[i] = betas[i,1]+betas[i,2]*t[1]+betas[i,3]*t[1]^2
149   }
150
151   return(temp_posterior_max)
152 }
153
154 hist(temp_posterior_max())
155
156 ##1d
157
158 lm7<-lm(temp~time+I(time^2)+I(time^3)+I(time^4)+I(time^5)+I(time^6)+I(time^7),data)
159 mu7<-lm7$coefficients
160 sigma7<-16
161 v7<-365
162 #First values
163 lambdaplot<-function(lambda){

```

```

164 lambda <-0.5
165 Sigma7<-diag(8)*lambda
166 sigma_prior7<-rinvchisq(1,df=v7,scale=sigma7)
167 regress_prior7<- matrix(nrow = 1000, ncol = 366)
168 for(i in 1:1000){
169   beta_prior7<-rmvnorm(n=1, mean =mu7, sigma = sigma_prior7*ginv(Sigma7))
170   regress_prior7[i,]<-beta_prior7 [1]+beta_prior7 [2]*data$time+beta_prior7 [3]*(data$time)^2+beta
     _prior7 [4]*(data$time)^3+beta_prior7 [5]*(data$time)^4+beta_prior7 [6]*(data$time)^5+beta_
     prior7 [7]*(data$time)^6+beta_prior7 [8]*(data$time)^7
171 }
172 }
173 return(colMeans(regress_prior7))
174 }
175 lambda05<-lambdaplot(lambda =0.5)
176 lambda1<-lambdaplot(lambda =1)
177 lambda5<-lambdaplot(lambda =10)
178 lambda10<-lambdaplot(lambda =100)
179
180 #lambda<-seq(from=0,to=10,by=1)
181 par(mfrow = c(2,2))
182 plot(y=data$temp,x=data$time, main = "lambda = 0.5")
183 lines(y=lambda05,x=data$time,col="red")
184 plot(y=data$temp,x=data$time, main = "lambda = 1")
185 lines(y=lambda1,x=data$time,col="red")
186 plot(y=data$temp,x=data$time, main = "lambda = 10")
187 lines(y=lambda5,x=data$time,col="red")
188 plot(y=data$temp,x=data$time, main = "lambda = 100 ")
189 lines(y=lambda10,x=data$time,col="red")

```

Question 2

```

1
2 #lab 2 Bayesian learning
3
4 #####
5 ###2a
6 Womenwork<-read.table("~/Google Drive/Kurser/Bayesian learning/Lab2/WomenWork.dat.txt",head=
   TRUE)
7 Womenwork<- read.table("C:/Users/Carles/Desktop/Bayesian learning/Part2/WomenWork.dat.txt",head
   =TRUE)
8
9
10 glmModel <- glm(Work ~ 0 + ., data = Womenwork, family = binomial)
11
12 summary(glmModel)
13
14 #####b
15 #install.packages("mvtnorm") # Loading a package that contains the multivariate normal pdf
16 library("mvtnorm") # This command reads the mvtnorm package into R's memory. NOW we can use
   dmvtorm function.
17
18 # Loading data from file
19 #Data<-read.csv("/home/carsa564/Desktop/Bayesian learning/WomenWork.dat.txt", sep = "")
20 Data<- read.csv("C:/Users/M/Desktop/Statistics and Data Mining Master/Semester 2/Bayesian
   Learning/lab2/WomenWork.dat.txt", sep = "")
21 Data<-read.csv("~/Google Drive/Kurser/Bayesian learning/Lab2/WomenWork.dat.txt",sep = "")
22 Data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part2/WomenWork.dat.txt",sep = "")
23
24 tau <- 10; # Prior scaling factor such that Prior Covariance = (tau^2)*I
25 chooseCov <- c(2:8) # Here we choose which covariates to include in the model
26
27 y <- as.vector(Data[,1]); # Data from the read.table function is a data frame. Let's convert y
   and X to vector and matrix.
28 X <- as.matrix(Data[,2:ncol(Data)]);
29 covNames <- names(Data)[2:length(names(Data))];
30 X <- X[,chooseCov]; # Here we pick out the chosen covariates.
31 covNames <- covNames[chooseCov];
32 nPara <- dim(X)[2];
33
34 # Setting up the prior
35 mu <- as.vector(rep(0,nPara)) # Prior mean vector
36 Sigma <- tau^2*diag(nPara);
37
38
39
40
41 LogPostLogistic <- function(betaVect,y,X,mu,Sigma){
42
43   nPara <- length(betaVect);
44   linPred <- X%*%betaVect;
45
46   # evaluating the log-likelihood
47   logLik <- sum( linPred*y -log(1 + exp(linPred)));

```

```

48   if (abs(logLik) == Inf) logLik = -20000; # Likelihood is not finite, steer the optimizer away
      from here!
49
50   # evaluating the prior
51   logPrior <- dmvnorm(betaVect, matrix(0,nPara,1), Sigma, log=TRUE);
52
53   # add the log prior and log-likelihood together to get log posterior
54   return(logLik + logPrior)
55 }
56
57 initVal <- as.vector(rep(0,dim(X)[2]));
58
59 OptimResults<-optim(initVal,LogPostLogistic,gr=NULL,y,X,mu,Sigma,method=c("BFGS"),control=list(
      fnscale=-1),hessian=TRUE)
60
61 BetaCoef<- OptimResults$par
62 J<--solve(OptimResults$hessian)
63
64 myconf95<-c(lowbound =BetaCoef[6]-1.96*sqrt(J[6,6]), highbound=BetaCoef[6]+1.96*sqrt(J[6,6]))
65
66
67
68 #####2c
69
70 ##J as the matrix of covariates for sigma is calculated before as well as BetaCoef which is mu
      or the BetaCoefficients
71 Woman<-c(10, 8, 10, 100, 40, 2,0)
72
73
74
75
76 logProv<- function(n, betas= BetaCoef, covariances= J,obs=Woman){
77   prob<- integer(n)
78   Pr<- integer(n)
79   for(i in 1:n){
80     myrandombetas<- as.vector(rmvnorm(1, mean= BetaCoef, sigma = covariances))
81     Pr[i]<- exp((obs)%*%myrandombetas)/(1 + exp((obs)%*%myrandombetas))
82     prob[i]<- rbinom(1,1,Pr[i])
83   }
84
85   return(list("Predictive results"= prob, "Probability results"= Pr))
86 }
87
88
89 distribution_of_the_predictive_results<- logProv(10000, betas= BetaCoef, covariances= J,obs=
      Woman)
90
91 hist(distribution_of_the_predictive_results[[1]],
92     xlab = "bernoulli results on 0, 1",
93     ylab = "frequency results",
94     main = "Distribution of the predictive results for 10000 cases")

```


732A91: Lab 3

Bayesian Learning

Sarah Alsaadi, Carles Sans Fuentes

May 22, 2017

Normal model, mixture of normal model with semi-conjugate prior

The data rainfall.dat consists of daily records, from the beginning of 1948 to the end of 1983, of precipitation (rain or snow in units of 1 inch, and records of zero 100 precipitation are excluded) at Snoqualmie Falls, Washington. Analyze the data using the following two models.

1. Assume the daily precipitation $\{y_1, \dots, y_n\}$ are independent normally distributed, $y_1, \dots, y_n | \mu, \sigma^2 \sim N(\mu, \sigma^2)$ where both μ and σ^2 are unknown. Let $\mu \sim N(\mu_0, \tau_0^2)$ independently of $\sigma^2 \sim \text{Inv}\chi^2(\nu_0, \sigma_0^2)$

- (a) Implement (code!) a Gibbs sampler that simulates from the joint posterior $p(\mu, \sigma^2 | y_1, \dots, y_n)$. Where the full conditional posteriors are given by:

$$\text{i. } \mu | \sigma^2, y_1, \dots, y_n \sim N(\mu_n, \tau_n^2) \text{ where } \mu_n = \frac{n/\sigma^2}{n/\sigma^2 + 1/\tau_0^2} \bar{y} + \frac{1/\tau_0^2}{n/\sigma^2 + 1/\tau_0^2} \mu_0 \text{ and } \tau_n^2 = \frac{1}{n/\sigma^2 + 1/\tau_0^2}$$

$$\text{ii. } \sigma^2 | \mu, y_1, \dots, y_n \sim \text{Inv} - \chi^2(v_0 + n, \frac{v_0 \sigma_0^2 + \sum_{i=1}^n (y_i - \mu)^2}{n + v_0}).$$

The initial values have been set up near to 0 but τ , which has been chosen as 10. Proving with different τ values showed us that small τ , which accounts for variance, can lead our distribution of posterior estimates to stuck in local minimas since it is not able to get out from an optima. Still, only results with $\tau = 10$ will be shown, but it is interesting to know that with an smaller τ our optimal posterior μ distribution was lower (at around 26-28).

- (b) Here below in figure 1 it can be seen how the posterior of our μ and σ converges to 31-33 and 1550 respectively. There is a burning period for σ and μ to converge to the distribution.

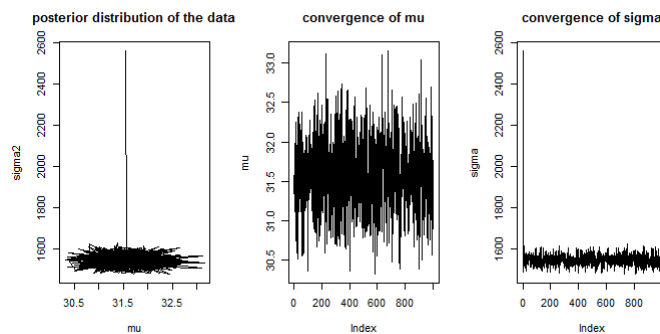


Figure 1: 3 plots of the posterior distribution of μ and σ with 1000 draws

In 2, we can see how the μ parameter is distributed with its cumulative mean that goes to 32 and how data is correlated. It can be seen that data is correlated just if we take each or every 2 observations. Thus, every third answer should be taken because it does not show any more autocorrelation or dependence.

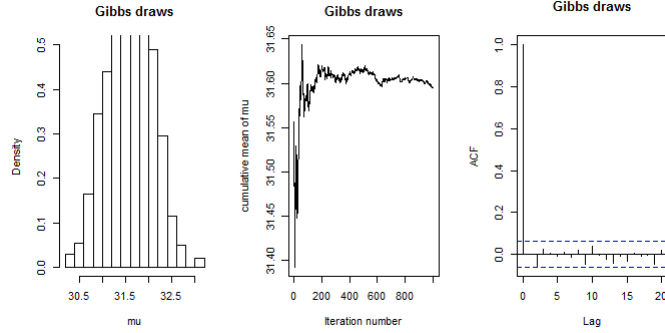


Figure 2: distribution of μ with 1000 draws, cumulative mean of the mean of μ and correlation lags

- (c) In 1b, we are asked to run a mixture of normal on the data for 2 μ and σ , assuming that data could come from two different points in Sweden.

In the following figure 3 below it can be seen the convergence of the mixture density model after 100 iterations.

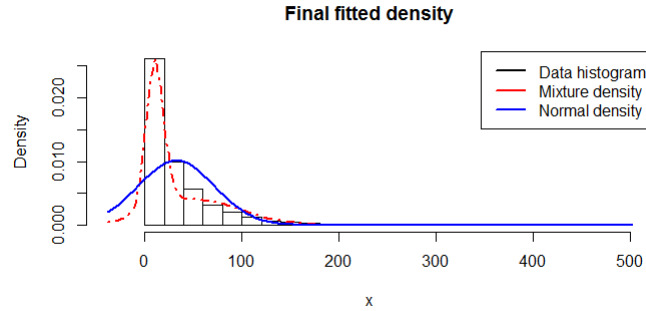


Figure 3: Convergence of the mixture density model after 100 iterations

Now (in 1c) we are asked to plot in one figure 4 to 1) a histogram or kernel density estimate of the data, 2) Normal density $N(\mu, \sigma^2)$ in (a) and 3) Mixture of normals in (b). This is done in the following graph being the red line the density 1000 random sample numbers from the distribution

$$N(\mu = \text{mean}(\mu_1), sd = \text{sqrt}(\text{mean}(\text{sigma2})))$$

and the yellow one a mixture from activity b also from a 1000 random generated data from 1000 points with from

$$\pi * N(\mu_1, \sigma_1^2) + (1 - \pi) * N(\mu_2, \sigma_2^2)$$

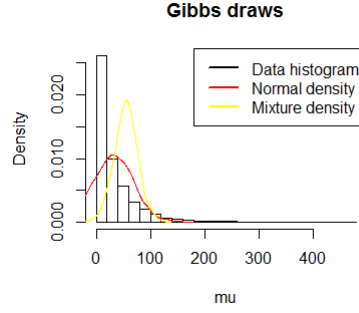


Figure 4: a histogram or kernel density estimate of the data, 2) Normal density $N(\mu, \sigma^2)$ in (a) and 3) Mixture of normals in(b)

Results show that none of both distributions is good for modeling the data. It is not a good idea just to take one or two normal distributions to model exponential data. More distributions would be necessary to get a better model.

Probit regression

1. Implement (code!) a data augmentation Gibbs sampler for the probit regression model

$$Pr(y = 1|\mathbf{x}) = \Phi(\mathbf{x}^T \beta)$$

The code is given in the appendix.

2. Compute the posterior of β in the probit regression for the WomenWork dataset from Lab 2 using the prior $\beta \sim N(0, \tau^2 I)$, with $\tau = 10$.
3. Do a normal approximation $\beta|\mathbf{y}, \mathbf{X} \sim \mathbf{N}(\tilde{\beta}, \mathbf{J}_{\mathbf{y}}^{-1}(\tilde{\beta}))$ of the posterior for β in the probit regression. Compare with the results from 2(b). Is the normal approximation accurate? In figure 5 and 6, you can see the distribution of our parameters. Also the the mean is provided with its 95% intervals below:

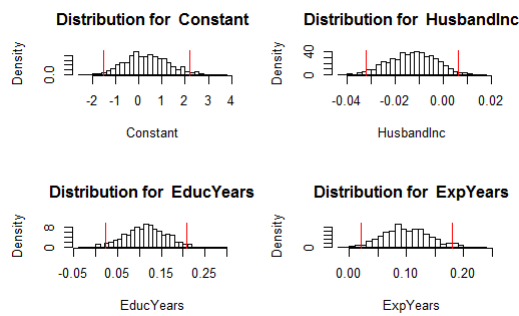


Figure 5: Distribution of the parameters with its 95% confidence interval (I)

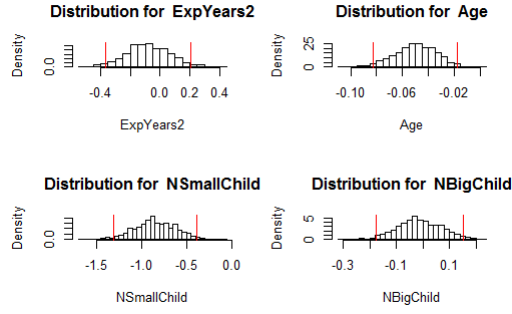


Figure 6: Distribution of the parameters with its 95% confidence interval (II)

```

1 > BetaFeat
2
3      Mean      upper      lower
4 Constant  0.34057957  2.180953926 -1.49979478
5 HusbandInc -0.01299479  0.006076275 -0.03206585
6 EducYears  0.11455707  0.207338016  0.02177613
7 ExpYears  0.10056627  0.179791545  0.02134099
8 ExpYears2 -0.07799728  0.207495176 -0.36348974
9 Age       -0.04989447 -0.017475928 -0.08231301
10 NSmallChild -0.84840269 -0.388974933 -1.30783046
11 NBigChild -0.01189881  0.152902525 -0.17670015

```

Now we have been asked to do a normal approximation which results can be seen here below. It can be seen that results are somehow similar, though not exact.

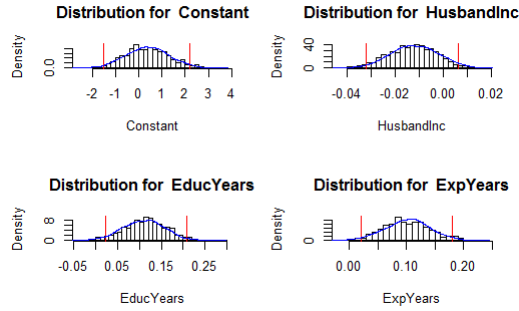


Figure 7: Distribution of the parameters with its 95% confidence interval and blue line for the normal approximation (I)

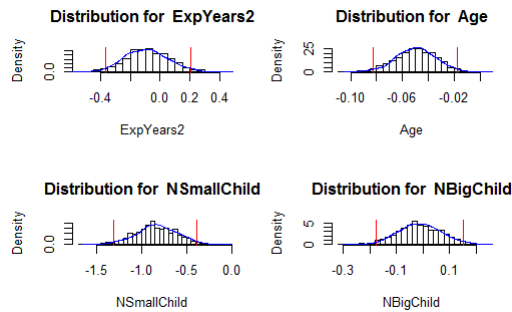


Figure 8: Distribution of the parameters with its 95% confidence interval and blue line for the normal approximation (II)

Also a table with the comparison of the different coefficients for both methods is reported.

```

1  > ConclTAbLe
2
3      GibbsMean      Gibbssd OptimalNBeta      Nsd
4 Constant      0.34057957 0.938966506 0.38451729 0.914146031
5 HusbandInc    -0.01299479 0.009730134 -0.01208712 0.009474897
6 EducYears     0.11455707 0.047337216 0.10840175 0.047218262
7 ExpYears      0.10056627 0.040421059 0.10221471 0.039459840
8 ExpYears2     -0.07799728 0.145659416 -0.09024608 0.140976385
9 Age           -0.04989447 0.016540073 -0.04962341 0.015890547
10 NSmallChild  -0.84840269 0.234401919 -0.81567500 0.226635007
    NBigChild    -0.01189881 0.084082316 -0.01241943 0.081506596

```

Contributions

All results and comments presented have been developed and discussed together by the members of the group.

Appendix

Question 1

```
1
2 library(geoR)
3 library(mvtnorm)
4
5 set.seed(12345)
6
7 data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part3/rainfall.dat.txt")
8
9
10 #####1a
11
12 ##Initial Data
13
14
15 nDraws<-1000
16 data<- unlist(data)
17 mu_0<- 0
18 v_0<- 1
19 sigma2_0<- 1
20 tau2_0<- 10
21
22
23 Normal_model<- function(nDraws, data,mu_0, v_0, tau2_0, sigma2_0){
24   require(geoR)
25   ##initializing data result
26   gibbsDraws <- matrix(0,nDraws,2)
27   colnames(gibbsDraws)<-c("mu", "sigma2")
28
29   #basic variables
30   n<- length(data)
31   v_n<- n +v_0
32
33   #initializing basic variables for the gibb sampling
34   mu<-mu_0
35
36   for(i in 1:nDraws){
37
38     sigma2 <- rinvchisq(1, df= v_n, scale = (v_0*sigma2_0+ sum((data-mu)**2))/(n +v_0))
39     gibbsDraws[i,2] <- sigma2
40
41     w<- (n/sigma2)/(n/sigma2+1/tau2_0)
42     mu_n<- w*mean(data)+(1-w)*mu_0
43     invtau2<- (1/(n/sigma2+1/tau2_0))
44
45     mu <-rnorm(1,mu_n, sd = sqrt(invtau2))
46     gibbsDraws[i,1]<- mu
47
48   }
49
50   return(gibbsDraws)
51
52 }
53
54
55 gibbsDraws<-Normal_model(nDraws=nDraws, data= data ,mu_0= mu_0, v_0= v_0, tau2_0= tau2_0,
56   sigma2_0= sigma2_0)
57 tail(gibbsDraws)
58 par(mfrow = c(1,3))
59 plot(gibbsDraws, type = "l", main = "posterior distribution of the data")
60 plot(gibbsDraws[,1], type = "l", ylab = "mu", main = "convergence of mu")
61 plot(gibbsDraws[,2], type = "l", ylab = "sigma", main = "convergence of sigma")
62
63
64
65 hist(gibbsDraws[,1], freq = FALSE, main='Gibbs draws', ylim = c(0,0.5), xlab= "mu")
66 lines(seq(-2,4,by=0.01),dnorm(seq(-2,4,by=0.01), mean = 1), col = "red", lwd = 3)
67 plot(cumsum(gibbsDraws[,1])/seq(1,nDraws),type="l", main='Gibbs draws', xlab='Iteration number',
68   , ylab='cumulative mean of mu')
69 lines(seq(1,nDraws),1*matrix(1,1,nDraws),col="red",lwd=3)
70 acf(gibbsDraws[,1], main='Gibbs draws', lag.max = 20)
71 par(mfrow = c(1,1))
72
73 #####1b
74 ##### BEGIN USER INPUT #####
75 # Data options
76 rawData <- data
77 x <- as.matrix(data)
78
79 # Model options
```

```

80 nComp <- 2      # Number of mixture components
81
82 # Prior options
83 alpha <- 10*rep(1,nComp) # Dirichlet(alpha)
84 muPrior <- rep(0,nComp) # Prior mean of theta
85 tau2Prior <- rep(10,nComp) # Prior std theta
86 sigma2_0 <- rep(var(x),nComp) # s20 (best guess of sigma2)
87 nu0 <- rep(4,nComp) # degrees of freedom for prior on sigma2
88
89 # MCMC options
90 nIter <- 100 # Number of Gibbs sampling draws
91
92 # Plotting options
93 plotFit <- TRUE
94 lineColors <- c("blue", "green", "magenta", 'yellow')
95 sleepTime <- 0.1 # Adding sleep time between iterations for plotting
96 ##### END USER INPUT #####
97
98 ##### Defining a function that simulates from the
99 rScaledInvChi2 <- function(n, df, scale){
100   return((df*scale)/rchisq(n,df=df))
101 }
102
103 ##### Defining a function that simulates from a Dirichlet distribution
104 rDirichlet <- function(param){
105   nCat <- length(param)
106   thetaDraws <- matrix(NA,nCat,1)
107   for (j in 1:nCat){
108     thetaDraws[j] <- rgamma(1,param[j],1)
109   }
110   thetaDraws = thetaDraws/sum(thetaDraws) # Dividing every column of ThetaDraws by the sum of the
111     elements in that column.
112   return(thetaDraws)
113 }
114
115 # Simple function that converts between two different representations of the mixture allocation
116 S2alloc <- function(S){
117   n <- dim(S)[1]
118   alloc <- rep(0,n)
119   for (i in 1:n){
120     alloc[i] <- which(S[i,] == 1)
121   }
122   return(alloc)
123 }
124
125 # Initial value for the MCMC
126 nObs <- length(x)
127 S <- t(rmultinom(nObs, size = 1, prob = rep(1/nComp,nComp))) # nObs-by-nComp matrix with
128   component allocations.
129 theta <- quantile(x, probs = seq(0,1,length = nComp))
130 sigma2 <- rep(var(x),nComp)
131 probObsInComp <- rep(NA, nComp)
132
133 # Setting up the plot
134 xGrid <- seq(min(x)-1*apply(x,2,sd),max(x)+1*apply(x,2,sd),length = 100)
135 xGridMin <- min(xGrid)
136 xGridMax <- max(xGrid)
137 mixDensMean <- rep(0,length(xGrid))
138 effIterCount <- 0
139 ylim <- c(0,2*max(hist(x)$density))
140
141 ##Recording mu and sigma2
142 matmu<- matrix(0, nrow = nIter, ncol = nComp)
143 colnames(matmu)<- c("mu1", "mu2")
144
145 matsigma2<- matrix(0, nrow = nIter, ncol = nComp)
146 colnames(matsigma2)<- c("sigma1", "sigma2")
147
148 matpi<- matrix(0, nrow = nIter, ncol = nComp)
149 colnames(matsigma2)<- c("pi1", "pi2")
150
151
152 for (k in 1:nIter){
153   message(paste('Iteration number:',k))
154   alloc <- S2alloc(S) # Just a function that converts between different representations of the
155     group allocations
156   nAlloc <- colSums(S)
157   print(nAlloc)
158   # Update components probabilities
159   w <- rDirichlet(alpha + nAlloc)
160   matpi[k,]<- w
161
162   # Update theta's
163   for (j in 1:nComp){

```



```

164     precPrior <- 1/tau2Prior[j]
165     precData <- nAlloc[j]/sigma2[j]
166     precPost <- precPrior + precData
167     wPrior <- precPrior/precPost
168     muPost <- wPrior*muPrior + (1-wPrior)*mean(x[alloc == j])
169
170     tau2Post <- 1/precPost
171     theta[j] <- rnorm(1, mean = muPost, sd = sqrt(tau2Post))
172 }
173 matmu[k,]<- muPost
174
175
176 # Update sigma2's
177 for (j in 1:nComp){
178     sigma2[j] <- rScaledInvChi2(1, df = nu0[j] + nAlloc[j], scale = (nu0[j]*sigma2_0[j] + sum((
179         x[alloc == j] - theta[j])^2))/(nu0[j] + nAlloc[j]))
180 }
181 matsigma2[k,]<- sigma2
182
183 # Update allocation
184 for (i in 1:nObs){
185     for (j in 1:nComp){
186         probObsInComp[j] <- w[j]*dnorm(x[i], mean = theta[j], sd = sqrt(sigma2[j]))
187     }
188     S[i,] <- t(rmultinom(1, size = 1, prob = probObsInComp/sum(probObsInComp)))
189 }
190
191 # Printing the fitted density against data histogram
192 if (plotFit && (k%1 == 0)){
193     effIterCount <- effIterCount + 1
194     hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = paste("Iteration
195         number",k), ylim = ylim)
196     mixDens <- rep(0,length(xGrid))
197     components <- c()
198     for (j in 1:nComp){
199         compDens <- dnorm(xGrid,theta[j],sd = sqrt(sigma2[j]))
200         mixDens <- mixDens + w[j]*compDens
201         lines(xGrid, compDens, type = "l", lwd = 2, col = lineColors[j])
202         components[j] <- paste("Component ",j)
203     }
204     mixDensMean <- ((effIterCount-1)*mixDensMean + mixDens)/effIterCount
205     lines(xGrid, mixDens, type = "l", lty = 2, lwd = 3, col = 'red')
206     legend("topleft", box.lty = 1, legend = c("Data histogram",components, 'Mixture'),
207         col = c("black",lineColors[1:nComp], 'red'), lwd = 2)
208     Sys.sleep(sleepTime)
209 }
210
211 }
212
213 hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = "Final fitted density")
214 lines(xGrid, mixDensMean, type = "l", lwd = 2, lty = 4, col = "red")
215 lines(xGrid, dnorm(xGrid, mean = mean(x), sd = apply(x,2,sd)), type = "l", lwd = 2, col = "blue
216 ")
217 legend("topright", box.lty = 1, legend = c("Data histogram","Mixture density","Normal density")
218 , col=c("black","red","blue"), lwd = 2)
219
220 #####C Graphical representation
221 #Data sets for the distribution of mu1, mu2, sigma2_1, sigma_2, pi1, pi2
222 matmu
223 matpi
224 matsigma2
225 #Data set for the distribution of mu1, sigma2_1 in the first case
226
227
228 sampled_simple<-rnorm(1000, mean = mean(gibbsDraws[,1]), sd = sqrt(mean(gibbsDraws[,2])))
229 mixture_sample<- mean(matpi[,1])*rnorm(1000, mean = mean(matmu[,1]), sd = sqrt(mean(matsigma2
230     [,1])))+ (1-mean(matpi[,1]))*rnorm(1000, mean = mean(matmu[,2]), sd = sqrt(mean(matsigma2
231     [,2])))
232
233 hist(data, freq = FALSE, main='Gibbs draws', xlab= "mu", breaks = 30)
234 lines(density(sampled_simple), col = "red")
235 lines(density(mixture_sample), col = "yellow")
236 legend("topright", box.lty = 1, legend = c("Data histogram","Normal density", "Mixture density"
237     ), col=c("black","red","yellow"), lwd = 2)

```

Question 2

```

1
2 #####Question 2

```

```

3 #install.packages("msm")
4 library(msm)
5
6 Data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part2/WomenWork.dat.txt", sep = "",
7   header = TRUE)
8 glmModel <- glm(Work ~ 0 + ., data = Data, family = binomial(link = "probit"))
9 summary(glmModel)
10
11 ##### Variables
12 tau <- 10; # Prior scaling factor such that Prior Covariance = (tau^2)*I
13 chooseCov <- c(1:8) # Here we choose which covariates to include in the model
14
15 y <- as.vector(Data[,1]); # Data from the read.table function is a data frame. Let's convert y
16   and X to vector and matrix.
17 X <- as.matrix(Data[,2:ncol(Data)]);
18 initVal <- as.vector(rep(0,dim(X)[2]));
19 covNames <- names(Data)[2:length(names(Data))];
20 X <- X[,chooseCov]; # Here we pick out the chosen covariates.
21 covNames <- covNames[chooseCov];
22 nPara <- dim(X)[2];
23
24
25 # Setting up the prior
26 mu_0 <- as.vector(rep(0,nPara)) # Prior mean vector
27 Omega_0 <- tau^2*diag(nPara);
28 beta_0 <- as.vector(rmvnorm(1, mean = mu_0, sigma =Omega_0))
29 nIter <- 1000
30
31
32
33 ProbitFunction<- function(beta_0, mu_0, Omega_0, X, y, n_iter){
34   posy_0 <- which(y == 0)
35   posy_1 <- which(y == 1)
36   nPara <- dim(X)[2]
37
38   beta<-matrix(0, ncol = nPara, nrow = n_iter)
39   beta[1,]<- beta_0
40   B<- solve(solve(Omega_0)+t(X)%*%X)
41
42   for(i in 1:n_iter){
43     beta[i, ]<- rmvnorm(1, as.vector(B%*(solve(Omega_0)%*%mu_0+t(X)%*%y)), sigma = B)
44
45     y[posy_1] <- rtorm(length(posy_1), mean = as.vector(X[posy_1,]%*%beta[i,]), sd = 1,lower =
46       0, upper = Inf)
47     y[posy_0] <- rtorm(length(posy_0), mean = as.vector(X[posy_0,]%*%beta[i,]), sd = 1,lower =
48       -Inf, upper = 0)
49
50   }
51   return(beta)
52 }
53
54
55
56 Betadist<-ProbitFunction(beta_0=beta_0, mu_0= mu_0, Omega_0= Omega_0, X= X, y= y , n_iter =
57   nIter)
58 colnames(Betadist)<- covNames
59 dim(Betadist)
60 BetaFeat<- data.frame(Mean = apply(Betadist, 2, mean),
61   upper = apply(Betadist, 2, mean)+1.96*apply(Betadist, 2, sd),
62   lower = apply(Betadist, 2, mean)-1.96*apply(Betadist, 2, sd))
63 par(mfrow = c(2,2))
64 for(i in 1: 4){
65   hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i],
66     breaks = 30, freq = FALSE)
67   abline(v = BetaFeat$lower[i], b = 0, col = "red")
68   abline(v = BetaFeat$upper[i], b = 0, col = "red")
69 }
70 par(mfrow = c(2,2))
71 for(i in 5:8){
72   hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i], breaks
73     = 30, freq = FALSE)
74   abline(v = BetaFeat$lower[i], b = 0, col = "red")
75   abline(v = BetaFeat$upper[i], b = 0, col = "red")
76 }
77
78 ###C
79
80 LogPostProbit <- function(betaVect,y,X,mu,Sigma){
81   nPara <- length(betaVect);
82   linPred <- X%*%betaVect;

```

```

83
84 # The following is a more numerically stable evaluation of the log-likelihood in my slides:
85 # logLik <- sum(y*log(pnorm(linPred)) + (1-y)*log(1-pnorm(linPred)))
86 logLik <- sum(y*pnorm(linPred, log.p = TRUE) + (1-y)*pnorm(linPred, log.p = TRUE, lower.tail
    = FALSE))
87
88 # evaluating the prior
89 logPrior <- dmvnorm(betaVect, matrix(0,nPara,1), Sigma, log=TRUE);
90
91 # add the log prior and log-likelihood together to get log posterior
92 return(logLik + logPrior)
93
94 }
95
96 OptimResults<-optim(initVal,LogPostProbit,gr=NULL,y,X,mu= mu_0,Sigma= Omega_0,method=c("BFGS"),
    control=list(fnscale=-1),hessian=TRUE)
97
98 BetaCoef<- OptimResults$par
99 names(BetaCoef)<- covNames
100 J<--solve(OptimResults$hessian)
101
102 myconf95<-c(lowbound =BetaCoef[6]-1.96*sqrt(J[6,6]), highbound=BetaCoef[6]+1.96*sqrt(J[6,6]))
103
104 sampleBetaProbit<- matrix(ncol = length(BetaCoef), nrow = 1000)
105 for(i in 1:length(BetaCoef)){
106 sampleBetaProbit[,i]<- rnorm(1000, mean = BetaCoef[i], sd = sqrt(J[i,i]))
107 }
108
109 par(mfrow = c(2,2))
110 for(i in 1: 4){
111 hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i], breaks
    = 30, freq = FALSE)
112 abline(v = BetaFeat$lower[i], b = 0, col = "red")
113 abline(v = BetaFeat$upper[i], b = 0, col = "red")
114 lines(density(sampleBetaProbit[,i]), col = "blue")
115 }
116 for(i in 5:8){
117 hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i], breaks
    = 30, freq = FALSE)
118 abline(v = BetaFeat$lower[i], b = 0, col = "red")
119 abline(v = BetaFeat$upper[i], b = 0, col = "red")
120 lines(density(sampleBetaProbit[,i]), col = "blue")
121
122 }
123 }
124
125 ConclTable<- data.frame(GibbsMean = apply(Betadist, 2, mean),
126 Gibbsstd = apply(Betadist, 2, sd),
127 OptimalBeta = BetaCoef,
128 Gibbsstd=sqrt(diag(J)))

```

732A91: Lab 4

Bayesian Learning

Sarah Alsaadi, Carles Sans Fuentes

May 24, 2017

Poisson regression-the MCMC way

Consider the following Poisson regression model

$$y_i|\beta \sim \text{Poisson}[\exp(\mathbf{x}_i^T \beta)], i = 1, \dots, n,$$

where y_i is the count for the i th observation in the sample and \mathbf{x}_i is the p -dimensional vector with covariate observations for the i th observation. The data set **eBayNumberOfBidderData.dat** contains observations from 1000 eBay auctions of coins. The response variable is **nBids** and records the number of bids in each auction. The remaining variables are features/covariates (\mathbf{x}):

- **Const** (for the intercept)
- **PowerSeller** (is the seller selling large volumes on eBay)
- **VerifyID** (is the seller verified by eBay?)
- **Sealed** (was the coin sold sealed in never opened envelope?)
- **MinBlem** (did the coin have a minor defect?)
- **MajBlem** (a major defect?)
- **LargNeg** (did the seller get a lot of negative feedback from customers?)
- **LogBook** (logarithm of the coins book value according to expert sellers. Standardized)
- **MinBidShare** (a variable that measures ratio of the minimum selling price (starting price) to the book value. Standardized).

(a) Using **glm** in R, we obtain the following results:

```
1 Call:
2 glm(formula = nBids ~ 0 + ., family = poisson, data = ebay)
3
4 Deviance Residuals:
5     Min       1Q   Median       3Q      Max
6  -3.5800  -0.7222  -0.0441   0.5269   2.4605
7
8 Coefficients:
9             Estimate Std. Error z value Pr(>|z|)
10 Const          1.07244    0.03077  34.848 < 2e-16 ***
11 PowerSeller    -0.02054    0.03678  -0.558  0.5765
12 VerifyID      -0.39452    0.09243  -4.268 1.97e-05 ***
13 Sealed         0.44384    0.05056   8.778 < 2e-16 ***
14 Minblem       -0.05220    0.06020  -0.867  0.3859
15 Majblem       -0.22087    0.09144  -2.416  0.0157 *
16 LargNeg        0.07067    0.05633   1.255  0.2096
17 LogBook       -0.12068    0.02896  -4.166 3.09e-05 ***
18 MinBidShare   -1.89410    0.07124 -26.588 < 2e-16 ***
19 ---
20 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
21
22 (Dispersion parameter for poisson family taken to be 1)
23
24 Null deviance: 6264.01 on 1000 degrees of freedom
```

25 Residual deviance: 867.47 on 991 degrees of freedom
 26 AIC: 3610.3
 27
 28 Number of Fisher Scoring iterations: 5

The output shows that the significant MLE of the β :s are those for the covariates Const, VerifyID, Sealed, MayBlem, LogBook and MinBidShare.

- (b) In this exercise we did a Bayesian analysis of the Poisson regression. We let $\beta \sim N[0, 100(\mathbf{X}^T \mathbf{X})^{-1}]$ apriori, where \mathbf{X} is the $n \times p$ covariate matrix. Next we assumed that the posterior density is approximately multivariate normal:

$$\beta|\mathbf{y}, \mathbf{X} \sim N(\tilde{\beta}, J_{\mathbf{y}}^{-1}(\tilde{\beta}))$$

where $\tilde{\beta}$ is the posterior mode and $J_{\mathbf{y}}(\tilde{\beta}) = -\frac{\partial^2 \ln p(\beta|\mathbf{y})}{\partial \beta \partial \beta^T} \big|_{\beta=\tilde{\beta}}$ is the observed Hessian evaluated at the posterior mode. To obtain $\tilde{\beta}$ and $J_{\mathbf{y}}^{-1}(\tilde{\beta})$ we used numerical optimization (**optim.R**). The results of the estimated β :s are given in the table below. We also made histograms, one for each variable, of 10000 draws of β , the marginal distribution looks normal, see Figure 1.

Variable	$\tilde{\beta}$
Const	1.06984118
PowerSeller	-0.02051246
VerifyID	-0.39300599
Sealed	0.44355549
MinBlem	-0.05246627
MajBlem	0.22123840
LargNeg	0.07069683
LogBook	-0.12021767
MinBidShare	-1.89198501

Table 1: The obtained β :s through maximization of the posterior distribution, one for each variable.




Figure 1: Histogram of the 10000 β :s drawn from the marginal posterior distribution

- (c) In this exercise we simulate from the actual posterior of β using the Metropolis algorithm and compare with the approximate results in b). We program a general function that uses

the Metropolis algorithm to generate random draws from an arbitrary posterior density. We use the multivariate normal density as proposal density:

$$\theta_p|\theta_c \sim N(\theta_c, \tilde{c} \cdot \Sigma)$$

where $\Sigma = J_{\mathbf{y}}^{-1}(\tilde{\beta})$ obtained in the previous exercise and θ_c is the current draw (hence the subscript c). The value \tilde{c} is a tuning parameter and is an input to our Metropolis function (so that a user can change it). The user of our Metropolis function is able to supply her own posterior density function, not necessarily for the Poisson regression, and still be able to use our Metropolis function.

First, one of the input arguments of our Metropolis function is called logPostFunc. logPostFunc is a function object that computes the log posterior density at any value of the parameter vector. This is needed when we compute the acceptance probability of the Metropolis algorithm. We program the log posterior density, since logs are more stable and avoids problems with too small or large numbers (overflow). Note that the ratio of posterior densities in the Metropolis acceptance probability can be written

$$\frac{p(\theta_p|\mathbf{y})}{p(\theta_c|\mathbf{y})} = \exp[\log(p(\theta_p|\mathbf{y})) - \log(p(\theta_c|\mathbf{y}))]$$

This is smart since the large or small common factors in $p(\theta_p|\mathbf{y})$ and $p(\theta_c|\mathbf{y})$ cancel out before we evaluate the exponential function (which can otherwise overflow).

Second, the first argument of our (log) posterior function is theta, the (vector) of parameters for which the posterior density is evaluated.

Third, the user's posterior density is also a function of the data and prior hyperparameters and those are supplied to the Metropolis function where we use the triple dot (...) argument which is like a wildcard for any parameters supplied by the user. This makes it possible to use the Metropolis function for any problem, even when a programmer don't know what the user's posterior density function looks like or what kind of data and hyperparameters being used in that particular problem.

Now, we use Metropolis function to sample from the posterior of β in the Poisson regression for the eBay dataset. We assess MCMC convergence by graphical methods. The parameters $\phi_j = \exp(\beta_j)$ are usually considered more interpretable than the β_j . We compute the posterior distribution of ϕ_j for all variables.

Figure 2 shows that the Metropolis algorithm seem to converge, with some burn in period for all the β :s.

In Table 2 we have the mean posterior of the simulated β :s which are very similar to the ones obtained through maximization.

Figure 3 shows the marginal posterior for each e^β obtained by the Metropolis algorit. The distribution looks normal.




Figure 2: Plot of the 10000 β :s drawn from the marginal posterior distribution obtained by the Metropolis algorithm.

Variable	$\tilde{\beta}$
Const	1.06538266
PowerSeller	-0.02529706
VerifyID	-0.38552635
Sealed	0.43553156
MinBlem	-0.05410433
MajBlem	-0.22019352
LargNeg	0.06493763
LogBook	-0.12021880
MinBidShare	-1.88132053

Table 2: The posterior mean of the simulated β :s, simulated using the Metropolis algorithm, the values are almost identical as the ones obtained through maximization

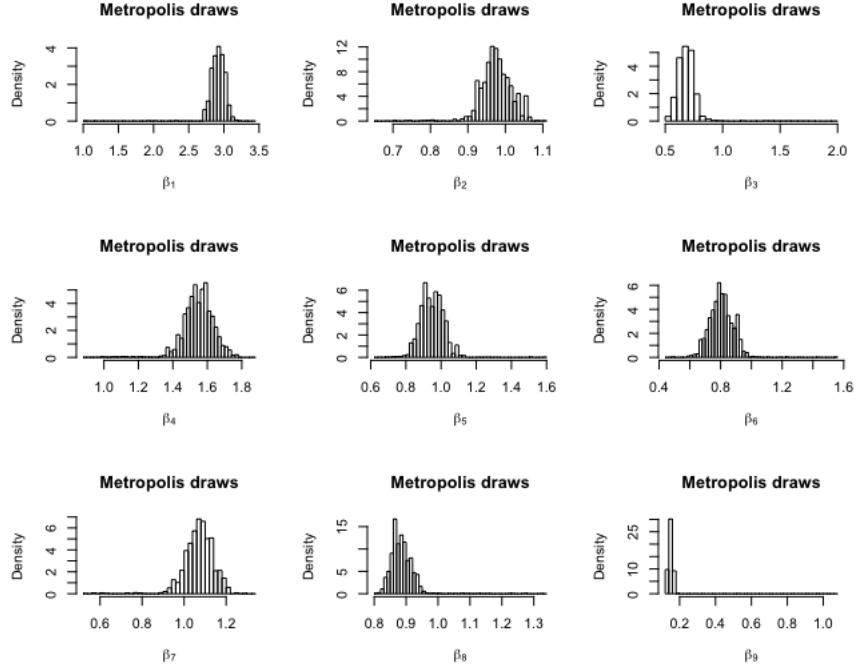


Figure 3: Histogram of the 10000 e^β drawn from the marginal posterior distribution obtained by the Metropolis algorithm.

(d) We use the MCMC draws from c) to simulate from the predictive distribution of the number of bidders in a new auction with the characteristics below. The histogram of the predictive distribution is shown in Figure 4. The probability of no bidders in this new auction is 0.3512.

- **PowerSeller= 1**
- **VerifyID= 1**
- **Sealed= 1**
- **MinBlem= 0**
- **MajBlem= 0**
- **LargNeg= 0**
- **LogBook= 1**
- **MinBidShare= 0.5**

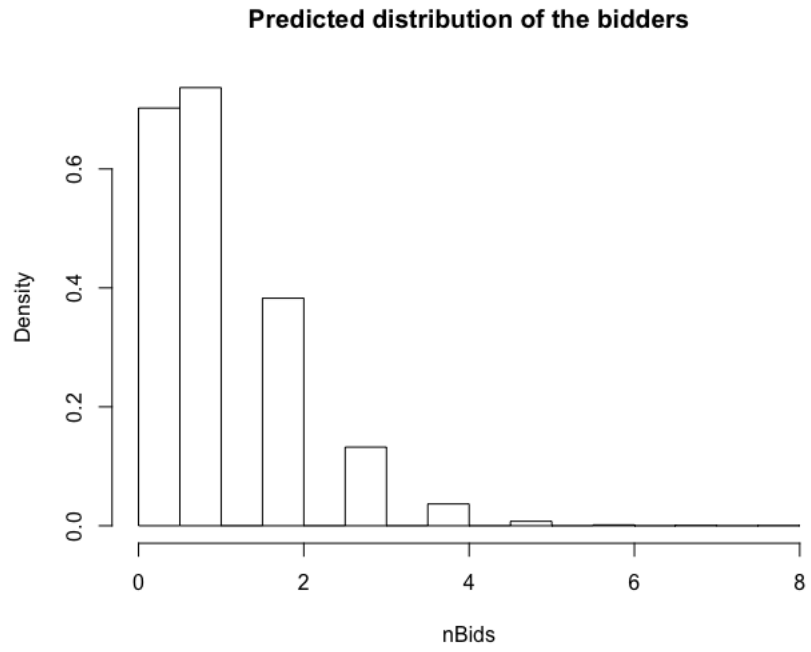


Figure 4: Histogram of 10000 simulated predicted values of the variable nBid with given characteristics

Contributions

All results and comments presented have been developed and discussed together by the members of the group.

Appendix

Poisson regression-the MCMC way

```
1
2 ### LAB 4
3
4 #1
5 Data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part4/eBayNumberOfBidderData.dat.txt", sep = "")
6
7 ##a
8 #### Variables
9
10 PoisModel<-glm(nBids~.-Const, family="poisson", data = Data)
11 logLik(PoisModel)
12
13
14 mysum<-summary(PoisModel)[["coefficients"]]
15
16
17
18 ##Significants at a 95% level
19 significant <- mysum[which(mysum[,4]<0.05),]
20 non_significant <- mysum[which(mysum[,4]>=0.05),]
21 list(significant = significant, nonsignificant = non_significant)
22
23
24 ##b
25 chooseCov <- 1:(length(names(Data))-1); # Here we choose which covariates to include in the model
26
27
28 y <- as.vector(Data[,1]); # Data from the read.table function is a data frame. Let's convert y and X to vector and matrix.
29 X <- as.matrix(Data[,2:length(names(Data))]);
30 covNames <- names(Data)[2:length(names(Data))];
31 X <- X[,chooseCov]; # Here we pick out the chosen covariates.
32 covNames <- covNames[chooseCov];
33 nPara <- dim(X)[2]
34
35
36 ##prior Beta
37 library(geoR)
38 library(mvtnorm)
39
40
41 mu_0 <- matrix(0, nPara,1)
42 Sigma_0 <- 100*solve(crossprod(X,X))
43 Beta_0 <- rmvnorm(1, mean = mu_0, sigma = Sigma_0)
44 InitVal <- matrix(0, ncol=nPara, nrow =1)
45
46
47 LogPostPois <- function(betaVect,y = y,X = X, mu =mu_0,Sigma= Sigma_0){
48   nPara <- length(betaVect);
49   linPred <- X%*%betaVect;
50
51   # The following is a more numerically stable evaluation of the log-likelihood in my slides:
52   # logLik <- sum(y*log(pnorm(linPred)) + (1-y)*log(1-pnorm(linPred)))
53   logLik <- sum(linPred*y-exp(linPred))
54
55   # evaluating the prior
56
57   logPrior <- dmvnorm(betaVect, mu, Sigma, log=TRUE);
58
59   # add the log prior and log-likelihood together to get log posterior
60   return(logLik + logPrior)
61 }
62
63
64 OptimResults<-optim(InitVal,LogPostPois,gr=NULL,y = y,X = X, mu =mu_0,Sigma= Sigma_0 ,method=c("BFGS"),control=list(fnscale=-1),hessian=TRUE)
65
66 BetaCoef<- OptimResults$par
67 colnames(BetaCoef)<- covNames
68 J<--solve(OptimResults$hessian)
69
70 mycoefvar<- data.frame(Coefficients= as.vector(BetaCoef), variance = diag(J))
71 rownames(mycoefvar)<- covNames
72 mycoefvar
73 #####C
74
75 set.seed(12345)
76
77 LogPostPoisson <- function(theta, priormu, priorsigma, X, Y, ...) {
```

```

78   require(mvtnorm)
79   likelihood <- dpois(Y, lambda = as.vector(exp((X) %*% t(theta))), log = TRUE)
80   prior <- dmvnorm(theta, mean = priormu, sigma = priorsigma, log=TRUE)
81   return(sum(likelihood) + prior)
82 }
83
84
85 metrop1<-function(logPostFunc, theta_0, constant, sigma, nIter,...){
86   #initialize chain
87   require(mvtnorm)
88   theta<- matrix(NA, nrow = nIter+1, ncol = dim(sigma)[1])
89   theta[1,]<- theta_0
90   rej_rate<-0
91   for(i in 2:(nIter+1)){
92     new_theta<- rmvnorm(1, mean = theta[i-1,], sigma = constant*sigma)
93     cur_theta<-t(as.matrix(theta[i-1,]))
94     U<-runif(1,0,1)
95     num<-logPostFunc(new_theta,...)+dmvnorm(cur_theta, mean =new_theta, sigma = sigma, log =
      TRUE)
96     den<-logPostFunc(cur_theta,...)+dmvnorm(new_theta, mean =cur_theta, sigma = sigma, log =
      TRUE)
97
98     if(U<min(1,exp(num-den))){
99       theta[i,] <-new_theta
100    }else{
101      theta[i,] <-cur_theta
102      rej_rate <-rej_rate+1
103    }
104  }
105
106  myres<- list(theta= theta, rej_rate= rej_rate/nIter)
107  return(myres)
108 }
109
110
111 res<-metrop1(logPostFunc=LogPostPoisson,
112             theta_0= rep(0,nPara),
113             constant= 0.6,
114             sigma = J ,
115             priormu= mu_0,
116             priorsigma= Sigma_0,
117             X= X, Y=y,
118             nIter= 10000)
119
120
121 res[[2]]
122
123 compbetas<- data.frame(OptimCoefficients= as.vector(BetaCoef), MCMCCoef =colMeans(res[[1]]))
124 rownames(compbetas)<- covNames
125 compbetas
126 ##d
127 Xpred <- matrix(c(1, 1, 1, 1, 0, 0, 0, 1, 0.5), nrow = 1)
128 predsamples <- rpois(10000, lambda = exp(Xpred %*% t(res[[1]])))
129 hist(predsamples, freq = FALSE)
130
131 ##probability of 0
132 length(which(predsamples==0))/length(predsamples)

```