

732A91: Lab 3

Bayesian Learning

Sarah Alsaadi, Carles Sans Fuentes

May 22, 2017

Normal model, mixture of normal model with semi-conjugate prior

The data rainfall.dat consists of daily records, from the beginning of 1948 to the end of 1983, of precipitation (rain or snow in units of 1 inch, and records of zero 100 precipitation are excluded) at Snoqualmie Falls, Washington. Analyze the data using the following two models.

1. Assume the daily precipitation $\{y_1, \dots, y_n\}$ are independent normally distributed, $y_1, \dots, y_n | \mu, \sigma^2 \sim N(\mu, \sigma^2)$ where both μ and σ^2 are unknown. Let $\mu \sim N(\mu_0, \tau_0^2)$ independently of $\sigma^2 \sim \text{Inv}\chi^2(\nu_0, \sigma_0^2)$

- (a) Implement (code!) a Gibbs sampler that simulates from the joint posterior $p(\mu, \sigma^2 | y_1, \dots, y_n)$. Where the full conditional posteriors are given by:

- i. $\mu | \sigma^2, y_1, \dots, y_n \sim N(\mu_n, \tau_n^2)$ where $\mu_n = \frac{n/\sigma^2}{n/\sigma^2 + 1/\tau_0^2} \bar{y} + \frac{1/\tau_0^2}{n/\sigma^2 + 1/\tau_0^2} \mu_0$ and $\tau_n^2 = \frac{1}{n/\sigma^2 + 1/\tau_0^2}$

- ii. $\sigma^2 | \mu, y_1, \dots, y_n \sim \text{Inv} - \chi^2(v_0 + n, \frac{v_0 \sigma_0^2 + \sum_{i=1}^n (y_i - \mu)^2}{n + v_0})$.

The initial values have been set up near to 0 but τ , which has been chosen as 10. Proving with different τ values showed us that small τ , which accounts for variance, can lead our distribution of posterior estimates to stuck in local minimas since it is not able to get out from an optima. Still, only results with $\tau = 10$ will be shown, but it is interesting to know that with an smaller τ our optimal posterior μ distribution was lower (at around 26-28).

- (b) Here below in figure 1 it can be seen how the posterior of our μ and σ converges to 31-33 and 1550 respectively. There is a burning period for σ and μ to converge to the distribution.

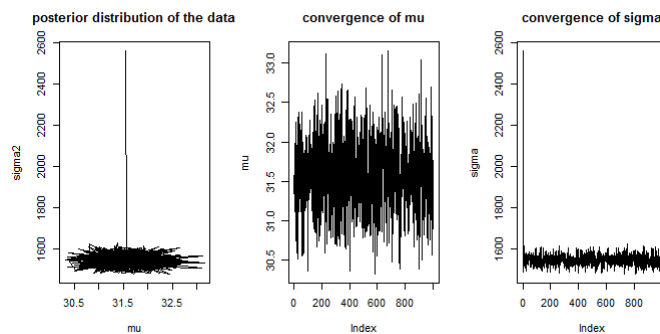


Figure 1: 3 plots of the posterior distribution of μ and σ with 1000 draws

In 2, we can see how the μ parameter is distributed with its cumulative mean that goes to 32 and how data is correlated. It can be seen that data is correlated just if we take each or every 2 observations. Thus, every third answer should be taken because it does not show any more autocorrelation or dependence.

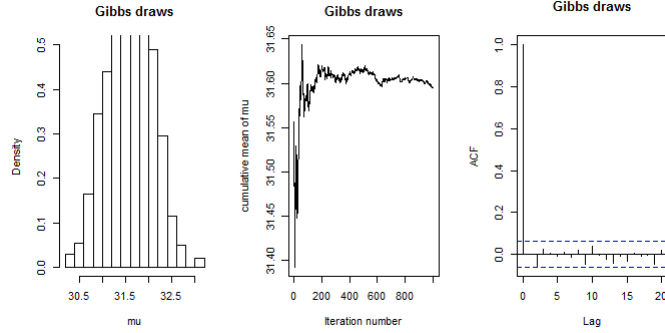


Figure 2: distribution of μ with 1000 draws, cumulative mean of the mean of μ and correlation lags

- (c) In 1b, we are asked to run a mixture of normal on the data for 2 μ and σ , assuming that data could come from two different points in Sweden.

In the following figure 3 below it can be seen the convergence of the mixture density model after 100 iterations.

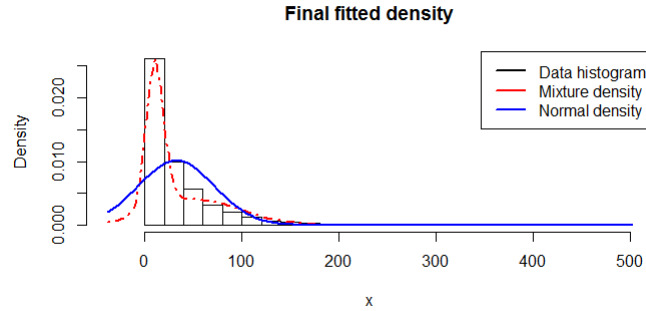


Figure 3: Convergence of the mixture density model after 100 iterations

Now (in 1c) we are asked to plot in one figure 4 to 1) a histogram or kernel density estimate of the data, 2) Normal density $N(\mu, \sigma^2)$ in (a) and 3) Mixture of normals in (b). This is done in the following graph being the red line the density 1000 random sample numbers from the distribution

$$N(\mu = \text{mean}(\mu_1), sd = \text{sqrt}(\text{mean}(\text{sigma2})))$$

and the yellow one a mixture from activity b also from a 1000 random generated data from 1000 points with from

$$\pi * N(\mu_1, \sigma_1^2) + (1 - \pi) * N(\mu_2, \sigma_2^2)$$

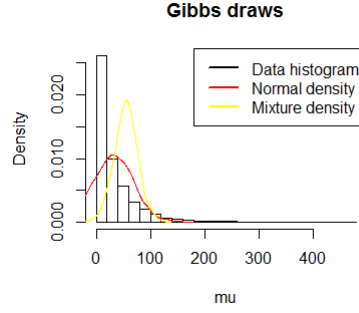


Figure 4: a histogram or kernel density estimate of the data, 2) Normal density $N(\mu, \sigma^2)$ in (a) and 3) Mixture of normals in(b)

Results show that none of both distributions is good for modeling the data. It is not a good idea just to take one or two normal distributions to model exponential data. More distributions would be necessary to get a better model.

Probit regression

1. Implement (code!) a data augmentation Gibbs sampler for the probit regression model

$$Pr(y = 1|\mathbf{x}) = \Phi(\mathbf{x}^T \beta)$$

The code is given in the appendix.

2. Compute the posterior of β in the probit regression for the WomenWork dataset from Lab 2 using the prior $\beta \sim N(0, \tau^2 I)$, with $\tau = 10$.
3. Do a normal approximation $\beta|\mathbf{y}, \mathbf{X} \sim \mathbf{N}(\tilde{\beta}, \mathbf{J}_{\mathbf{y}}^{-1}(\tilde{\beta}))$ of the posterior for β in the probit regression. Compare with the results from 2(b). Is the normal approximation accurate? In figure 5 and 6, you can see the distribution of our parameters. Also the the mean is provided with its 95% intervals below:

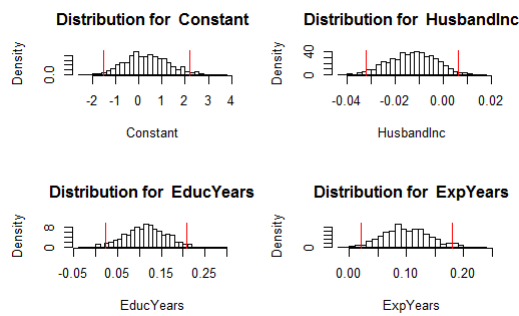


Figure 5: Distribution of the parameters with its 95% confidence interval (I)

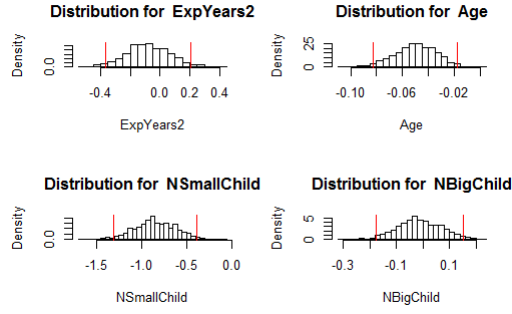


Figure 6: Distribution of the parameters with its 95% confidence interval (II)

```

1 > BetaFeat
2
3      Mean      upper      lower
4 Constant  0.34057957  2.180953926 -1.49979478
5 HusbandInc -0.01299479  0.006076275 -0.03206585
6 EducYears  0.11455707  0.207338016  0.02177613
7 ExpYears   0.10056627  0.179791545  0.02134099
8 ExpYears2  -0.07799728  0.207495176 -0.36348974
9 Age        -0.04989447 -0.017475928 -0.08231301
10 NSmallChild -0.84840269 -0.388974933 -1.30783046
11 NBigChild  -0.01189881  0.152902525 -0.17670015

```

Now we have been asked to do a normal approximation which results can be seen here below. It can be seen that results are somehow similar, though not exact.

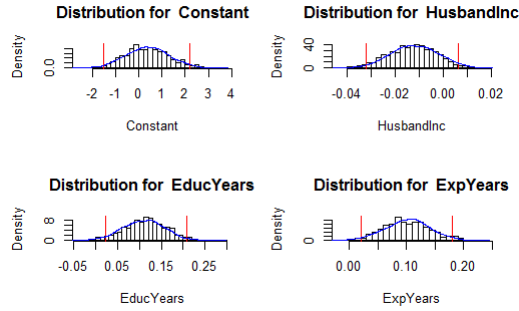


Figure 7: Distribution of the parameters with its 95% confidence interval and blue line for the normal approximation (I)

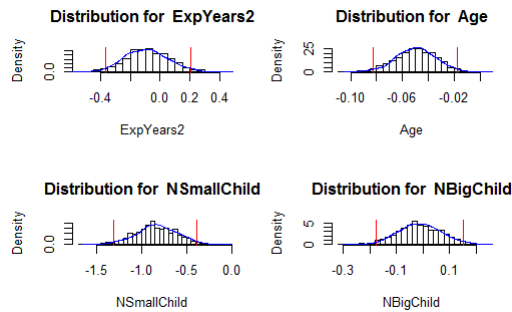


Figure 8: Distribution of the parameters with its 95% confidence interval and blue line for the normal approximation (II)

Also a table with the comparison of the different coefficients for both methods is reported.

```

1  > ConclTable
2      GibbsMean      Gibbssd OptimalNBeta      Nsd
3 Constant      0.34057957 0.938966506 0.38451729 0.914146031
4 HusbandInc    -0.01299479 0.009730134 -0.01208712 0.009474897
5 EducYears      0.11455707 0.047337216 0.10840175 0.047218262
6 ExpYears      0.10056627 0.040421059 0.10221471 0.039459840
7 ExpYears2     -0.07799728 0.145659416 -0.09024608 0.140976385
8 Age           -0.04989447 0.016540073 -0.04962341 0.015890547
9 NSmallChild   -0.84840269 0.234401919 -0.81567500 0.226635007
10 NBigChild    -0.01189881 0.084082316 -0.01241943 0.081506596

```

Contributions

All results and comments presented have been developed and discussed together by the members of the group.

Appendix

Question 1

```
1
2 library(geoR)
3 library(mvtnorm)
4
5 set.seed(12345)
6
7 data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part3/rainfall.dat.txt")
8
9
10 #####1a
11
12 ##Initial Data
13
14
15 nDraws<-1000
16 data<- unlist(data)
17 mu_0<- 0
18 v_0<- 1
19 sigma2_0<- 1
20 tau2_0<- 10
21
22
23 Normal_model<- function(nDraws, data,mu_0, v_0, tau2_0, sigma2_0){
24   require(geoR)
25   ##initializing data result
26   gibbsDraws <- matrix(0,nDraws,2)
27   colnames(gibbsDraws)<-c("mu", "sigma2")
28
29   #basic variables
30   n<- length(data)
31   v_n<- n +v_0
32
33   #Initializing basic variables for the gibb sampling
34   mu<-mu_0
35
36   for(i in 1:nDraws){
37
38     sigma2 <- rinvchisq(1, df= v_n, scale = (v_0*sigma2_0+ sum((data-mu)**2))/(n +v_0))
39     gibbsDraws[i,2] <- sigma2
40
41     w<- (n/sigma2)/(n/sigma2+1/tau2_0)
42     mu_n<- w*mean(data)+(1-w)*mu_0
43     invtau2<- (1/(n/sigma2+1/tau2_0))
44
45     mu <-rnorm(1,mu_n, sd = sqrt(invtau2))
46     gibbsDraws[i,1]<- mu
47
48   }
49
50   return(gibbsDraws)
51
52 }
53
54
55 gibbsDraws<-Normal_model(nDraws=nDraws, data= data ,mu_0= mu_0, v_0= v_0, tau2_0= tau2_0,
56   sigma2_0= sigma2_0)
57 tail(gibbsDraws)
58
59 par(mfrow = c(1,3))
60 plot(gibbsDraws, type = "l", main = "posterior distribution of the data")
61 plot(gibbsDraws[,1], type = "l", ylab = "mu", main = "convergence of mu")
62 plot(gibbsDraws[,2], type = "l", ylab = "sigma", main = "convergence of sigma")
63
64
65 hist(gibbsDraws[,1], freq = FALSE, main='Gibbs draws', ylim = c(0,0.5), xlab= "mu")
66 lines(seq(-2,4,by=0.01),dnorm(seq(-2,4,by=0.01), mean = 1), col = "red", lwd = 3)
67 plot(cumsum(gibbsDraws[,1])/seq(1,nDraws),type="l", main='Gibbs draws', xlab='Iteration number',
68   ylab='cumulative mean of mu')
69 lines(seq(1,nDraws),1*matrix(1,1,nDraws),col="red",lwd=3)
70 acf(gibbsDraws[,1], main='Gibbs draws', lag.max = 20)
71 par(mfrow = c(1,1))
72
73 #####1b
74
75 ##### BEGIN USER INPUT #####
76 # Data options
77 rawData <- data
78 x <- as.matrix(data)
79
80 # Model options
```

```

80 nComp <- 2      # Number of mixture components
81
82 # Prior options
83 alpha <- 10*rep(1,nComp) # Dirichlet(alpha)
84 muPrior <- rep(0,nComp) # Prior mean of theta
85 tau2Prior <- rep(10,nComp) # Prior std theta
86 sigma2_0 <- rep(var(x),nComp) # s20 (best guess of sigma2)
87 nu0 <- rep(4,nComp) # degrees of freedom for prior on sigma2
88
89 # MCMC options
90 nIter <- 100 # Number of Gibbs sampling draws
91
92 # Plotting options
93 plotFit <- TRUE
94 lineColors <- c("blue", "green", "magenta", 'yellow')
95 sleepTime <- 0.1 # Adding sleep time between iterations for plotting
96 ##### END USER INPUT #####
97
98 ##### Defining a function that simulates from the
99 rScaledInvChi2 <- function(n, df, scale){
100   return((df*scale)/rchisq(n,df=df))
101 }
102
103 ##### Defining a function that simulates from a Dirichlet distribution
104 rDirichlet <- function(param){
105   nCat <- length(param)
106   thetaDraws <- matrix(NA,nCat,1)
107   for (j in 1:nCat){
108     thetaDraws[j] <- rgamma(1,param[j],1)
109   }
110   thetaDraws = thetaDraws/sum(thetaDraws) # Dividing every column of ThetaDraws by the sum of the
111     elements in that column.
112   return(thetaDraws)
113 }
114
115 # Simple function that converts between two different representations of the mixture allocation
116 S2alloc <- function(S){
117   n <- dim(S)[1]
118   alloc <- rep(0,n)
119   for (i in 1:n){
120     alloc[i] <- which(S[i,] == 1)
121   }
122   return(alloc)
123 }
124
125 # Initial value for the MCMC
126 nObs <- length(x)
127 S <- t(rmultinom(nObs, size = 1, prob = rep(1/nComp,nComp))) # nObs-by-nComp matrix with
128   component allocations.
129 theta <- quantile(x, probs = seq(0,1,length = nComp))
130 sigma2 <- rep(var(x),nComp)
131 probObsInComp <- rep(NA, nComp)
132
133 # Setting up the plot
134 xGrid <- seq(min(x)-1*apply(x,2,sd),max(x)+1*apply(x,2,sd),length = 100)
135 xGridMin <- min(xGrid)
136 xGridMax <- max(xGrid)
137 mixDensMean <- rep(0,length(xGrid))
138 effIterCount <- 0
139 ylim <- c(0,2*max(hist(x)$density))
140
141 ##Recording mu and sigma2
142 matmu<- matrix(0, nrow = nIter, ncol = nComp)
143 colnames(matmu)<- c("mu1", "mu2")
144
145 matsigma2<- matrix(0, nrow = nIter, ncol = nComp)
146 colnames(matsigma2)<- c("sigma1", "sigma2")
147
148 matpi<- matrix(0, nrow = nIter, ncol = nComp)
149 colnames(matsigma2)<- c("pi1", "pi2")
150
151
152 for (k in 1:nIter){
153   message(paste('Iteration number:',k))
154   alloc <- S2alloc(S) # Just a function that converts between different representations of the
155     group allocations
156   nAlloc <- colSums(S)
157   print(nAlloc)
158   # Update components probabilities
159   w <- rDirichlet(alpha + nAlloc)
160   matpi[k,]<- w
161
162   # Update theta's
163   for (j in 1:nComp){

```



```

164     precPrior <- 1/tau2Prior[j]
165     precData <- nAlloc[j]/sigma2[j]
166     precPost <- precPrior + precData
167     wPrior <- precPrior/precPost
168     muPost <- wPrior*muPrior + (1-wPrior)*mean(x[alloc == j])
169
170     tau2Post <- 1/precPost
171     theta[j] <- rnorm(1, mean = muPost, sd = sqrt(tau2Post))
172 }
173 matmu[k,]<- muPost
174
175
176 # Update sigma2's
177 for (j in 1:nComp){
178     sigma2[j] <- rScaledInvChi2(1, df = nu0[j] + nAlloc[j], scale = (nu0[j]*sigma2_0[j] + sum((
179         x[alloc == j] - theta[j])^2))/(nu0[j] + nAlloc[j]))
180 }
181 matsigma2[k,]<- sigma2
182
183 # Update allocation
184 for (i in 1:nObs){
185     for (j in 1:nComp){
186         probObsInComp[j] <- w[j]*dnorm(x[i], mean = theta[j], sd = sqrt(sigma2[j]))
187     }
188     S[i,] <- t(rmultinom(1, size = 1, prob = probObsInComp/sum(probObsInComp)))
189 }
190
191 # Printing the fitted density against data histogram
192 if (plotFit && (k%1 == 0)){
193     effIterCount <- effIterCount + 1
194     hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = paste("Iteration
195         number",k), ylim = ylim)
196     mixDens <- rep(0,length(xGrid))
197     components <- c()
198     for (j in 1:nComp){
199         compDens <- dnorm(xGrid,theta[j],sd = sqrt(sigma2[j]))
200         mixDens <- mixDens + w[j]*compDens
201         lines(xGrid, compDens, type = "l", lwd = 2, col = lineColors[j])
202         components[j] <- paste("Component ",j)
203     }
204     mixDensMean <- ((effIterCount-1)*mixDensMean + mixDens)/effIterCount
205     lines(xGrid, mixDens, type = "l", lty = 2, lwd = 3, col = 'red')
206     legend("topleft", box.lty = 1, legend = c("Data histogram",components, 'Mixture'),
207         col = c("black",lineColors[1:nComp], 'red'), lwd = 2)
208     Sys.sleep(sleepTime)
209 }
210
211 }
212
213 hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = "Final fitted density")
214 lines(xGrid, mixDensMean, type = "l", lwd = 2, lty = 4, col = "red")
215 lines(xGrid, dnorm(xGrid, mean = mean(x), sd = apply(x,2,sd)), type = "l", lwd = 2, col = "blue
216 ")
217 legend("topright", box.lty = 1, legend = c("Data histogram","Mixture density","Normal density")
218 , col=c("black","red","blue"), lwd = 2)
219
220 #####C Graphical representation
221 #Data sets for the distribution of mu1, mu2, sigma2_1, sigma_2, pi1, pi2
222 matmu
223 matpi
224 matsigma2
225 #Data set for the distribution of mu1, sigma2_1 in the first case
226
227
228 sampled_simple<-rnorm(1000, mean = mean(gibbsDraws[,1]), sd = sqrt(mean(gibbsDraws[,2])))
229 mixture_sample<- mean(matpi[,1])*rnorm(1000, mean = mean(matmu[,1]), sd = sqrt(mean(matsigma2
230     [,1])))+ (1-mean(matpi[,1]))*rnorm(1000, mean = mean(matmu[,2]), sd = sqrt(mean(matsigma2
231     [,2])))
232
233 hist(data, freq = FALSE, main='Gibbs draws', xlab= "mu", breaks = 30)
234 lines(density(sampled_simple), col = "red")
235 lines(density(mixture_sample), col = "yellow")
236 legend("topright", box.lty = 1, legend = c("Data histogram","Normal density", "Mixture density"
237     ), col=c("black","red","yellow"), lwd = 2)

```

Question 2

```

1
2 #####Question 2

```

```

3 #install.packages("msm")
4 library(msm)
5
6 Data<- read.csv("C:/Users/Carles/Desktop/Bayesian learning/Part2/WomenWork.dat.txt", sep = "",
7   header = TRUE)
8 glmModel <- glm(Work ~ 0 + ., data = Data, family = binomial(link = "probit"))
9 summary(glmModel)
10 ##### Variables
11 tau <- 10; # Prior scaling factor such that Prior Covariance = (tau^2)*I
12 chooseCov <- c(1:8) # Here we choose which covariates to include in the model
13
14
15 y <- as.vector(Data[,1]); # Data from the read.table function is a data frame. Let's convert y
   and X to vector and matrix.
16 X <- as.matrix(Data[,2:ncol(Data)]);
17 initVal <- as.vector(rep(0,dim(X)[2]));
18 covNames <- names(Data)[2:length(names(Data))];
19 X <- X[,chooseCov]; # Here we pick out the chosen covariates.
20 covNames <- covNames[chooseCov];
21 nPara <- dim(X)[2];
22
23
24
25 # Setting up the prior
26 mu_0 <- as.vector(rep(0,nPara)) # Prior mean vector
27 Omega_0 <- tau^2*diag(nPara);
28 beta_0 <- as.vector(rmvnorm(1, mean = mu_0, sigma =Omega_0))
29 nIter <- 1000
30
31
32
33 ProbitFunction<- function(beta_0, mu_0, Omega_0, X, y, n_iter){
34   posy_0 <- which(y == 0)
35   posy_1 <- which(y == 1)
36   nPara <- dim(X)[2]
37
38   beta<-matrix(0, ncol = nPara, nrow = n_iter)
39   beta[1,]<- beta_0
40   B<- solve(solve(Omega_0)+t(X)%*%X)
41
42   for(i in 1:n_iter){
43     beta[i, ]<- rmvnorm(1, as.vector(B%*(solve(Omega_0)%*%mu_0+t(X)%*%y)), sigma = B)
44
45     y[posy_1] <- rtnorm(length(posy_1), mean = as.vector(X[posy_1,]%*%beta[i,]), sd = 1,lower =
46       0, upper = Inf)
47     y[posy_0] <- rtnorm(length(posy_0), mean = as.vector(X[posy_0,]%*%beta[i,]), sd = 1,lower =
48       -Inf, upper = 0)
49
50   }
51   return(beta)
52 }
53
54
55
56 Betadist<-ProbitFunction(beta_0=beta_0, mu_0= mu_0, Omega_0= Omega_0, X= X, y= y , n_iter =
   nIter)
57 colnames(Betadist)<- covNames
58 dim(Betadist)
59 BetaFeat<- data.frame(Mean = apply(Betadist, 2, mean),
60   upper = apply(Betadist, 2, mean)+1.96*apply(Betadist, 2, sd),
61   lower = apply(Betadist, 2, mean)-1.96*apply(Betadist, 2, sd))
62 par(mfrow = c(2,2))
63 for(i in 1: 4){
64   hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i],
65     breaks = 30, freq = FALSE)
66   abline(v = BetaFeat$lower[i], b = 0, col = "red")
67   abline(v = BetaFeat$upper[i], b = 0, col = "red")
68 }
69 par(mfrow = c(2,2))
70 for(i in 5:8){
71   hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i], breaks
72     = 30, freq = FALSE)
73   abline(v = BetaFeat$lower[i], b = 0, col = "red")
74   abline(v = BetaFeat$upper[i], b = 0, col = "red")
75 }
76
77
78 ###C
79
80 LogPostProbit <- function(betaVect,y,X,mu,Sigma){
81   nPara <- length(betaVect);
82   linPred <- X%*%betaVect;

```

```

83
84 # The following is a more numerically stable evaluation of the log-likelihood in my slides:
85 # logLik <- sum(y*log(pnorm(linPred)) + (1-y)*log(1-pnorm(linPred)))
86 logLik <- sum(y*pnorm(linPred, log.p = TRUE) + (1-y)*pnorm(linPred, log.p = TRUE, lower.tail
    = FALSE))
87
88 # evaluating the prior
89 logPrior <- dmvnorm(betaVect, matrix(0,nPara,1), Sigma, log=TRUE);
90
91 # add the log prior and log-likelihood together to get log posterior
92 return(logLik + logPrior)
93
94 }
95
96 OptimResults<-optim(initVal,LogPostProbit,gr=NULL,y,X,mu= mu_0,Sigma= Omega_0,method=c("BFGS"),
    control=list(fnscale=-1),hessian=TRUE)
97
98 BetaCoef<- OptimResults$par
99 names(BetaCoef)<- covNames
100 J<--solve(OptimResults$hessian)
101
102 myconf95<-c(lowbound =BetaCoef[6]-1.96*sqrt(J[6,6]), highbound=BetaCoef[6]+1.96*sqrt(J[6,6]))
103
104 sampleBetaProbit<- matrix(ncol = length(BetaCoef), nrow = 1000)
105 for(i in 1:length(BetaCoef)){
106 sampleBetaProbit[,i]<- rnorm(1000, mean = BetaCoef[i], sd = sqrt(J[i,i]))
107 }
108
109 par(mfrow = c(2,2))
110 for(i in 1: 4){
111 hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i], breaks
    = 30, freq = FALSE)
112 abline(v = BetaFeat$lower[i], b = 0, col = "red")
113 abline(v = BetaFeat$upper[i], b = 0, col = "red")
114 lines(density(sampleBetaProbit[,i]), col = "blue")
115 }
116 for(i in 5:8){
117 hist(Betadist[,i], main = paste("Distribution for " ,covNames[i]), xlab = covNames[i], breaks
    = 30, freq = FALSE)
118 abline(v = BetaFeat$lower[i], b = 0, col = "red")
119 abline(v = BetaFeat$upper[i], b = 0, col = "red")
120 lines(density(sampleBetaProbit[,i]), col = "blue")
121
122 }
123 }
124
125 ConclTable<- data.frame(GibbsMean = apply(Betadist, 2, mean),
126 Gibbsstd = apply(Betadist, 2, sd),
127 OptimalBeta = BetaCoef,
128 Gibbsstd=sqrt(diag(J)))

```