

Databases for Big Data – part 1

Valentina Ivanova

IDA, Linköping University

Outline – Today – Part 1

- RDBMS → NoSQL → NewSQL
 - DBMS – OLAP vs OLTP (ACID)
 - NoSQL Concepts and Techniques
 - Horizontal scalability
 - Consistency models
 - CAP theorem: BASE vs ACID
 - Consistent hashing
 - Vector clocks
 - NoSQL Systems - Types and Applications
 - Hadoop Distributed File System - HDFS
-

Outline – Next Lecture – Part 2

- Amazon DynamoDB
- HBase
- Hive
- Shark

DB rankings – September 2016

| Rank | | | DBMS | Database Model | Score | | |
|----------|----------|----------|----------------------|-------------------|----------|----------|----------|
| Sep 2016 | Aug 2016 | Sep 2015 | | | Sep 2016 | Aug 2016 | Sep 2015 |
| 1. | 1. | 1. | Oracle | Relational DBMS | 1425.56 | -2.16 | -37.81 |
| 2. | 2. | 2. | MySQL + | Relational DBMS | 1354.03 | -3.01 | +76.28 |
| 3. | 3. | 3. | Microsoft SQL Server | Relational DBMS | 1211.55 | +6.51 | +113.72 |
| 4. | ↑ 5. | ↑ 5. | PostgreSQL | Relational DBMS | 316.35 | +1.10 | +30.18 |
| 5. | ↓ 4. | ↓ 4. | MongoDB + | Document store | 316.00 | -2.49 | +15.43 |
| 6. | 6. | 6. | DB2 | Relational DBMS | 181.19 | -4.70 | -27.95 |
| 7. | 7. | ↑ 8. | Cassandra + | Wide column store | 130.49 | +0.26 | +2.89 |
| 8. | 8. | ↓ 7. | Microsoft Access | Relational DBMS | 123.31 | -0.74 | -22.68 |
| 9. | 9. | 9. | SQLite | Relational DBMS | 108.62 | -1.24 | +0.97 |
| 10. | 10. | 10. | Redis | Key-value store | 107.79 | +0.47 | +7.14 |
| 11. | 11. | ↑ 14. | Elasticsearch + | Search engine | 96.48 | +3.99 | +24.93 |
| 12. | 12. | ↑ 13. | Teradata | Relational DBMS | 73.06 | -0.57 | -1.20 |
| 13. | 13. | ↓ 11. | SAP Adaptive Server | Relational DBMS | 69.16 | -1.88 | -17.36 |
| 14. | 14. | ↓ 12. | Solr | Search engine | 66.96 | +1.19 | -14.98 |
| 15. | 15. | 15. | HBase | Wide column store | 57.81 | +2.30 | -1.22 |
| 16. | 16. | ↑ 17. | FileMaker | Relational DBMS | 55.35 | +0.34 | +4.35 |
| 17. | 17. | ↑ 18. | Splunk | Search engine | 51.29 | +2.38 | +9.06 |
| 18. | 18. | ↓ 16. | Hive | Relational DBMS | 48.82 | +1.01 | -4.71 |
| 19. | 19. | 19. | SAP HANA + | Relational DBMS | 43.42 | +0.68 | +5.22 |
| 20. | 20. | ↑ 25. | MariaDB | Relational DBMS | 38.53 | +1.65 | +14.31 |
| 21. | 21. | 21. | Neo4j + | Graph DBMS | 36.37 | +0.80 | +2.83 |
| 22. | ↑ 24. | ↑ 24. | Couchbase + | Document store | 28.54 | +1.14 | +2.28 |
| 23. | 23. | ↓ 22. | Memcached | Key-value store | 28.43 | +0.74 | -3.99 |
| 24. | ↓ 22. | ↓ 20. | Informix | Relational DBMS | 28.19 | -0.86 | -9.76 |
| 25. | 25. | ↑ 28. | Amazon DynamoDB + | Document store | 27.42 | +0.82 | +7.43 |

<http://db-engines.com/en/ranking>

RDBMS → NoSQL → NewSQL

DBMS history (Why NoSQL?)

- 1960 – Navigational databases
- 1970 – Relational databases (RDBMS)
- 1990 – Object-oriented databases and Data Warehouses
- 2000 – XML databases
- Mid 2000 – first NoSQL
- 2011 – NewSQL

RDBMS

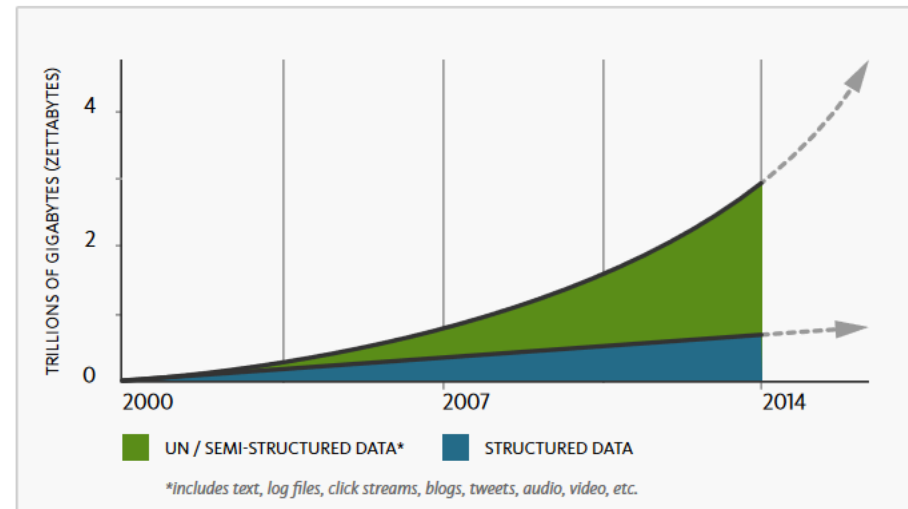
- Established technology
- Transactions support & ACID properties
- Powerful query language - SQL
- Experiences administrators
- Many vendors

Table: Item

| item id | name | color | size |
|---------|-------|-------|------|
| 45 | skirt | white | L |
| 65 | dress | red | M |

But ... – One Size Does Not Fit All^[1]

- Requirements have changed:
 - Frequent schema changes, management of unstructured and semi-structured data
 - Huge datasets
 - RDBMSs are not designed to be
 - distributed
 - continuously available
 - High read and write scalability
 - Different applications have different requirements^[1]



[1] “One Size Fits All”: An Idea Whose Time Has Come and Gone https://cs.brown.edu/~ugur/fits_all.pdf

Figure from: <http://www.couchbase.com/sites/default/files/uploads/all/whitepapers/NoSQL-Whitepaper.pdf>

NoSQL (not-only-SQL)

- A broad category of disparate solutions
- Simple and flexible non-relational data models
- High availability & relax data consistency requirement (CAP theorem)
 - BASE vs ACID
- Easy to distribute – horizontal scalability
- Data are replicated to multiple nodes
 - Down nodes easily replaced
 - No single point of failure
- Cheap & easy (or not) to implement (open source)

But ...

- No support for SQL → Low level programming → data analysts need to write custom programs
- No ACID
- Huge investments already made in SQL systems and experienced developers
- NoSQL systems do not provide interfaces to existing tools

NewSQL^[DataMan]

- First mentioned in 2011
- Supports the relational model
 - with horizontal scalability & fault tolerance
- Query language - SQL
- ACID
- Different data representation internally
- VoltDB, NuoDB, Clustrix, Google Spanner

NewSQL Applications^[DataMan]

- RBDMS applicable scenarios
 - transaction and manipulation of more than one object, e.g., financial applications
 - strong consistency requirements, e.g., financial applications
 - schema is known in advance and unlikely to change a lot
- But also Web-based applications^[1]
 - with different collection of OLTP requirements
 - multi-player games, social networking sites
 - real-time analytics (vs traditional business intelligence requests)

[1] <http://cacm.acm.org/blogs/blog-cacm/109710-new-sql-an-alternative-to-nosql-and-old-sql-for-new-oltp-apps/fulltext>

DBMS – OLAP and OLTP

DBMS applications – OLAP and OLTP

- OLTP – Online transaction processing - RDBMS
 - university database; bank database; a database with cars and their owners; online stores
- OLAP – Online analytical processing - Data warehouses
 - Summaries of multidimensional data

Example: sale (item, color, size, quantity)

What color/type of clothes is popular this season?

DBMS applications – OLTP

Table: Order

| order id | customer |
|----------|----------|
| 1 | 22 |
| 2 | 33 |

Table: Cart

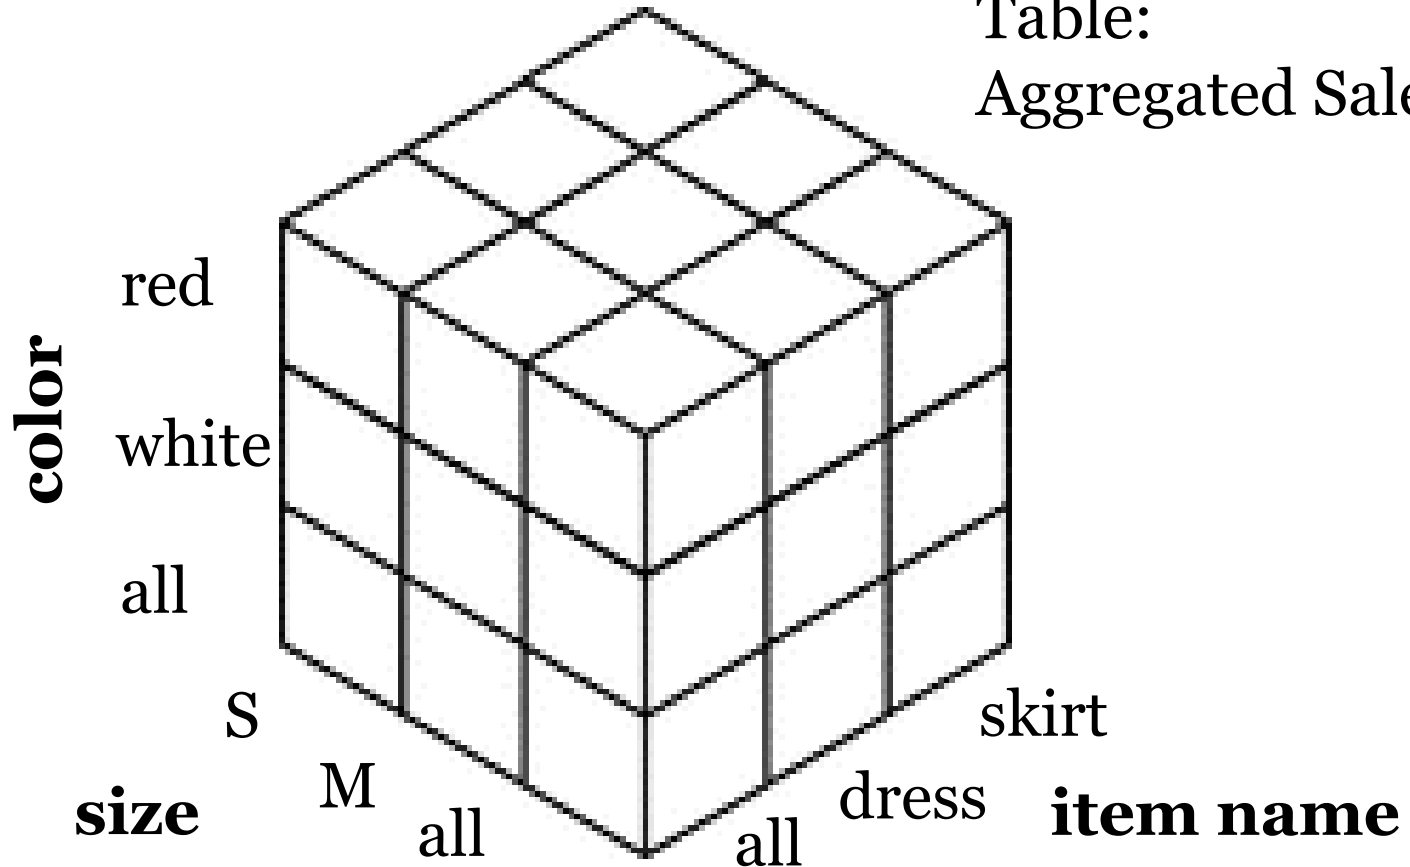
| order id | Item id | quantity |
|----------|---------|----------|
| 1 | 45 | 1 |
| 1 | 55 | 1 |
| 1 | 65 | 2 |
| 2 | 65 | 1 |

Table: Item

| item id | name | color | size |
|---------|-------|-------|------|
| 45 | skirt | white | L |
| 65 | dress | red | M |

DBMS applications – OLAP

Table:
Aggregated Sales



DBMS applications – OLAP and OLTP

- Relational DBMS vs Data Warehouse

<http://datawarehouse4u.info/OLTP-vs-OLAP.html>

| | RDBMS (OLTP) | Data Warehouse (OLAP) |
|--------------------|---|---|
| Source of data | Operational data; OLTPs are the original source of the data. | Consolidation data; OLAP data comes from the various OLTP DBs |
| Purpose of data | To control and run fundamental business tasks | To help with planning, problem solving, and decision support |
| What the data | Reveals a snapshot of ongoing business processes | Multi-dimensional views of various kinds of business activities |
| Inserts & Updates | Short and fast inserts and updates initiated by end users | Periodic long-running batch jobs refresh the data |
| Queries | Relatively standardized and simple queries returning relatively few records | Often complex queries involving aggregations |
| Processing Speed | Typically very fast | Depends on the amount of data involved |
| Space Requirements | Can be relatively small if historical data is archived | Larger due to the existence of aggregation structures and history data; |
| Database Design | Highly normalized, many tables | Typically de-normalized, fewer tables |
| Backup & Recovery | Highly important | Reloading from OLTPs |

DBMS applications – OLTP

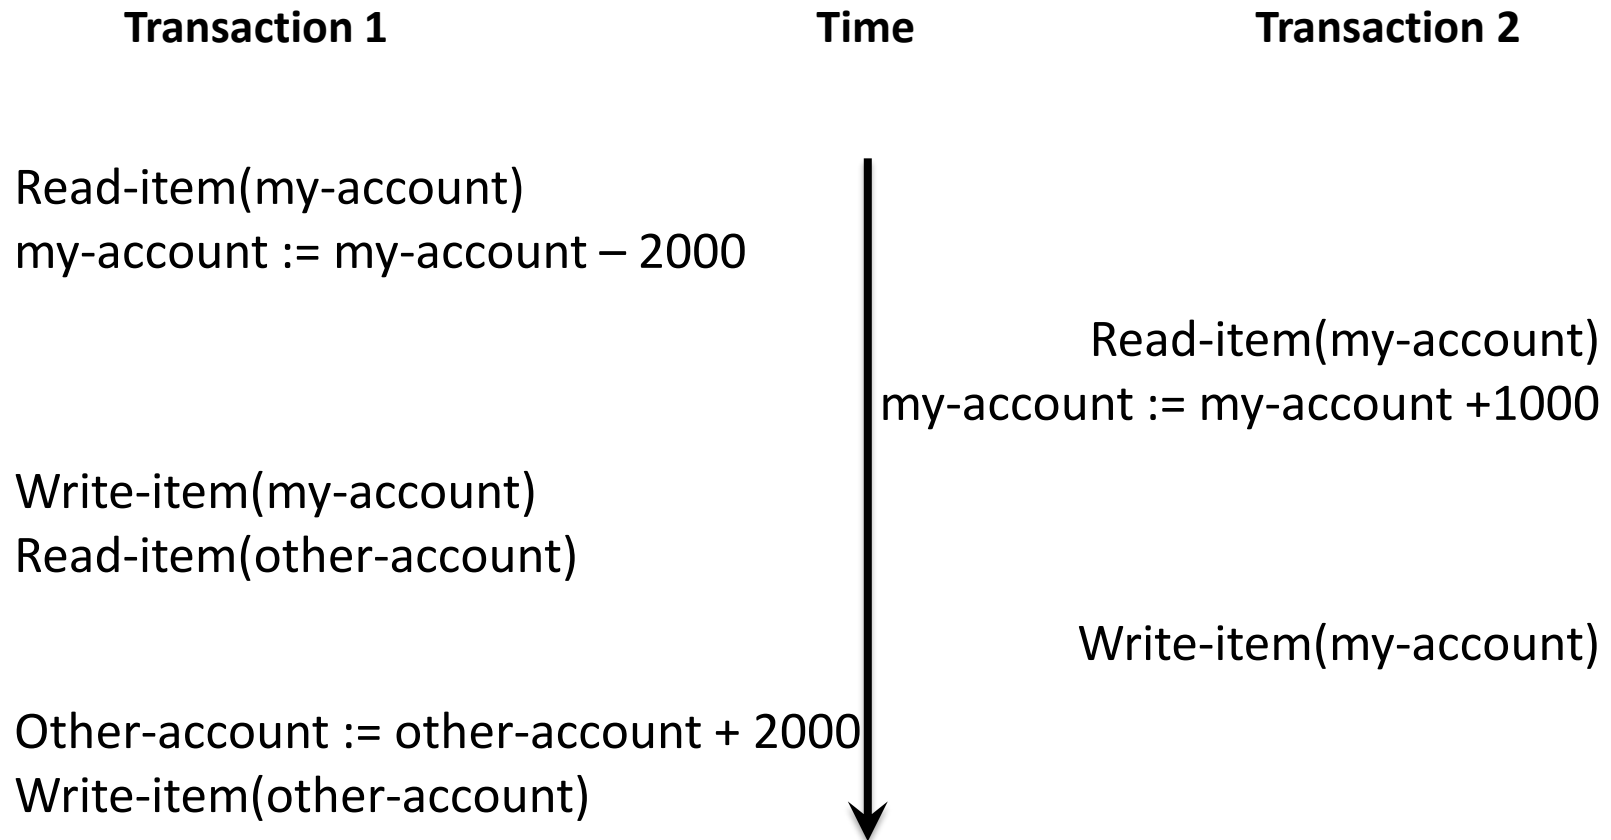
- OLTP – Online transaction processing
 - large number of data reads, writes and updates → transactions!

```
Read-item(my-account)
my-account := my-account - 2000
Write-item(my-account)
Read-item(other-account)
other-account := other-account + 2000
Write-item(other-account)
```

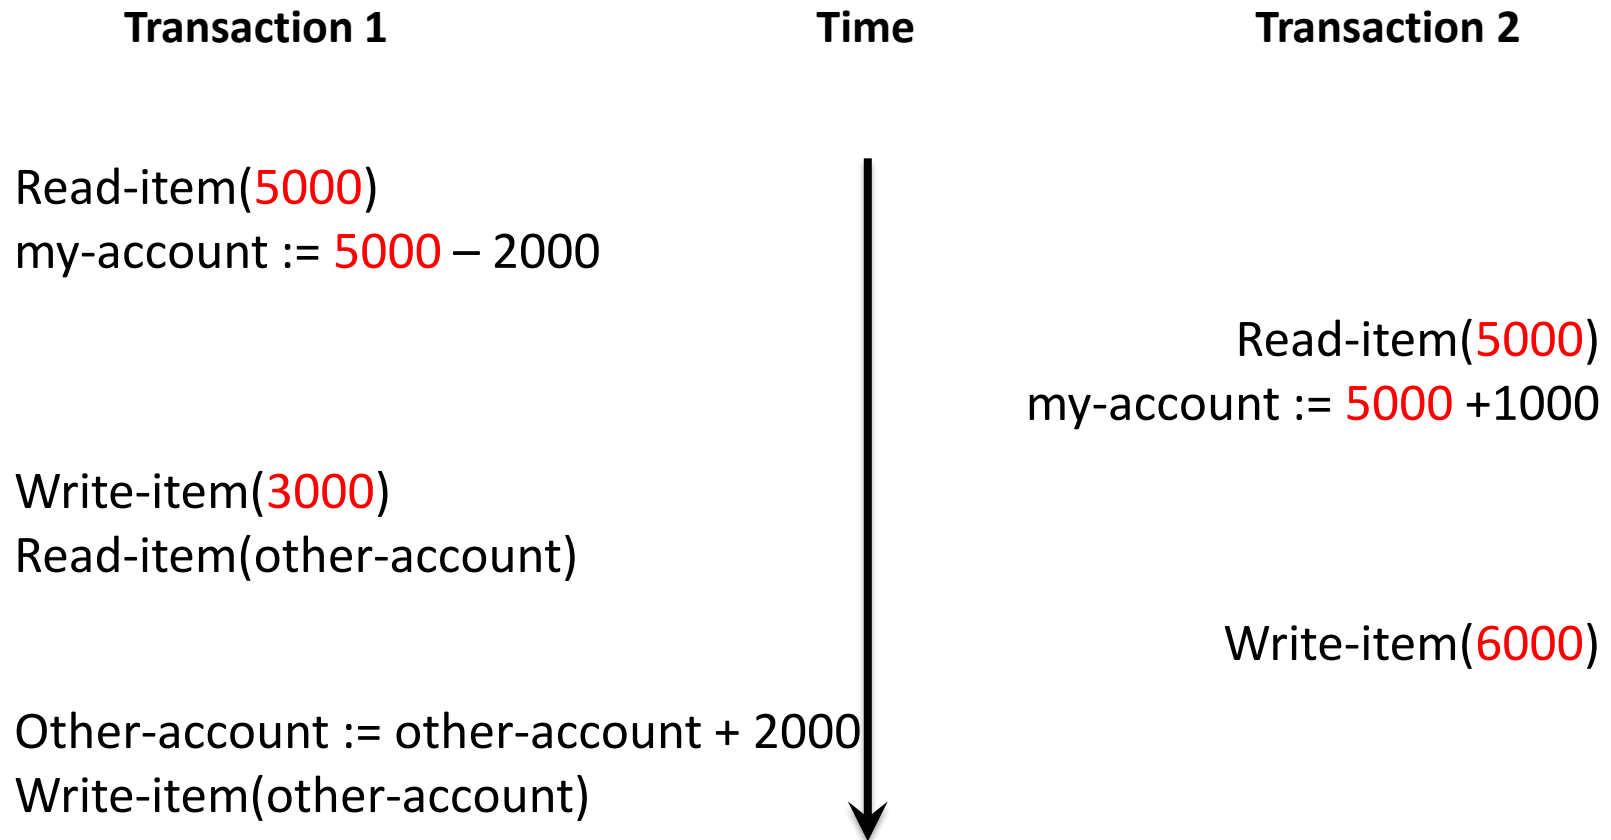
DBMS applications – Transactions

- A transaction is a logical unit of database processing and consists of one or several operations.
 - Begin transaction
 - Reads and writes
 - Commit or rollback
 - End transaction
- It leaves the database in a consistent state

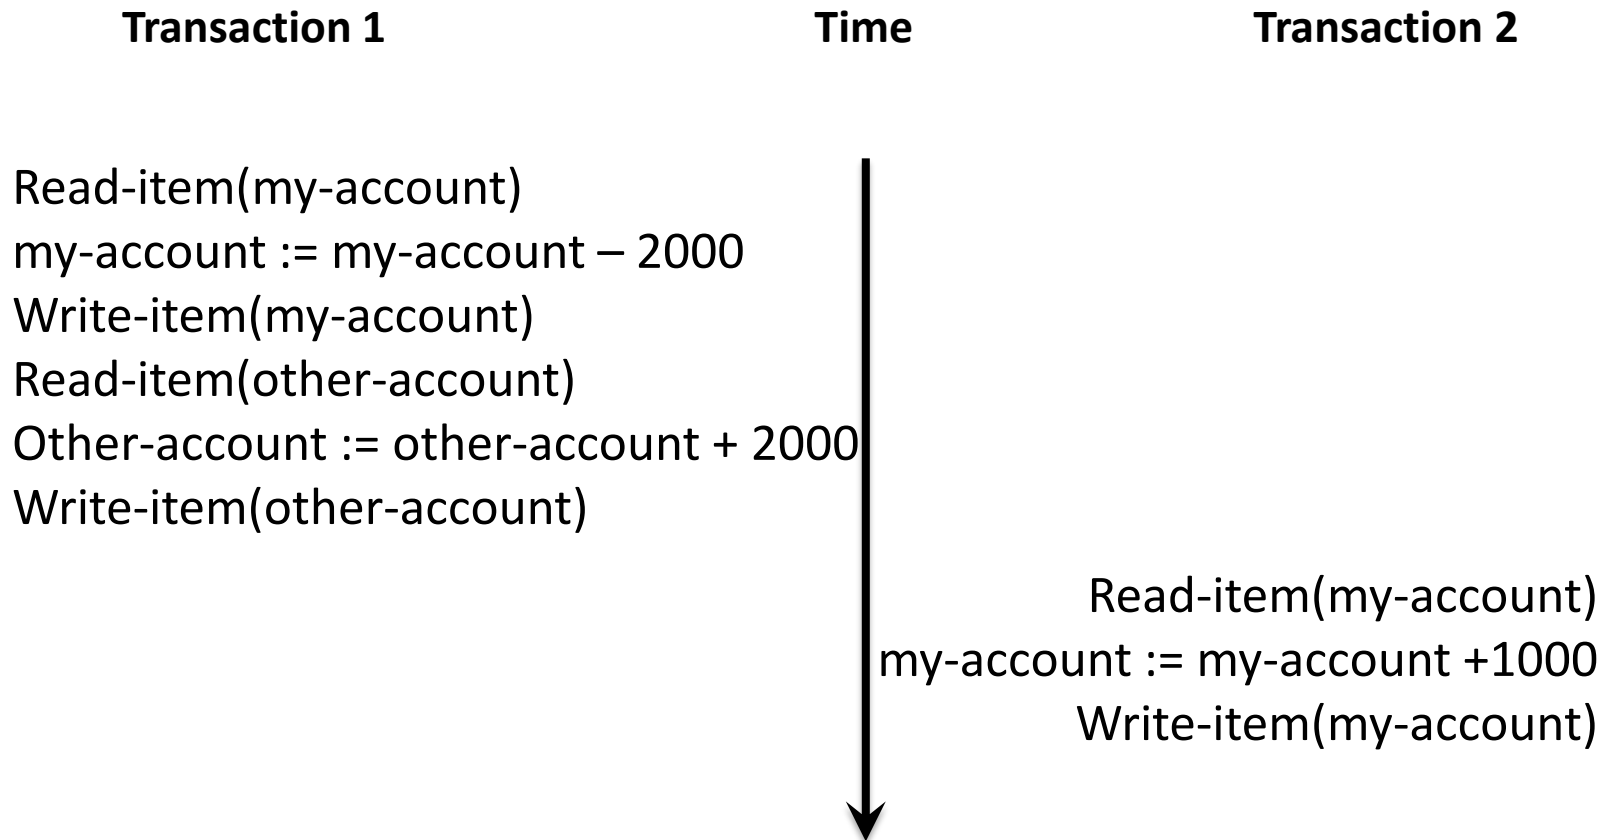
DBMS applications – Transactions



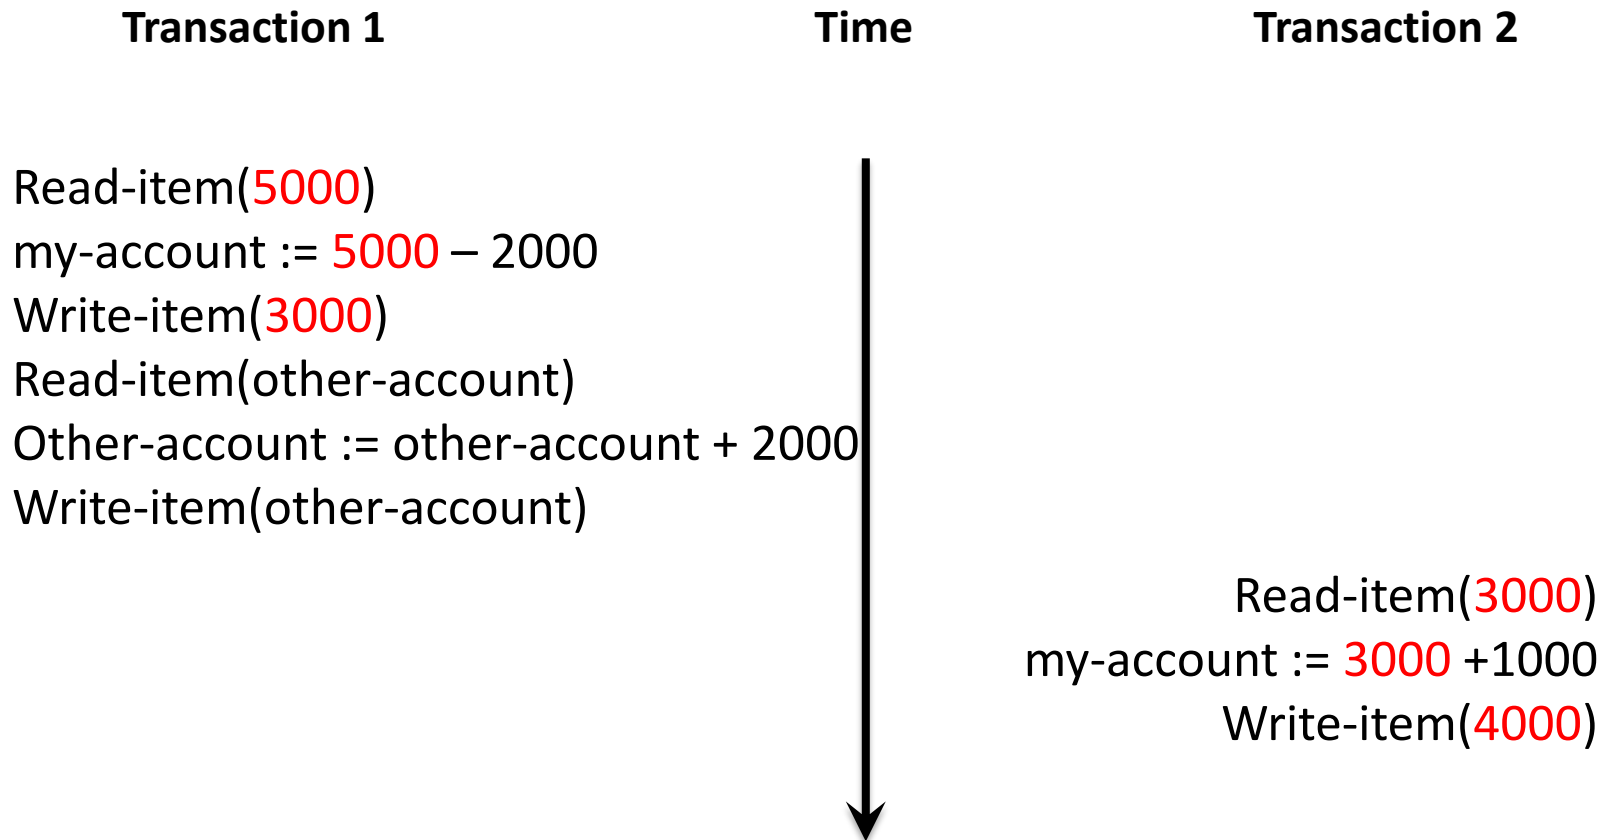
DBMS applications – Transactions



DBMS applications – Transactions



DBMS applications – Transactions



Transactions - ACID properties

- **Atomicity** → A transaction is an atomic unit: It is either executed completely or not at all
- **Consistency** → A database that is in a consistent state before the execution of a transaction, is also in a consistent state after the execution of the transaction
- **Isolation** → A transaction should act as if it is executed in isolation from the other transactions
- **Durability** → Changes in the database made by a committed transaction are permanent

NoSQL Concepts and Techniques

NoSQL Databases (not only SQL)

nosql-database.org

NoSQL Definition:

Next Generation Databases mostly addressing some of the points: being ***non-relational, distributed, open source*** and ***horizontally scalable***.

The original intention has been modern web-scale databases. ... Often more characteristics apply as: ***schema-free, easy replication support, simple API, eventually consistent/BASE*** (not ACID), a ***huge data amount***, and more.

NoSQL: Concepts

Scalability: system can handle growing amounts of data without losing performance.

- **Vertical Scalability (scale up)**
 - add resources (more CPUs, more memory) to a single node
 - using more threads to handle a local problem
- **Horizontal Scalability (scale out)**
 - add nodes (more computers, servers) to a distributed system
 - gets more and more popular due to low costs for commodity hardware
 - often surpasses scalability of vertical approach

Distributed (Data Management) Systems

- Number of processing nodes interconnected by a computer network
 - Data is stored, replicated, updated and processed across the nodes
 - Networks failures are given, not an exception
 - Network is partitioned
 - Communication between nodes is an issue
- Data consistency vs Availability

Consistency models^[Vogels]

- A distributed system through the developers' eyes
 - Storage system as a black box
 - Independent processes that write and read to the storage
- Strong consistency – after the update completes, any subsequent access will return the updated value.
- Weak consistency – the system does not guarantee that subsequent accesses will return the updated value.
 - inconsistency window

Consistency models^[Vogels]

- Weak consistency
 - Eventual consistency – if no new updates are made to the object, eventually all accesses will return the last updated value
 - Popular example: DNS

Consistency models^[Vogels]

- Server side view of a distributed system – Quorum
 - N – number of nodes that store replicas
 - R – number of nodes for a successful read
 - W – number of nodes for a successful write

Consistency models^[Vogels]

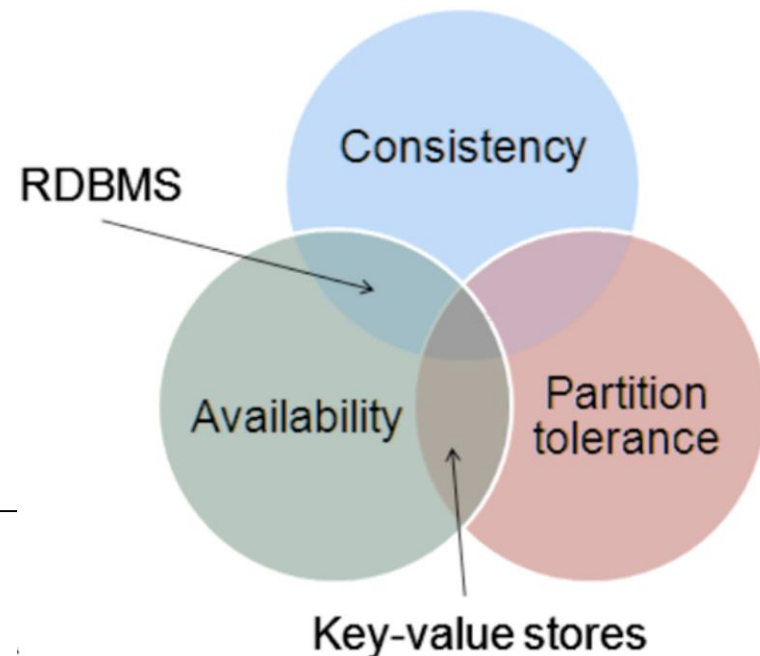
- Server side view of a distributed system – Quorum
 - High read loads – hundreds of N, $R=1$
 - Fault tolerance/availability (& relaxed consistency) $W=1$
 - $R + W > N$ strong consistency
 - Consistency (& reduced availability) $W=N$
 - $R + W \leq N$ eventual consistency
 - Inconsistency window – the period until all replicas have been updated in a lazy manner

NoSQL: Concepts

CAP Theorem: Consistency, Availability, Partition Tolerance ^[Brewer]

Theorem

(Gilbert, Lynch SIGACT'2002):
only 2 of the 3 guarantees
can be given in a shared-data
system.



NoSQL: Concepts

CAP Theorem: Consistency, Availability, Partition Tolerance^[Brewer]

- **Consistency**

- after an update, all readers in a distributed system see the same data
- all nodes are supposed to contain the same data at all times

- **Example**

- single database instance will always be consistent
- if multiple instances exist, all writes must be duplicated before write operation is completed

NoSQL: Concepts

CAP Theorem: Consistency, Availability, Partition Tolerance^[Brewer]

- **Availability**
 - all requests will be answered, regardless of crashes or downtimes
- **Example**
 - a single instance has an availability of 100% or 0%, two servers may be available 100%, 50%, or 0%

NoSQL: Concepts

CAP Theorem: Consistency, Availability, Partition Tolerance^[Brewer]

- **Partition Tolerance**

- system continues to operate, even if two sets of servers get isolated

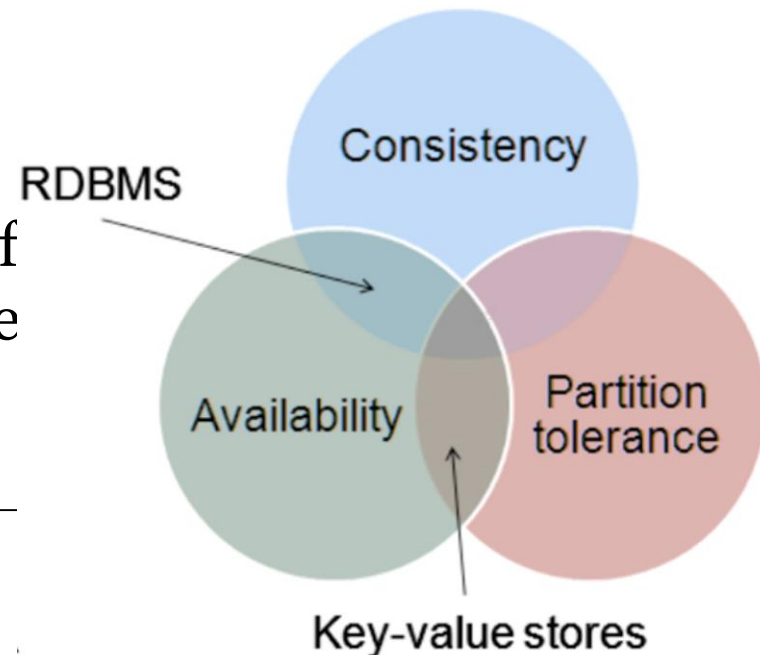
- **Example**

- system gets partitioned if connection between server clusters fails
- failed connection will not cause troubles if system is tolerant

NoSQL: Concepts

CAP Theorem: Consistency, Availability, Partition Tolerance^[Brewer]

- (Positive) consequence: we can concentrate on two challenges
- **ACID** properties needed to guarantee consistency and availability
- **BASE** properties come into play if availability and partition tolerance is favored



NoSQL: Concepts

ACID: Atomicity, Consistency, Isolation, Durability

- **A**tomicity → all operations in a transaction will complete, or none will
- **C**onsistency → before and after the transaction, the database will be in a consistent state
- **I**solation → operations cannot access data that is currently modified
- **D**urability → data will not be lost upon completion of a transaction

NoSQL: Concepts

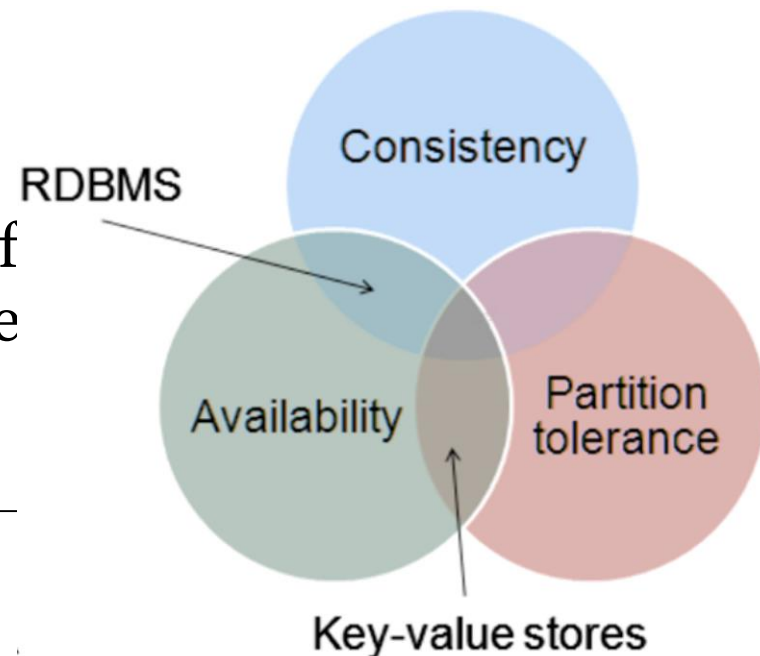
BASE: Basically Available, Soft State, Eventual Consistency [Fox]

- **Basically Available** → an application works basically all the time (despite partial failures)
- **Soft State** → is in flux and non-deterministic (changes all the time)
- **Eventual Consistency** → will be in some consistent state (at some time in future)

NoSQL: Concepts

CAP Theorem: Consistency, Availability, Partition Tolerance^[Brewer]

- (Positive) consequence: we can concentrate on two challenges
- **ACID** properties needed to guarantee consistency and availability
- **BASE** properties come into play if availability and partition tolerance is favored



NoSQL: Techniques

Basic techniques (widely applied in NoSQL systems)

- distributed data storage, replication (how to distribute the data) → Consistent hashing
- distributed query strategy (horizontal scalability) → MapReduce (in the MapReduce lecture)
- recognize order of distributed events and potential conflicts → Vector clock (later in this lecture)

NoSQL: Techniques – Consistent Hashing [Karger]

Task

- find machine that stores data for a specified key k
- trivial hash function to distribute data on n nodes:
 $h(k; n) = k \bmod n$
- if number of nodes changes, all data will have to be redistributed!

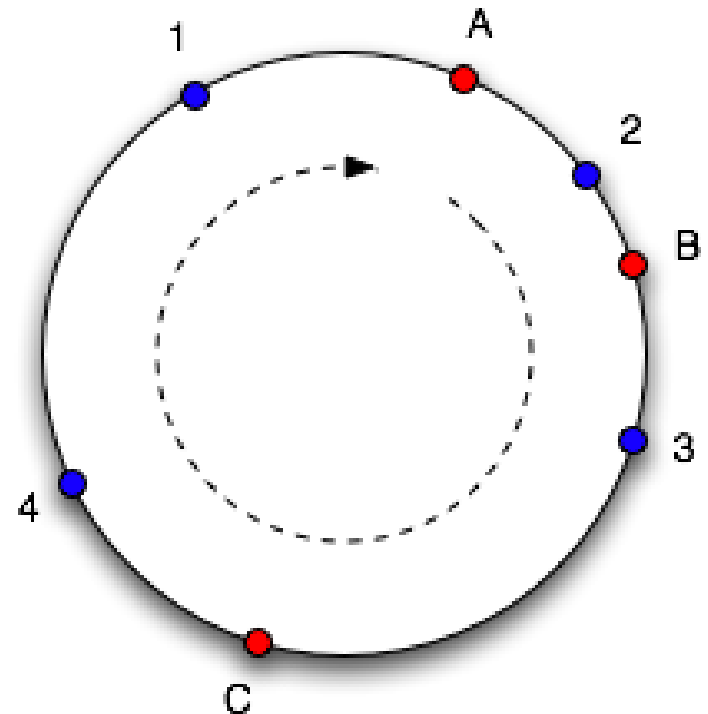
Challenge

- minimize number of nodes to be copied after a configuration change
- incorporate hardware characteristics into hashing model

NoSQL: Techniques – Consistent Hashing [Karger]

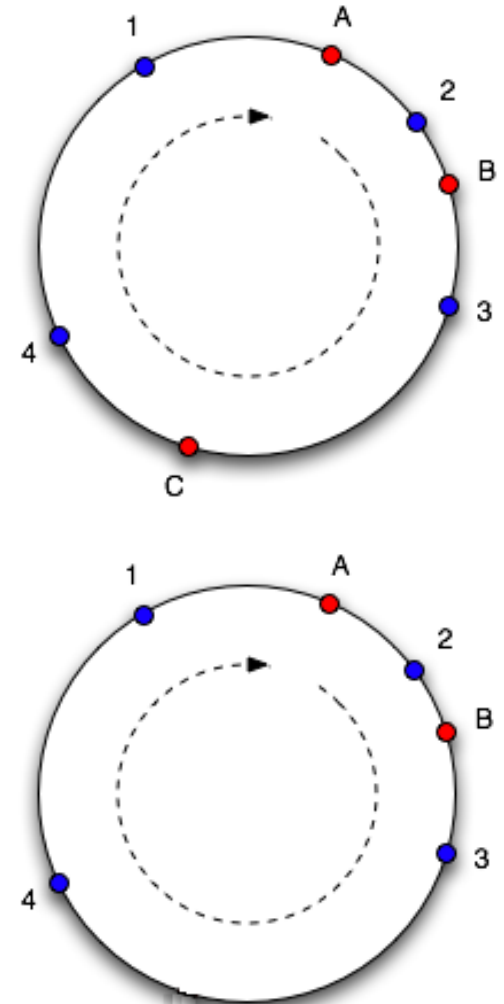
Basic idea

- arrange the nodes in a ring and each node is in charge of the hash values in the range between its neighbor node
- include hash values of all nodes in hash structure
- calculate hash value of the key to be added/retrieved
- choose node which occurs next clockwise in the ring



NoSQL: Techniques – Consistent Hashing [Karger]

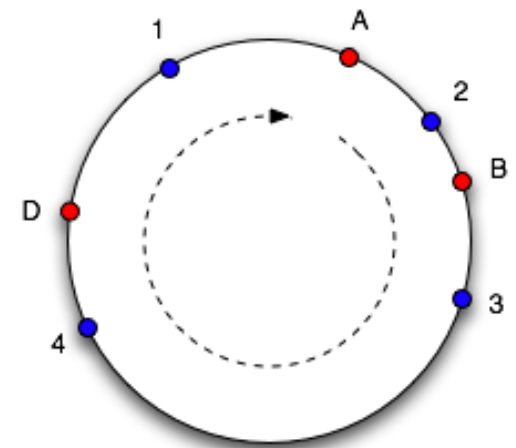
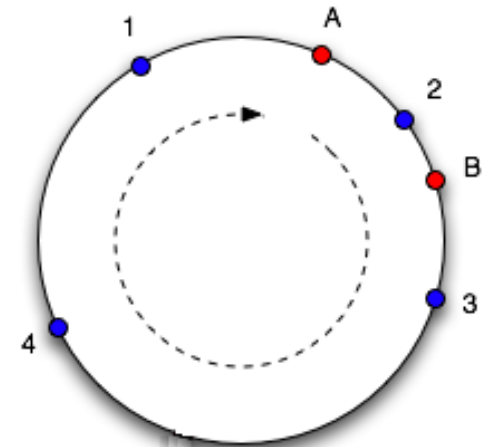
- include hash values of all nodes in hash structure
- calculate hash value of the key to be added/retrieved
- choose node which occurs next clockwise in the ring
- if node is dropped or gets lost, missing data is redistributed to adjacent nodes (replication issue)



NoSQL: Techniques – Consistent Hashing [Karger]

- if a new node is added, its hash value is added to the hash table
- the hash realm is repartitioned, and hash data will be transferred to new neighbor

→ no need to update remaining nodes!



NoSQL: Techniques – Consistent Hashing [Karger]

- A replication factor r is introduced: not only the next node but the next r nodes in clockwise direction become responsible for a key
- Number of added keys can be made dependent on node characteristics (bandwidth, CPU, ...)

NoSQL: Techniques – Logical Time

Challenge

- recognize order of distributed events and potential conflicts
- most obvious approach: attach timestamp (ts) of system clock to each

event $e \rightarrow ts(e)$

→ error-prone, as clocks will never be fully synchronized

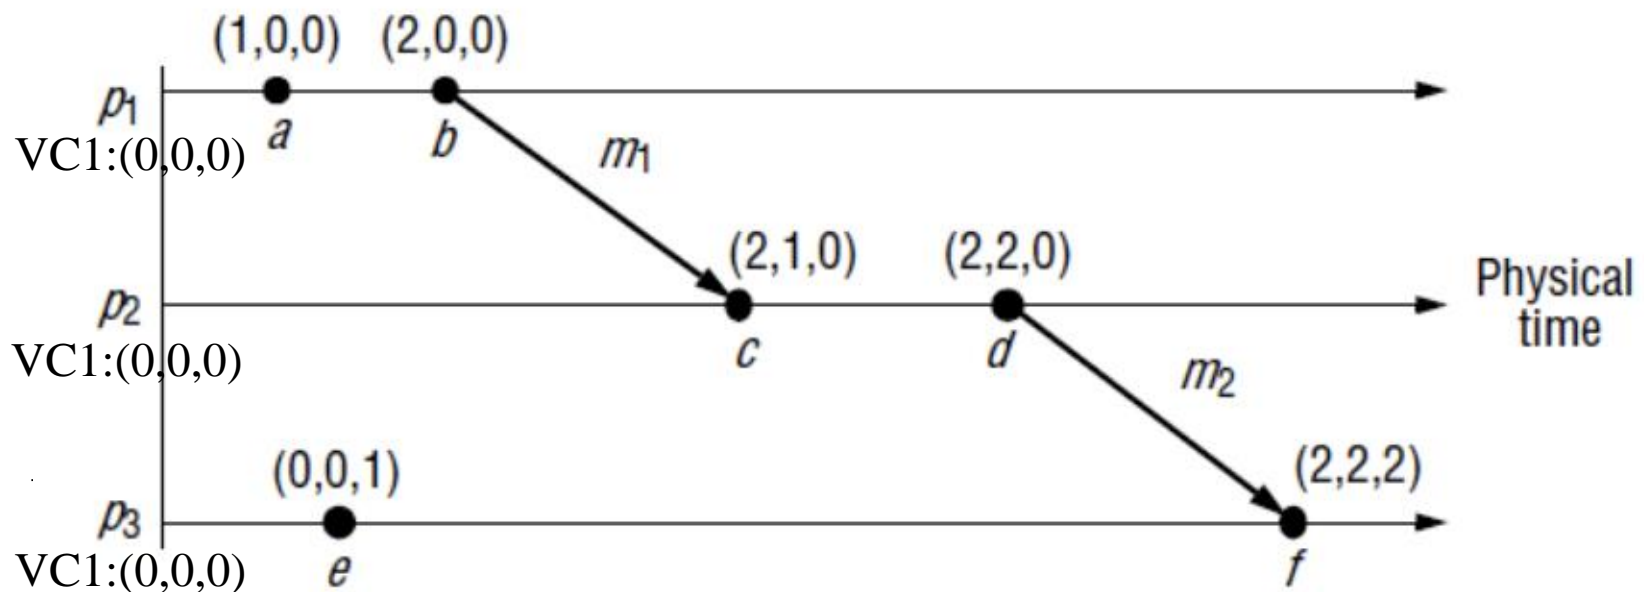
→ insufficient, as we cannot catch causalities (needed to detect conflicts)

NoSQL: Techniques – Vector Clock^[Coulouris]

- A vector clock for a system of N nodes is an array of N integers.
- Each process keeps its own vector clock, V_i , which it uses to timestamp local events.
- Processes piggyback vector timestamps on the messages they send to one another, and there are simple rules for updating the clocks:
 - VC1: Initially, $V_i[j] = 0$, for $i, j = 1, 2, \dots, N$
 - VC2: Just before p_i timestamps an event, it sets $V_i[i] := V_i[i] + 1$
 - VC3: p_i includes the value $t = V_i$ in every message it sends
 - VC4: When p_i receives a timestamp t in a message, it sets $V_i[j] := \max(V_i[j]; t[j])$, for $j = 1, 2, \dots, N$

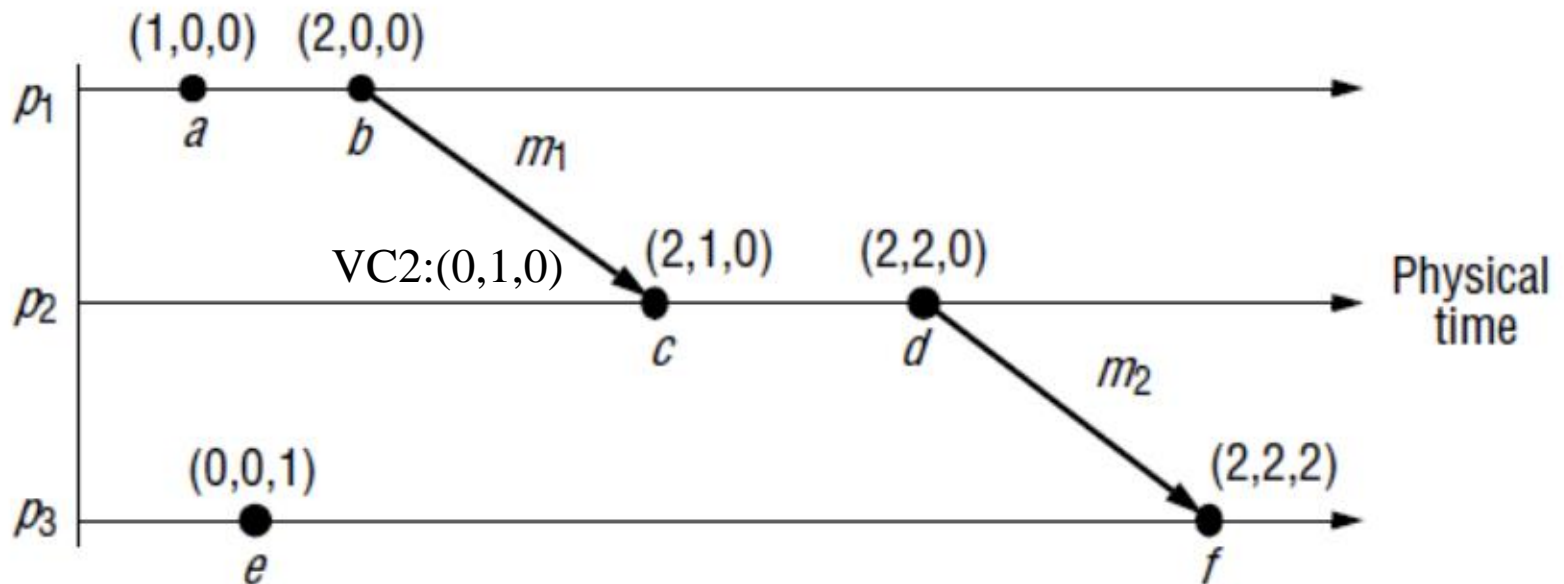
NoSQL: Techniques – Vector Clock^[Coulouris]

- VC1: Initially, $V_i[j] = 0$, for $i, j = 1, 2, \dots, N$
- VC2: Just before p_i timestamps an event, it sets $V_i[i] := V_i[i] + 1$



NoSQL: Techniques – Vector Clock^[Coulouris]

- VC3: p_i includes the value $t = V_i$ in every message it sends
- VC4: When p_i receives a timestamp t in a message, it sets $V_i[j] := \max(V_i[j]; t[j])$, for $j = 1, 2, \dots, N$



NoSQL: Techniques – Vector Clock^[Coulouris]

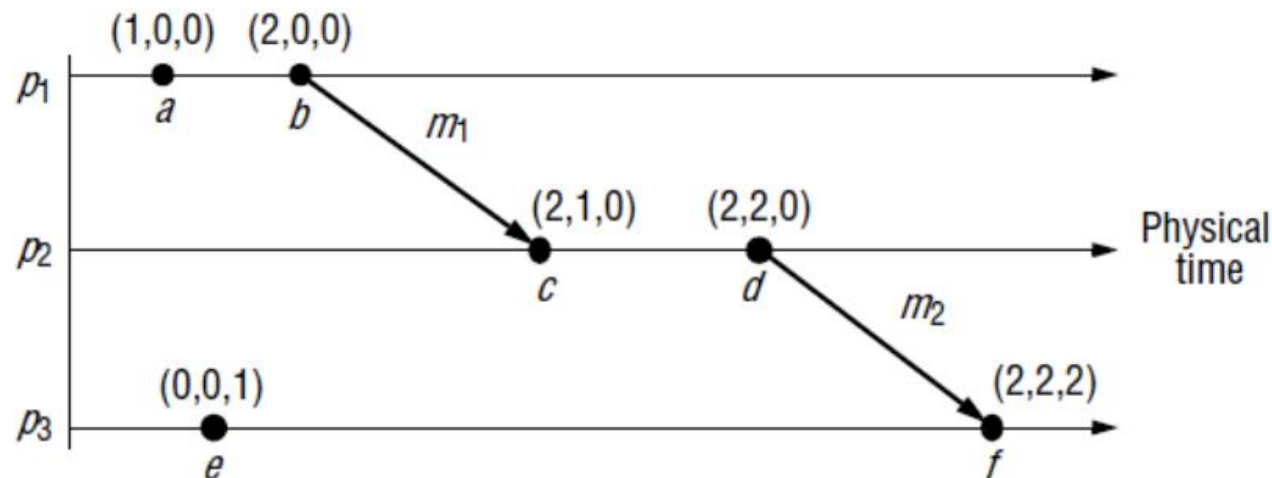
Properties:

- $V = V'$ iff $V[j] = V'[j]$ for $j = 1, 2, \dots, N$
- $V \leq V'$ iff $V[j] \leq V'[j]$ for $j = 1, 2, \dots, N$
- $V < V'$ iff $V \leq V'$ and $V \neq V'$

two events e and e' : that $e \rightarrow e' \leftrightarrow V(e) < V(e')$

→ Conflict detection! ($c \parallel e$ since neither $V(c) \leq V(e)$ nor $V(e) \leq V(c)$)

c & e are concurrent



NoSQL Systems – Types and Applications

NoSQL Classification Dimensions^[HBase]

- Data model – how the data is stored
- Storage model – in-memory vs persistent
- Consistency model – strict, eventual consistent, etc.
 - Affects reads and writes requests
- Physical model – distributed vs single machine
- Read/Write performance – what is the proportion between reads and writes
- Secondary indexes - sort and access tables based on different fields and sorting orders

NoSQL Classification Dimensions^[HBase]

- Failure handling – how to address machine failures
- Compression – result in substantial savings in raw storage
- Load balancing – how to address high read or write rate
- Atomic read-modify-write – difficult to achieve in a distributed system
- Locking, waits and deadlocks – locking models and version control

NoSQL Data Models

- Key-Value Stores
 - Document Stores
 - Column-Family Stores
 - Graph Databases
-
- Impacts application, querying, scalability

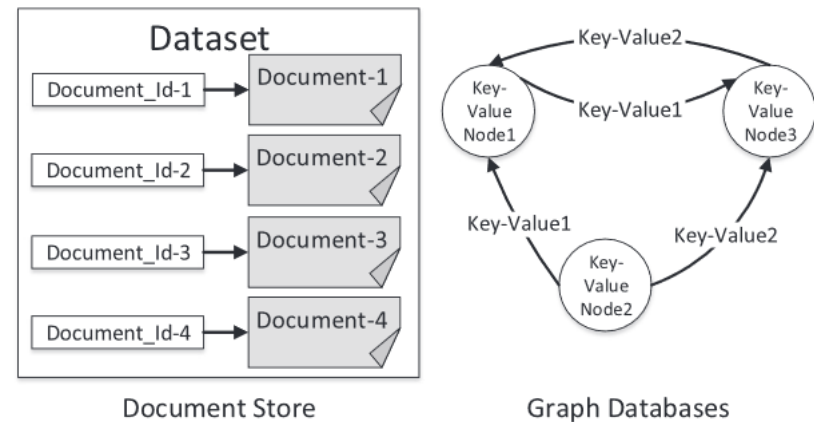
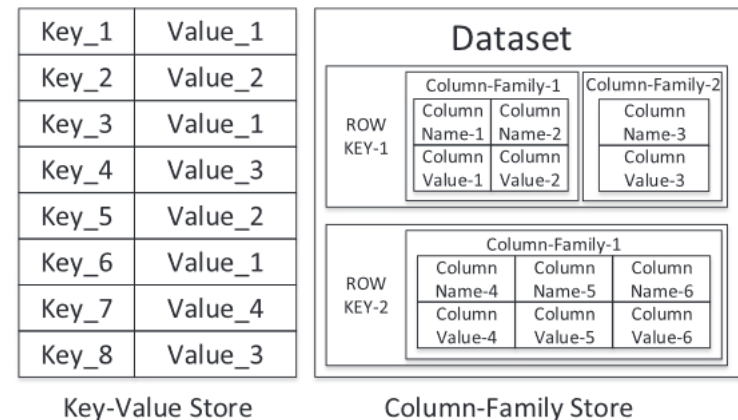


figure from [DataMan]

DBs not referred as NoSQL

- Object DBs
- XML DBs
- Special purpose DBs
 - Stream processing

Key-Value Stores^[DataMan]

- Schema-free
 - Keys are unique
 - Values of arbitrary types
- Efficient in storing distributed data
- (very) Limited query facilities and indexing
 - get(key), put(key, value)
 - Value → opaque to the data store → no data level querying and indexing

| | |
|-------|---------|
| Key_1 | Value_1 |
| Key_2 | Value_2 |
| Key_3 | Value_1 |
| Key_4 | Value_3 |
| Key_5 | Value_2 |
| Key_6 | Value_1 |
| Key_7 | Value_4 |
| Key_8 | Value_3 |

Key-Value Store

Key-Value Stores^[DataMan]

- Types
 - In-memory stores – Memcached, Redis
 - Persistent stores – BerkeleyDB, Voldemort, RiakDB

| | |
|-------|---------|
| Key_1 | Value_1 |
| Key_2 | Value_2 |
| Key_3 | Value_1 |
| Key_4 | Value_3 |
| Key_5 | Value_2 |
| Key_6 | Value_1 |
| Key_7 | Value_4 |
| Key_8 | Value_3 |

Key-Value Store

- Not suitable for
 - structures and relations
 - accessing multiple items (since the access is by key and often no transactional capabilities)

Key-Value Stores^[DataMan]

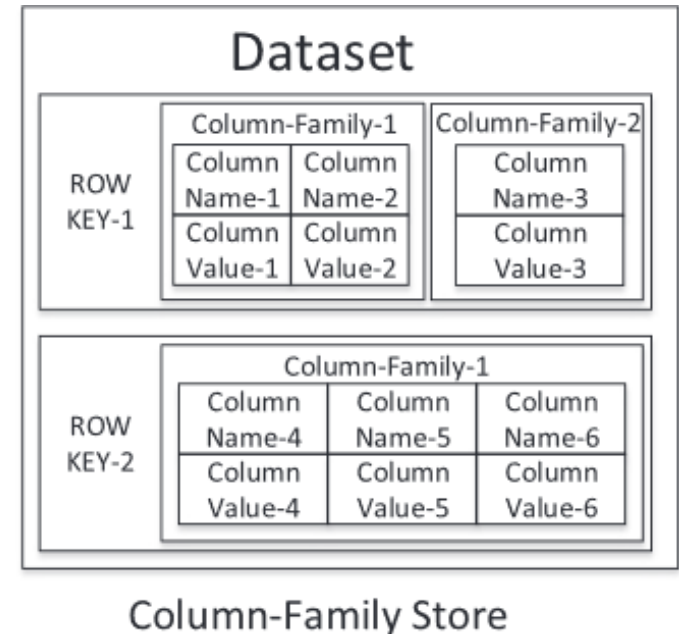
- Applications:
 - Storing web session information
 - User profiles and configuration
 - Shopping cart data
 - Using them as a caching layer to store results of expensive operations (create a user-tailored web page)

| | |
|-------|---------|
| Key_1 | Value_1 |
| Key_2 | Value_2 |
| Key_3 | Value_1 |
| Key_4 | Value_3 |
| Key_5 | Value_2 |
| Key_6 | Value_1 |
| Key_7 | Value_4 |
| Key_8 | Value_3 |

Key-Value Store

Column-Family Stores^[DataMan]

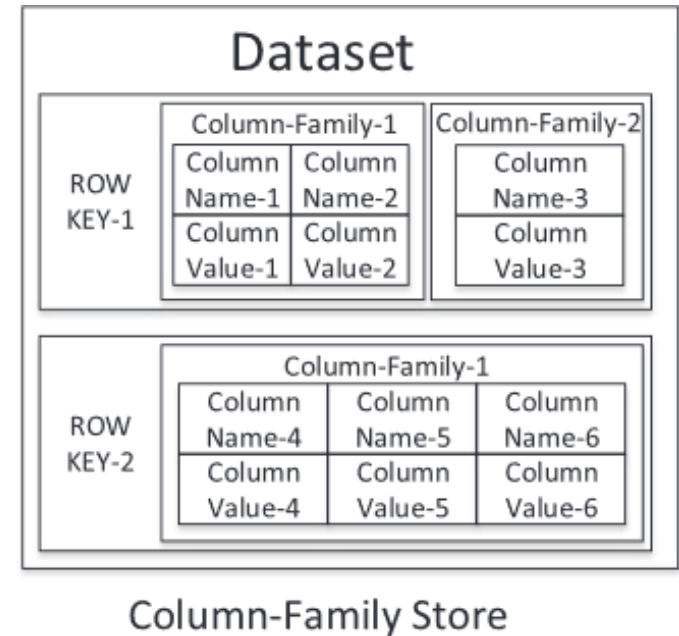
- Schema-free
 - Rows have unique keys
 - Values are varying column families and act as keys for the columns they hold
 - Columns consist of key-value pairs



- Better than key-value stores for querying and indexing

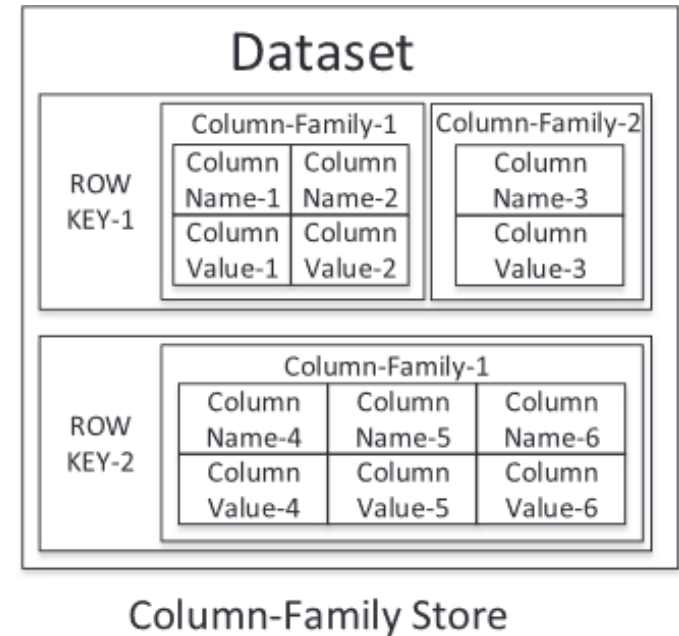
Column-Family Stores^[DataMan]

- Types
 - Googles BigTable, Hadoop HBase
 - No column families – Amazon SimpleDB, DynamoDB
 - Supercolumns - Cassandra
- Not suitable for
 - structures and relations
 - highly dynamic queries (HBase and Cassandra)



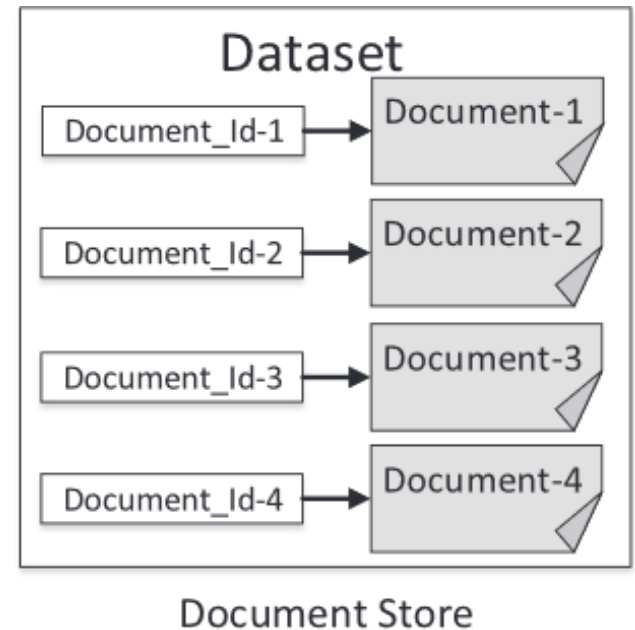
Column-Family Stores^[DataMan]

- Applications:
 - Document stores applications
 - Analytics scenarios – HBase and Cassandra
 - Web analytics
 - Personalized search
 - Inbox search



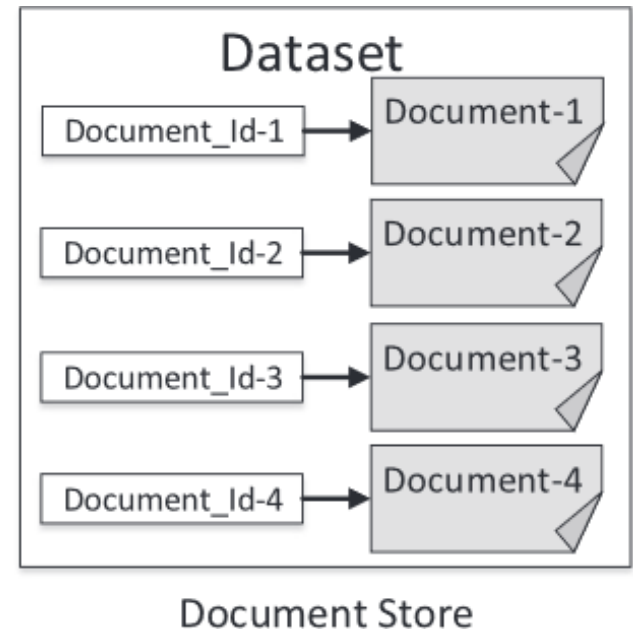
Document Stores^[DataMan]

- Schema-free
 - Keys are unique
 - Values are documents – complex (nested) data structures in JSON, XML, binary (BSON), etc.
- Indexing and querying based on primary key and content
- The content needs to be representable as a document
- MongoDB, CouchDB, Couchbase



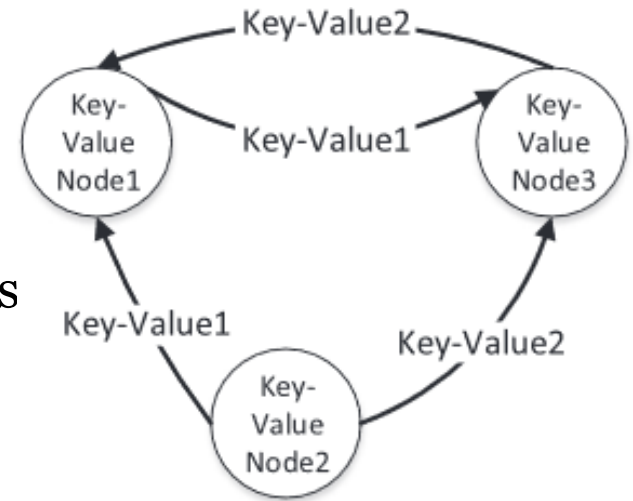
Document Stores^[DataMan]

- Applications:
 - Items with similar nature but different structure
 - Blogging platforms
 - Content management systems
 - Event logging
 - Fast application development



Graph Databases^[DataMan]

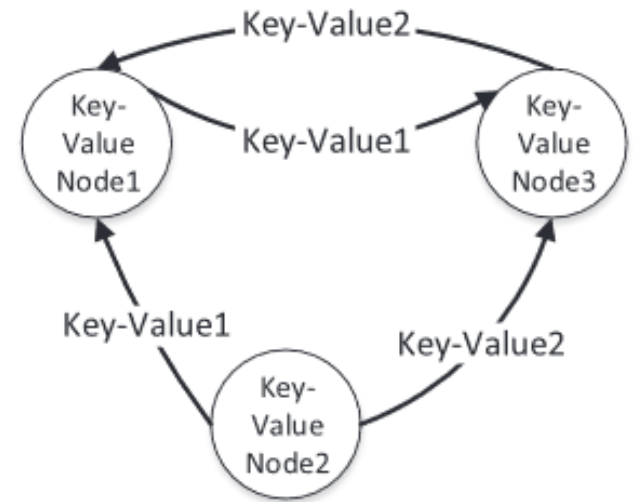
- Graph model
 - Nodes/vertices and links/edges
 - Properties consisting of key-value pairs
- Suitable for very interconnected data since they are efficient in traversing relationships
- Not as efficient
 - as other NoSQL solutions for non-graph applications
 - horizontal scaling
- Neo4J, HyperGraphDB



Graph Databases

Graph Databases^[DataMan]

- Applications:
 - location-based services
 - recommendation engines
 - complex network-based applications
 - social, information, technological, and biological network
 - memory leak detection



Graph Databases

Big Data Analytics Stack

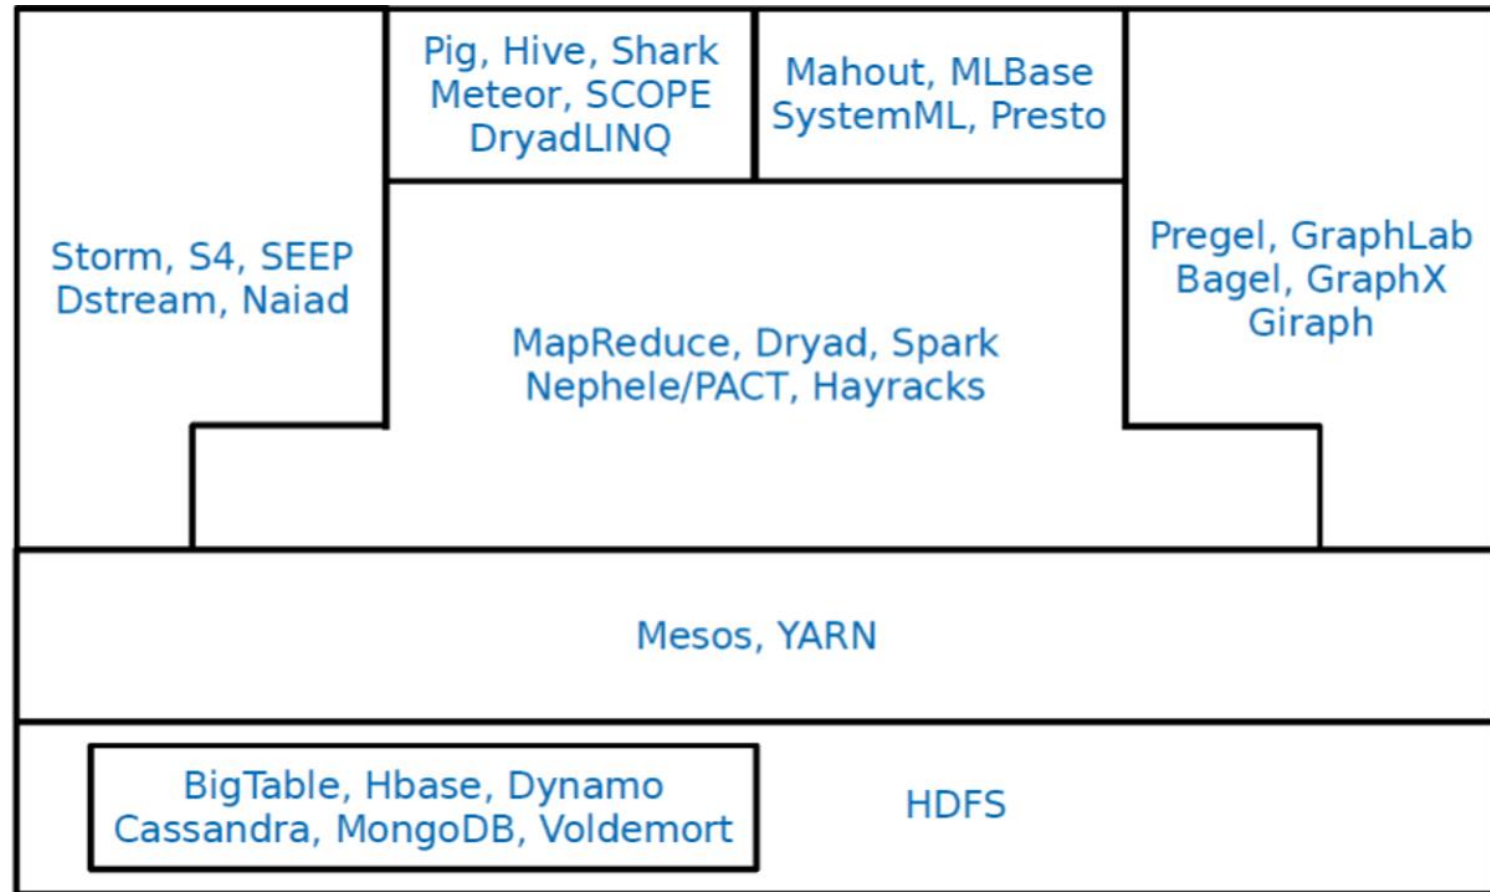


figure from: <https://www.sics.se/~amir/dic.htm>

HDFS^{[Hadoop][HDFS][HDFSpaper]} Hadoop Distributed File System



Compute Nodes^[Massive]

- Compute node – processor, main memory, cache and local disk
- Organized into racks
- Intra-rack connection typically gigabit speed
- Inter-rack connection slower by a small factor

HDFS (Hadoop Distributed File System)

- Runs on top of the native file system
 - Files are very large divided into 128 MB chunks/blocks
 - To minimize the cost of seeks
 - Caching blocks is possible
 - Single writer, multiple readers
 - Exposes the locations of file blocks via API
 - Fault tolerance and availability to address disk/node failures
 - Usually replicated three times on different compute nodes
- Based on GFS (Google File System - proprietary)

HDFS is Good for ...

- Store very large files – GBs and TBs
- Streaming access
 - Write-once, read many times
 - Time to read the entire dataset is more important than the latency in reading the first record.
- Commodity hardware
 - Clusters are built from commonly available hardware
 - Designed to continue working without a noticeable interruption in case of failure

HDFS is currently Not Good for ...

- Low-latency data access
 - HDFS is optimized for delivering high throughput of data
- Lots of small files
 - the amount of files is limited by the memory of the namenode; blocks location is stored in memory
- Multiple writers and arbitrary file modifications
 - HDFS files are append only – write always at the end of the file

HDFS Organization

- Namenode (master)
 - Manages the filesystem namespace and metadata
 - Stores in memory the location of all blocks for a given file
- Datanodes (workers)
 - Store and retrieve blocks
 - Send heartbeat to the namenode
- Secondary namenode
 - Periodically merges the namespace image with the edit log
 - **Not** a backup for a namenode, only a checkpoint

HDFS Organization

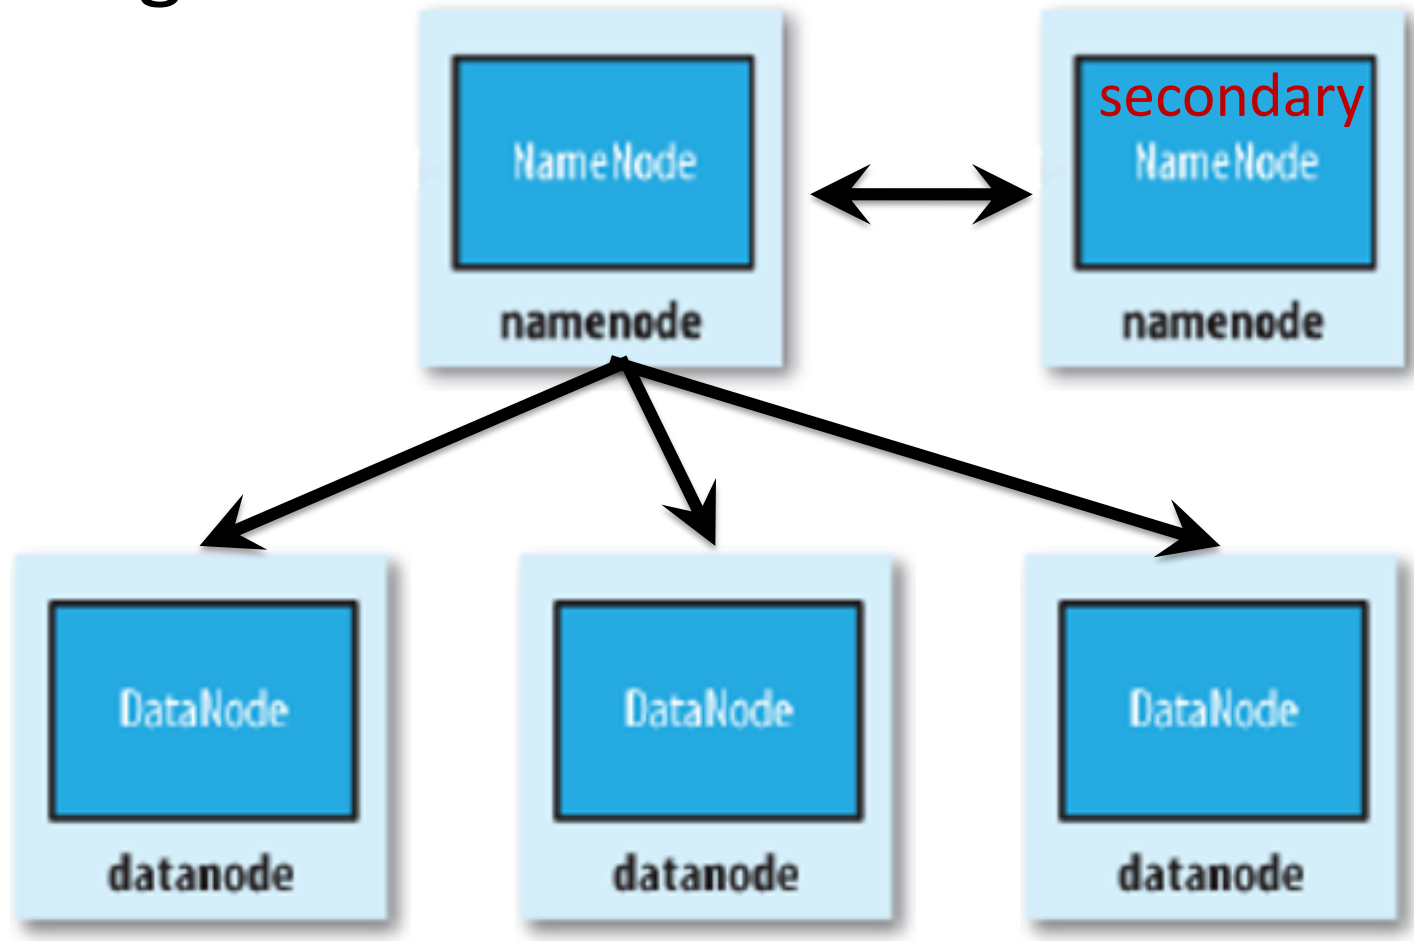
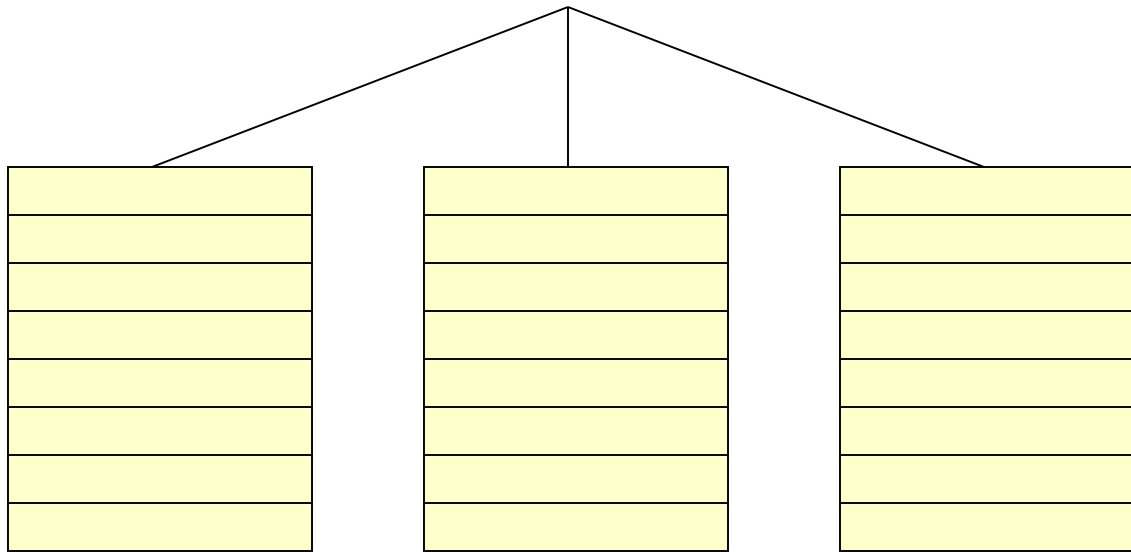


figure based on a figure from [Hadoop]

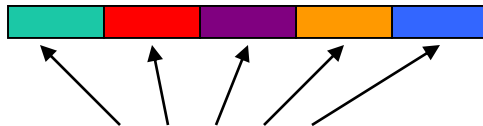
Block Placement and Replication

- Aim – improve data reliability, availability and network bandwidth utilization
- Default replica placement policy
 - No Datanode contains more than one replica
 - No rack contains more than two replicas of the same block
- Namenode ensures the number of replicas is reached
- Balancer tool – balances the disk space usage
- Block scanner – periodically verifies checksums



Racks of Compute Nodes

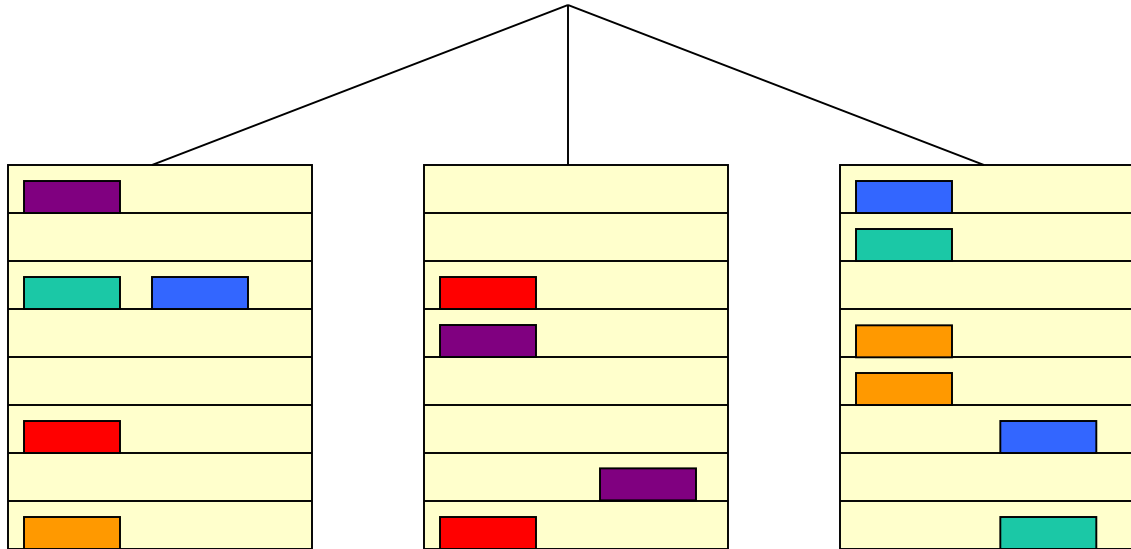
File



Chunks

Source: J. D. Ullman invited talk EDBT 2011

Default HDFS Block Placement Policy



- 1st replica located on the writer node
- 2nd and 3rd replicas on two different nodes in a different rack
- The other replicas are located on random nodes

HDFS – File Reads

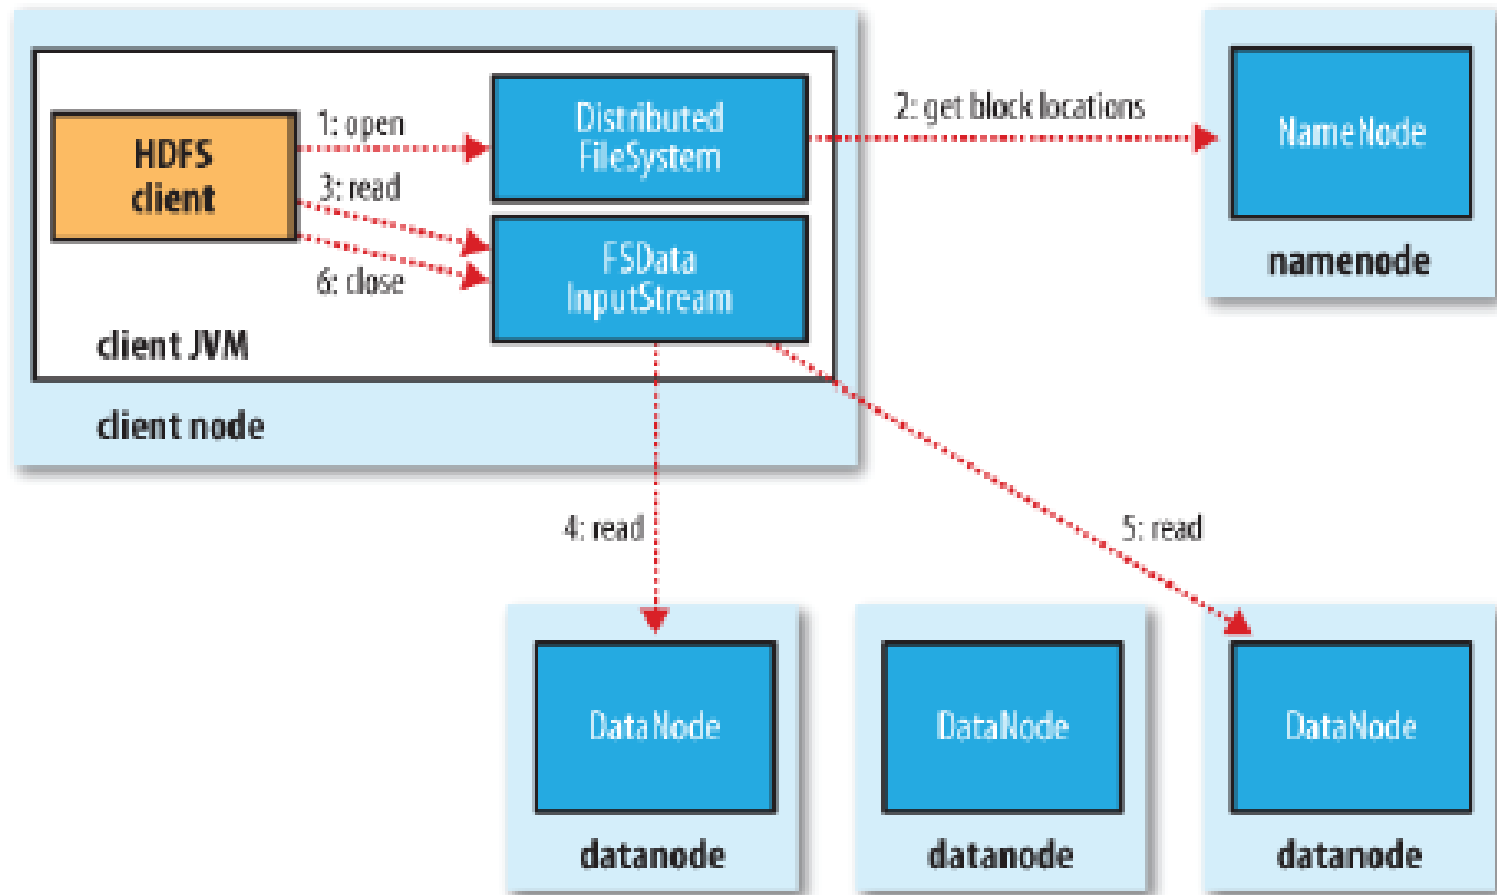


figure from [Hadoop]

HDFS – File Writes

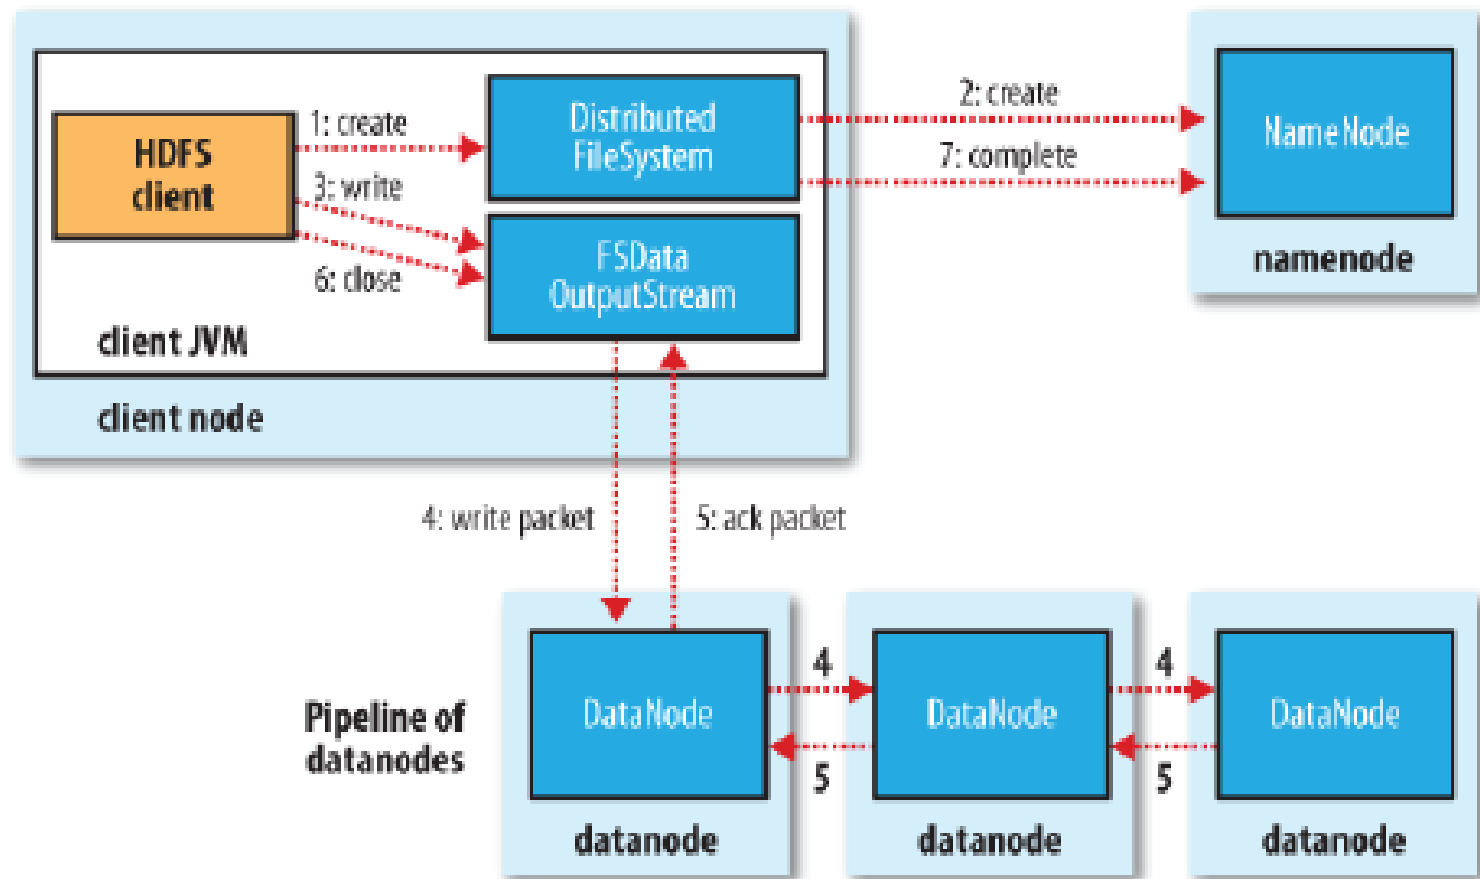


figure from [Hadoop]

HDFS – High Availability

- The namenode is single point of failure:
 - If a namenode crashes the cluster is down
- Secondary node
 - periodically merges the namespace image with the edit log to prevent the edit log from becoming too large.
 - lags the state of the primary prevents data loss but does **not** provide high availability
 - time for cold start 30 minutes
- In practice, the case for planned downtime is more important

HDFS – High Availability

- Pair of namenodes in an active stand-by configuration:
 - Highly available shared storage for the shared edit log
 - Datanodes send block reports to all namenodes
 - Clients must provide transparent to the user mechanism to handle failover
 - The standby node takes checkpoints of the active namenode namespace instead of the secondary node

HDFS commands

- List all options for the hdfs dfs
 - `hdfs dfs -help`
 - `dfs` – run a filesystem command
- Create a new folder
 - `hdfs dfs -mkdir /BigDataAnalytics`
- Upload a file from the local file system to the HDFS
 - `hdfs dfs -put bigdata /BigDataAnalytics`

HDFS commands

- List the files in a folder

- `hdfs dfs -ls /BigDataAnalytics`

- Determine the size of a file

- `hdfs dfs -du -h /BigDataAnalytics/bigdata`

- Print the first 5 lines from a file

- `hdfs dfs -cat /BigDataAnalytics/bigdata | head -n 5`

- Copy a file to another folder

- `hdfs dfs -cp /BigDataAnalytics/bigdata /BigDataAnalytics/AnotherFolder`

HDFS commands

- Copy a file to a local filesystem and rename it
 - `hdfs dfs -get /BigDataAnalytics/bigdata bigdata_localcopy`
- Scan the entire HDFS for problems
 - `hdfs fsck /`
- Delete a file from HDFS
 - `hdfs dfs -rm /BigDataAnalytics/bigdata`
- Delete a folder from HDFS
 - `hdfs dfs -rm -r /BigDataAnalytics`

References

- A comparison between several NoSQL databases with comments and notes by Bogdan George Tudorica, Cristian Bucur
 - nosql-databases.org
 - Scalable SQL and NoSQL data stores by Rick Cattell
 - [Brewer] Towards Robust Distributed Systems @ACM PODC'2000
 - [12 years later] CAP Twelve Years Later: How the "Rules" Have Changed, Eric A. Brewer, @Computer Magazine 2012. <https://www.infoq.com/articles/cap-twelve-years-later-how-the-rules-have-changed>
 - [Fox et al.] Cluster-Based Scalable Network Services @SOSP'1997
 - [Karger et al.] Consistent Hashing and Random Trees @ACM STOC'1997
 - [Coulouris et al.] Distributed Systems: Concepts and Design, Chapter: Time & Global States, 5th Edition
 - [DataMan] Data Management in cloud environments: NoSQL and NewSQL data stores.
-

References

- NoSQL Databases - Christof Strauch – University of Stuttgart
- The Beckman Report on Database Research
- [Vogels] Eventually Consistent by Werner Vogels, doi:10.1145/1435417.1435432
- [Hadoop] Hadoop The Definitive Guide, Tom White, 2011
- [Hive] Hive - a petabyte scale data warehouse using Hadoop
- <https://github.com/Prokopp/the-free-hive-book>
- [Massive] Mining of Massive Datasets
- [HiveManual]
<https://cwiki.apache.org/confluence/display/Hive/LanguageManual>
- [Shark] Shark: SQL and Rich Analytics at Scale
- [SparkSQLHistory] <https://databricks.com/blog/2014/07/01/shark-spark-sql-hive-on-spark-and-the-future-of-sql-on-spark.html>

References

- [HDFS] The Hadoop Distributed File System
- [Dynamo] Dynamo: Amazon's Highly Available Key-value Store, 2007
- [HBaseInFacebook] Apache hadoop goes realtime at Facebook
- [HBase] HBase The Definitive Guide, 2011
- [HDFSpaper] The Hadoop Distributed File System @MSST2010