

Neural Networks and Learning Systems
TBM126 / 732A55
2018

Lecture 7
Reinforcement Learning

Magnus Borga
magnus.borga@liu.se




Article

Mastering the game of Go without human knowledge

David Silver , Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel & Demis Hassabis

[David Silver about Alpha Go Zero](#)

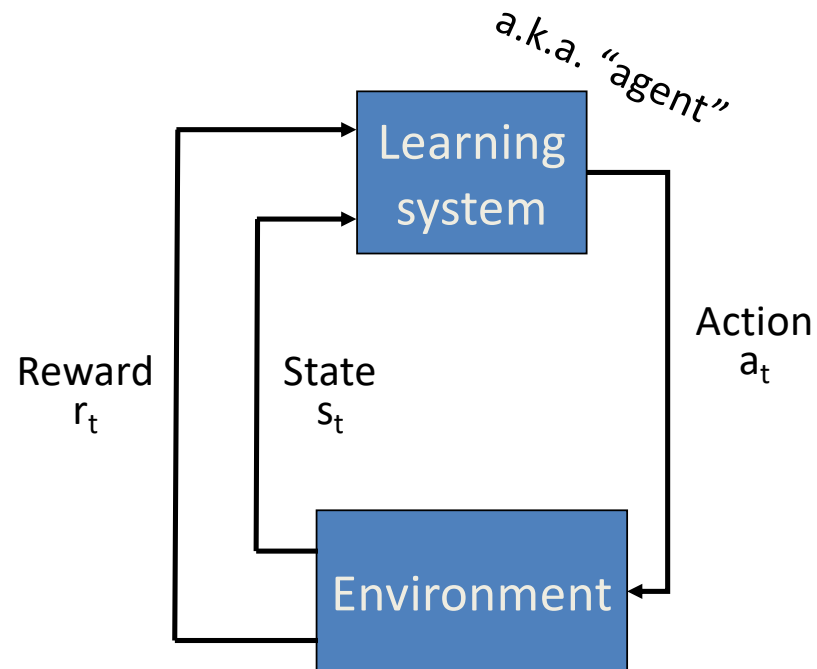
Go



Board size n×n	3^{n^2}	Percent legal	L (legal positions)
1×1	3	33.33%	1
2×2	81	70.37%	57
3×3	19,683	64.40%	12,675
4×4	43,046,721	56.49%	24,318,165
5×5	847,288,609,443	48.90%	414,295,148,741
9×9	$4.43426488243 \times 10^{38}$	23.44%	$1.03919148791 \times 10^{38}$
13×13	$4.30023359390 \times 10^{80}$	8.66%	$3.72497923077 \times 10^{79}$
19×19	$1.74089650659 \times 10^{172}$	1.20%	$2.08168199382 \times 10^{170}$


Reinforcement learning

Learn by interacting with the environment!



Examples

Board games




State: The board position.

Action: Placing a stone in a valid location.

Reward: At the end of the game (win/loss).

Exploring maps




<http://rumpus.rubyforge.org/>

State: Current location (x,y).

Action: Move in a valid direction.

Reward: When the way between point A and point B is found. Negative reward for each move that is made.

Balancing a pole






State: $x, \dot{x}, \theta, \dot{\theta}$

Action: Apply force F.

Reward: Negative reward if the pole falls.


Differences to other methods

- Difference to supervised learning
 - Time!
 - Feedback is given as a scalar reward, not as the correct action to make.
 - Feedback is usually not immediate but is given after many actions – delayed feedback!
 - Can become better than the system designer, unlike a supervised system that can never become better than the teacher.
- Difference to control theory
 - No physical model of the world, e.g., in pole balancing




Discretize

Discretize state-space and time!



Markov Decision Process – Stochastic state transitions




Policy

- Defines which action to take in each state
- Can be represented with a look-up table:

State:	s_1	s_2	s_3	s_4
Action:	a_4	a_1	a_{10}	a_7	

Reward

$r(\mathbf{x}, \mathbf{a})$ – the reward for making action \mathbf{a} in state \mathbf{x} .






Reinforcement learning goal

- The goal in reinforcement learning is to find an optimal policy, i.e., state-action pairs that maximize the reward.
- To do this, we have to solve the following problems:
 1. How to evaluate how good a policy is?
 2. Which policies should we explore?

Value function

How good is a policy?



A function $V(s)$ of the state that tells us the value of being in the state given a policy, i.e., the expected amount of reward we get from this state by following the policy:

$$\text{Alt 1: } V(s_t) = \sum_{k=0}^{\infty} r_{t+k}$$


In general, a random variable

$$\text{Alt 2: } V(s_t) = \sum_{k=0}^{\infty} \gamma^k r_{t+k} \text{ where } 0 \leq \gamma \leq 1$$

Makes immediate rewards more important than distant rewards

Accumulated reward


$V(\mathbf{x})$




$\gamma = 0.9$

A non-stochastic example

Example 1: Bad policy



Example 2: Optimal policy




Value function, cont.

How good is a policy?

Note 1: The value function tells us how good a policy is. The optimal policy has maximum $V(s)$ for all states.


Note 2: When we start learning we explore different policies and we have to learn $V(s)$ for each policy as it is unknown before we start exploring the environment.

Note 3: When programming intelligence using other approaches, the value function is usually specified by the designer, i.e., not learned by the agent.



The Credit Assignment Problem

How good is a policy? - How to learn $V(s)$?




The big question:

How to distribute the reward in the chain of states (and actions) that led to the reward?

For example: Was there a "genius" move that led to the win?

The Credit Assignment Problem

How good is a policy? - How to learn $V(s)$?




Possible solution:

A state is awarded reward according to how many transitions m from the terminal state it is, i.e., with a discount factor γ^m where $0 < \gamma < 1$.

A Monte Carlo approach

How good is a policy? - How to learn $V(s)$?



- Generate MANY episodes, possibly starting from different states.
- The state-space chains will generally be different because of the stochastic state transitions.
- After a reward r , update the value functions for the visited states:

$$\hat{V}(s_k) \leftarrow (1 - \alpha) \hat{V}(s_k) + \alpha \gamma^m r$$



Learning rate $0 < \alpha < 1$ (points to α)

Discount factor $0 < \gamma < 1$ (points to γ)

states from reward (points to m)

Cliff walk example

Windy: 25% chance of stumbling up or down




$r = -100$

$r = +100$




Actions:
up, down, right


Policy 1



Policy 2



Cliff walk - Monte Carlo evaluation




\hat{V}_0

0	0	0	0	0
0	0	0	0	0
0	Cliff			G

$$\alpha = 0.1$$

$$\gamma = 1$$


$$\hat{V}(s_k) \leftarrow (1 - 0.1)\hat{V}(s_k) + 0.1r$$



\hat{V}_1

0	0	0	0	0
-10	-10	0	0	0
-10	Cliff			G

$$r = -100$$



\hat{V}_2

0	0	0	0	0
1	1	10	10	10
1	Cliff			G

$$r = +100$$

Cliff walk – Monte Carlo evaluation

α decreasing towards 0

Policy 1

→	→	→	→	↓
→	→	→	→	↓
S ↑	Cliff			G

$\hat{V}_{100,000}$

71	81	89	96	100
40	40	56	75	100
40	Cliff			G

Policy 2

→	→	→	→	↓
↑	↑	↑	↑	↓
S ↑	Cliff			G


$\hat{V}_{100,000}$

88	88	92	96	100
88	65	68	73	100
88	Cliff			G

Temporal Difference approach

How good is a policy? - How to learn $V(s)$?


Update $V(s)$ after each state transition!



$$\hat{V}(s_k) \leftarrow (1 - \alpha)\hat{V}(s_k) + \alpha(r_k + \gamma \hat{V}(s_{k+1}))$$

1. $V(s_k)$ is the expected reward in state s_k .
2. $r_k + \gamma V(s_{k+1})$ is a more accurate estimate of the exp. reward in s_k because we have seen r_k , and $V(s_{k+1})$ is closer to the end.
3. Update $V(s_k)$ with a weighted average using α !

Reinforcement learning



How to learn $V(s)$ – Summary

How good is a policy?

- For a given policy, the value (expected reward) $V(s)$ of each state is unknown before we learn it by interacting with the environment.
- $V(s)$ is found iteratively, starting for example with $V(s) \leftarrow 0$, using the Monte Carlo or Temporal Difference methods.
- The Temporal Difference method generally converges much faster.

But our goal is to find the best policy!

- Brute force – Test all possible policies and choose the one with best value function.
 - Only possible for very small toy problems
- Focus the search more on policies that seem promising, i.e., variations of policies that have already been found to give good value functions.
 - Exploration-Exploitation Dilemma

Exploration-Exploitation Dilemma

How much should we explore?

- How much should we explore new policies and how much should we exploit what we already learned?
- Example: Multi-armed bandit

Each arm gives a win with a certain probability.

Start by testing both arms to learn which one gives the best win ratio.

When do you stop exploring and start exploiting (pulling only one arm) to maximize your expected winnings?



Q-function – Add an action dimension

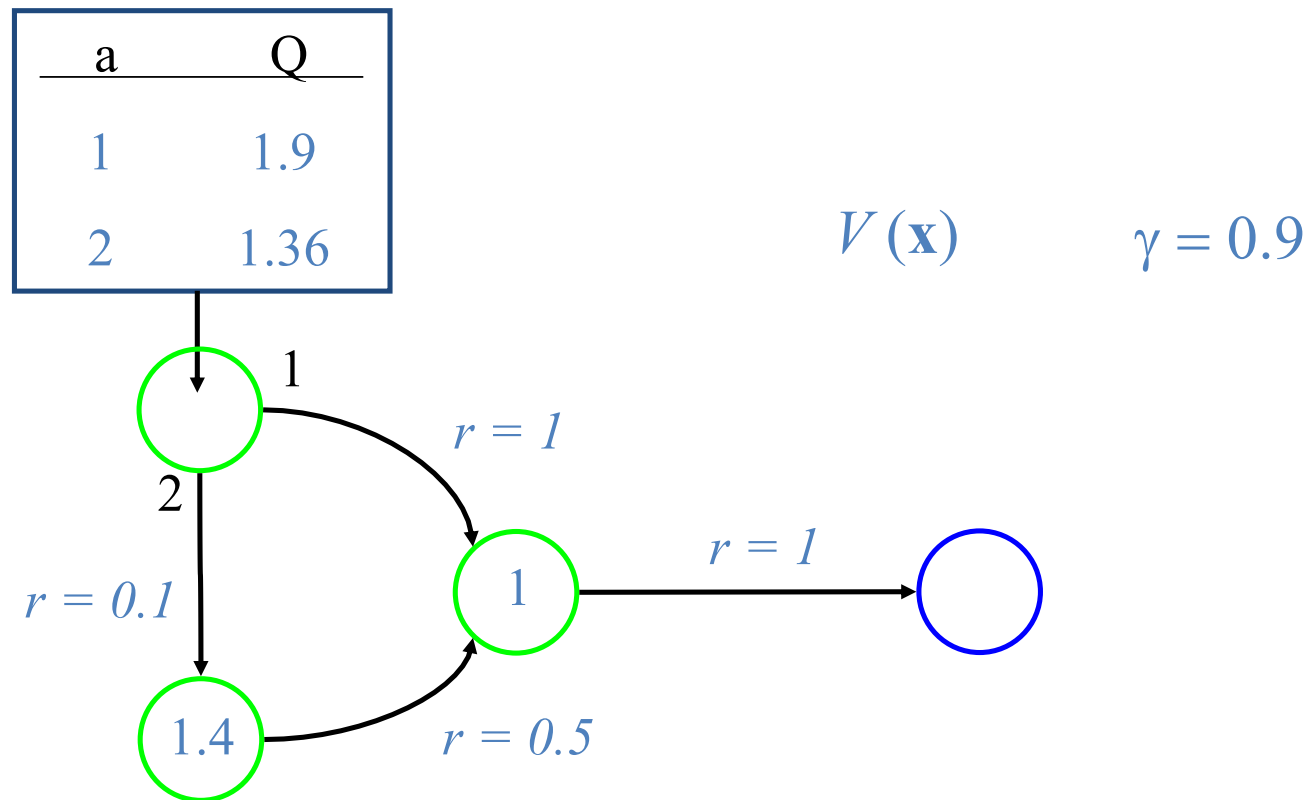
Which policies should we explore?

- Let $V^*(s)$ denote the value function for the optimal policy.
- $Q(s,a)$: expected future reward of doing action a in state s and then following the optimal policy:

$$Q(s_k, a) = r(s_k, a) + \gamma V^*(s_{k+1})$$

- $Q(s,a)$ indirectly encodes the optimal policy and its value function: $V^*(s) = \max_a Q(s,a)$.
- $Q(s,a)$ is unknown – it must be learned!

Q-learning



Why the Q-function?

Which policies should we explore?

- Seems that we have just made the problem more complex!?

Added an extra dimension
in the look-up table!


$Q(s,a)$	s_1	s_2	s_3	...
a_1	0.5	0.9	-	...
a_2	1.2	-	0.3	...
...

The Q-function is an instrument for experimentation around the best policies during learning.

Q-Learning

Explores policies and value functions simultaneously!

Learn the Q-function using an iteration
similar to the Temporal Difference update!



$$\hat{Q}(s_k, a_j) \leftarrow (1 - \alpha) \hat{Q}(s_k, a_j) + \alpha \left(r + \gamma \hat{V}(s_{k+1}) \right) =$$
$$(1 - \alpha) \hat{Q}(s_k, a_j) + \alpha \left(r + \gamma \max_a \hat{Q}(s_{k+1}, a) \right)$$

Q-Learning algorithm

Initialize a look-up table for $Q(s,a)$ with random values.

for each episode

 Init a start state s

repeat for each step k in the episode

 Choose an action a_j (see next slides)

 Take action a_j and observe reward r and next state s_{k+1}

 Update estimated Q-function:

$$\hat{Q}(s_k, a_j) \leftarrow (1 - \alpha) \hat{Q}(s_k, a_j) + \alpha \left(r + \gamma \max_a \hat{Q}(s_{k+1}, a) \right)$$

end

end

ϵ -greedy exploration

- Make a random action (explore) with probability ϵ .
- Make a greedy action (exploit) with probability $1-\epsilon$:

$$\arg \max_a \hat{Q}(s, a)$$

- May want to explore more in the beginning (large ϵ) of the training phase and less towards the end.

Parameter summary


$$\hat{Q}(s_k, a_j) \leftarrow (1 - \alpha) \hat{Q}(s_k, a_j) + \alpha \left(r + \gamma \max_a \hat{Q}(s_{k+1}, a) \right)$$

- **$0 < \alpha < 1$**
The learning rate. A value close to 0 puts more emphasis on already learned experience and a value close to 1 will overwrite previous experience with new information. Good value to start with: 0.1-0.5
- **$0 < \gamma < 1$**
The discount factor. A value close to 0 will seek to maximize short-term rewards whereas a value close to 1 will focus the learning on long term rewards. Good value to start with: around 0.9
- **$0 < \epsilon < 1$**
Exploration factor – the probability of choosing a random action. A value close to 1 will make the learning focus on exploration and a value close to 0 will make the system take actions based on already learned experience. Explore a lot in the beginning (large ϵ) and focus the search more around the good policies towards the end (small ϵ).

Summary Q-Learning

- Solves the two problems simultaneously:
 1. How to evaluate how good a policy is?
 2. Which policies should we explore?
- The look-up table quickly becomes (too) big.
- Lots of demos and applets on the internet.

Method summary



Reinforcement Learning - Summary

- Learning by trial-and-error
- Often hard to tell *how* a task should be solved but easy to tell *if* and *how well* it has been solved.
- Still mainly a research field
- Strong links to dynamic optimization and optimal control fields.

Another example...

Pancake flipping



http://www.youtube.com/watch?v=W_gxLKSsSIE